

AI Model Landscape

April 2026

Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Artificial Analysis | April 2026 | Source: artificialanalysis.ai
For: AI Engineering Leads & Product Managers


Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Intelligence Index: composite score across reasoning, knowledge & instruction-following benchmarks (higher = stronger)

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33


gpt-oss-120B Tops Output Speed at 218 Tokens/s — More Than 4x Faster Than Claude Opus 4.6

Output speed in tokens per second; higher = faster inference (Muse Spark excluded — no published speed data)




From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models

API pricing per million tokens; Muse Spark excluded (no published pricing data)



No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

Intelligence Index vs. output tokens consumed during evaluation — models in the upper-left are most efficient



Proprietary vs. Open-Weight Leaders

Metric	Proprietary Leader	Open-Weight Leader
Intelligence	Gemini 3.1 Pro (57)	GLM-5.1 (51)
Speed (tokens/s)	Gemini 3 Flash (153)	gpt-oss-120B (218)
Price (\$/1M tokens)	Gemini 3 Flash (\$1.10)	gpt-oss-120B (\$0.30)

Key Insight: Open-weight models dominate speed and price; proprietary models lead intelligence by only 6 pts (57 vs 51)

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Actionable guidance for AI engineering leads and product managers

Intelligence Leaders

Gemini 3.1 Pro & GPT-5.4
Score 57 — co-lead the Intelligence Index
Anthropic's Claude Opus 4.6 trails at 53

Speed & Cost Leaders

gpt-oss-120B & DeepSeek V3.2 at \$0.30/1M
Tokens — 33x cheaper than Claude Opus 4.6
Open-weight models dominate throughput

The Open-Weight Signal

GLM-5.1 scores 51 (within 10% of leaders)
at \$2.10/1M tokens — best intelligence/price
ratio among open-weight models

Model Selection Guidance

Latency-critical → gpt-oss-120B
Complex reasoning → Gemini 3.1 Pro / GPT-5.4
Budget-constrained → GLM-5.1 or DeepSeek V3.2