AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Co

Source: Artificial Analysis (captured April 10, 2026) — 11-model benchmark across intelligence, speed, and pricing

Intelligence leaders

57 / 57

Fastest output speed

218 tok/s

Pricing spread

33× gap

Frontier is a two-way tie on intelligence (Gemini 3.1 Pro Preview and GPT-5.4 at 57).
Open-weight and smaller models dominate throughput (gpt-oss-120B at 218 tokens/s).
Cost ranges from $0.30 to $10 per 1M tokens across models with published pricing.
Core takeaway: model selection is workload-specific; no single model wins across all axes.

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Artificial Analysis Intelligence Index (higher is better) — full 11-model ranking

Rank	Model	Provider	Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

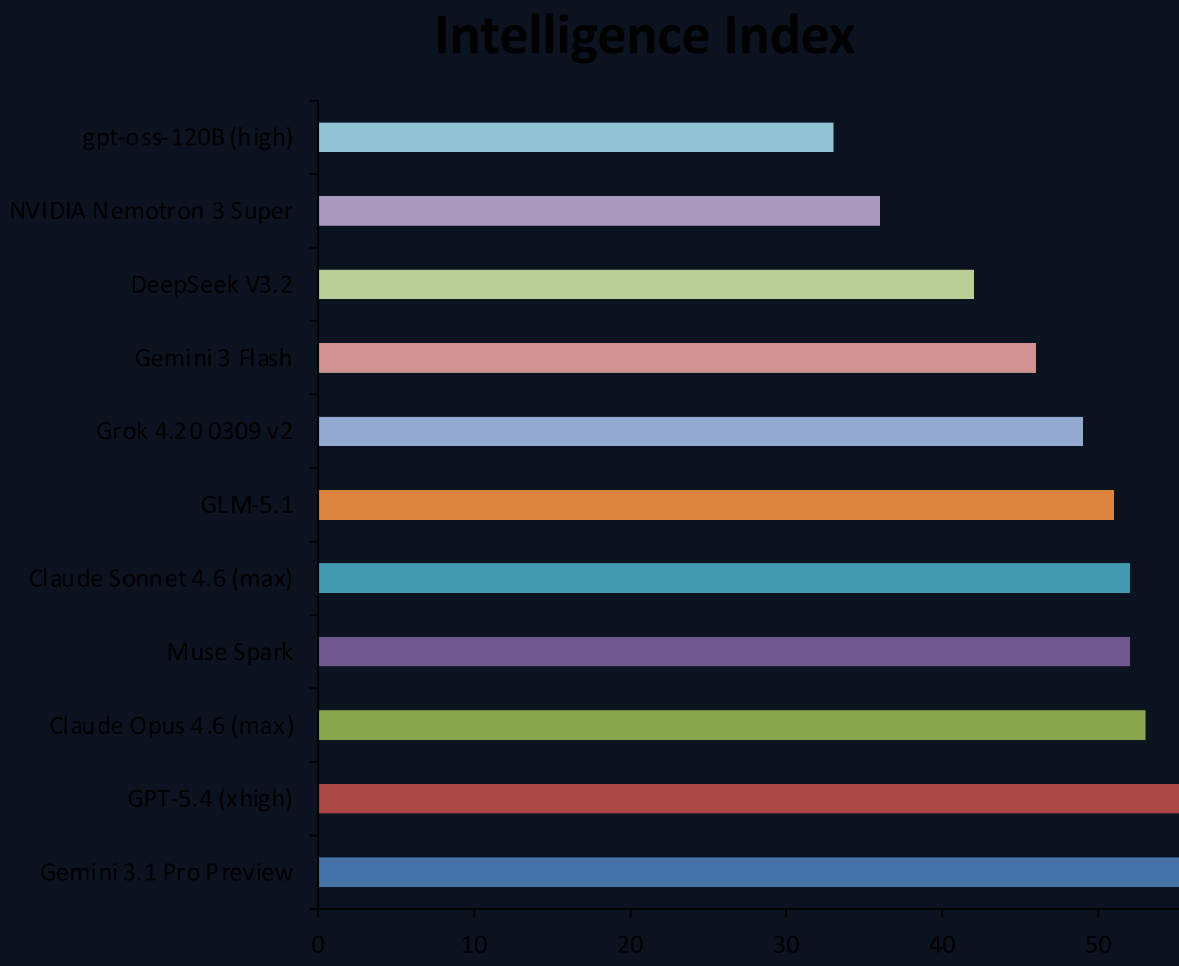

Spread: 57 → 33 (24 points)

Artificial Analysis Intelligence Index • April 2026

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster T

Output tokens per second (Muse Spark excluded — no speed data in source)

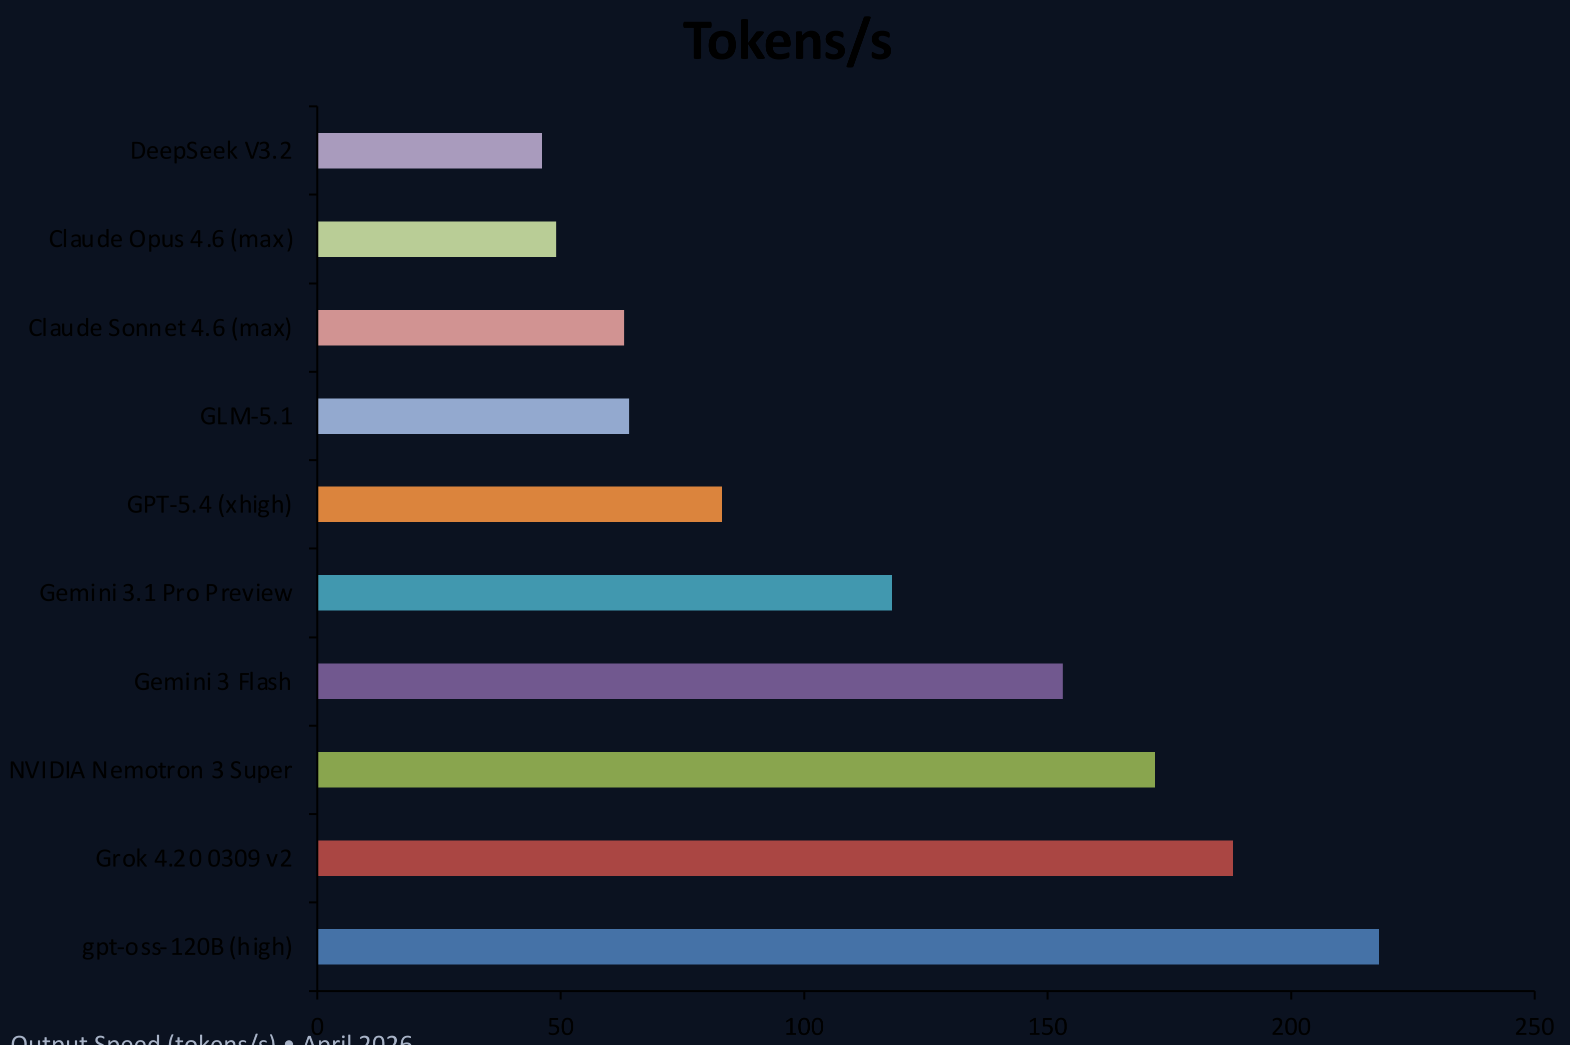

Open-weight models dominate throughput.
Fastest: gpt-oss-120B (high) at 218 tokens/s.
Claude Opus 4.6 (max): 49 tokens/s (4× slower t
Leaders in intelligence sit mid-pack on speed: Ge
Pattern: higher intelligence often trades off with

Output Speed (tokens/s) • April 2026

From $0.30 to $10 per Million Tokens: A 33x Pricing Gap Across Models

USD per 1M tokens (Muse Spark excluded — no published pricing in source)

Rank	Model	$/1M tokens
1	gpt-oss-120B (high)	$0.30
2	DeepSeek V3.2	$0.30
3	NVIDIA Nemotron 3 Super	$0.40
4	Gemini 3 Flash	$1.10
5	GLM-5.1	$2.10
6	Grok 4.20 0309 v2	$3.00
7	Gemini 3.1 Pro Preview	$4.50
8	GPT-5.4 (xhigh)	$5.60
9	Claude Sonnet 4.6 (max)	$6.00
10	Claude Opus 4.6 (max)	$10.00

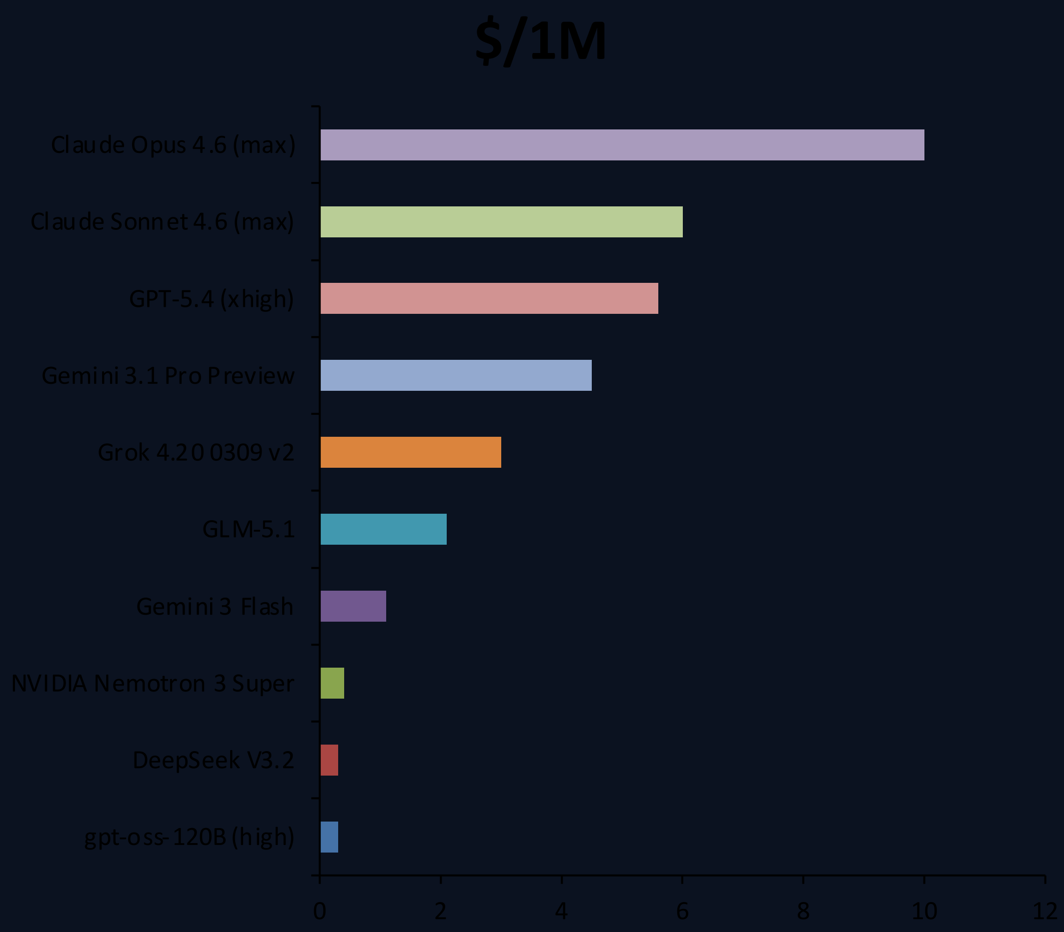

Cheapest: $0.30 • Most expensive: $10.00 (≈33×)

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-C

Scatter plot uses output tokens consumed during evaluation (token efficiency), not tokens/s

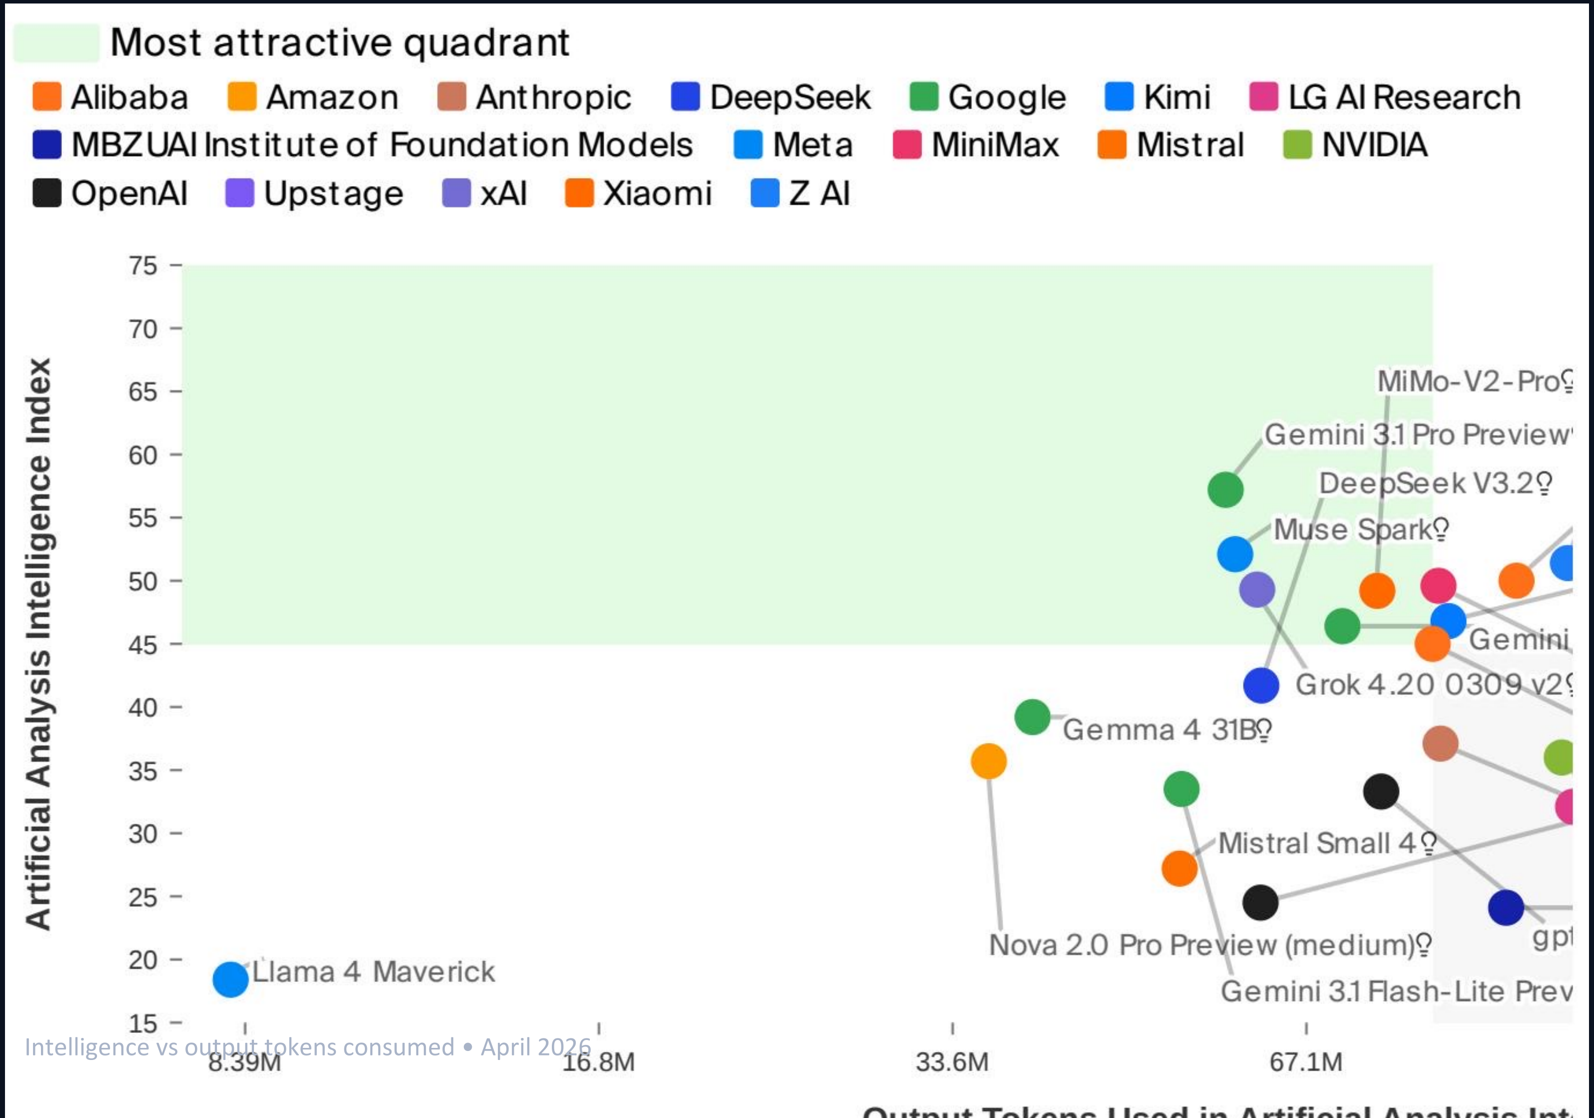

Intelligence vs output tokens consumed • April 2026

Metric	Prop leader	Score	Open leader	S
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	5
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	2
Price ($/1M tokens)	Gemini 3 Flash	$1.10	gpt-oss-120B (high)	$ 0

Intelligence, token-efficiency, and price pull in differ
Use-case fit beats global rank: optimize for your bot
Open-weight options win on speed and cost, while

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at $2.10/M

Actionable guidance for evaluation and routing

No single model dominates intelligence, speed, and price simultaneously.
Open weights are closing the gap: GLM-5.1 reaches 51 intelligence at $2.10/M
Route by workload:
- Quality-critical: prioritize top intelligence (57-tier).
- Latency-critical: prioritize high tokens/s leaders (218–153 tier).
- Budget at scale: prioritize $0.30–$1.10/M token tier.

Operationalize selection:
- Maintain a small model portfolio and switch by task class.
- Re-benchmark periodically; frontier shifts quickly.

Recommendations derived from April 2026 Artificial Analysis tables

Closest open-weight to frontier

51 vs 57

Best speed/price combo

$0.30 & 218

Balanced value callout

153 tok/s @ $1.10