


AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost


Competitive landscape of frontier AI models as of April 2026



Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33


gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6




Key takeaways

Open-weight and smaller models dominate raw throughput
gpt-oss-120B outputs 218 tokens/s (>4x Claude Opus 4.6).
Intelligence co-leaders sit mid-pack on speed (118 and 83 tokens/s)

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost F



Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Token

Actionable guidance for model selection

No universal winner: choose based on workload priorities.

Intelligence-critical workflows: consider top-tier proprietary leaders.

Latency/cost-sensitive flows: open-weight leaders often outperform.

Adopt portfolio routing across task classes and budgets.

Re-benchmark quarterly as open-weight models rapidly close gaps.