

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Gemini 3.1 Pro and GPT-5.4 tie at 57 on the Intelligence Index while a 33x pricing gap separates the cheapest and most expensive frontier models.

CO-LEADERS

57

Gemini 3.1 Pro + GPT-5.4

COST SPREAD

33x

\$0.30 to \$10.00 / 1M tokens

FASTEST OUTPUT

218

gpt-oss-120B tokens/s

Evaluation lens



Intelligence



Speed




Pricing

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57




The full 11-model ranking shows a tied frontier and a compressed mid-tier where open-weight models are narrowing the gap.

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6



From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models



Rank	Model	USD / 1M
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00

CHEAPEST

\$0.30

gpt-oss-120B + DeepSeek V3.2

Source: Artificial Analysis; captured April 10, 2026

MOST EXPENSIVE

\$10.00

Claude Opus 4.6


MID-RANGE VALUE

\$1.10

Gemini 3 Flash

Pricing data covers 10 models; Muse Spark is excluded because pricing data was not provided in the brief.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier



Proprietary vs. open-weight leaders

Metric	Proprietary Leader	Score	Open-Weight Leader	Score
Intelligence	Gemini 3.1 Pro Preview	57	GLM-5.1	51
Speed (tokens/s)	Gemini 3 Flash	153	gpt-oss-120B (high)	218
Price (\$/1M tokens)	Gemini 3 Flash	\$1.10	gpt-oss-120B (high)	\$0.30

Interpretation


High intelligence with fewer output tokens indicates better engineering and more efficient architectures.

Proprietary models still lead raw intelligence, but GLM-5.1 is only 6 points behind the co-leaders.

Open-weight gpt-oss-120B leads both speed and price, but scores 33 on intelligence.

Core insight: every workload makes an explicit trade-off among intelligence, token efficiency, and cost.

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens



- Prioritize intelligence** Use Gemini 3.1 Pro Preview or GPT-5.4 when reasoning, knowledge, or instruction following dominates value.
- Prioritize latency** Use gpt-oss-120B, Grok 4.20, NVIDIA Nemotron 3 Super, or Gemini 3 Flash for speed-sensitive experiences.
- Prioritize cost** Benchmark gpt-oss-120B, DeepSeek V3.2, NVIDIA Nemotron 3 Super, and Gemini 3 Flash for budget-sensitive scale.

Decision rule: select models by workload priority, then validate on your own prompts, latency budgets, and token-cost profile.