

AI Model Landscape April 2026: Benchmarking Intelligence, Speed, and Cost Across 11 Frontier Models

Independent analysis from Artificial Analysis | April 2026

Gemini 3.1 Pro Preview and GPT-5.4 Co-Lead the Intelligence Index at 57

Google and OpenAI are neck-and-neck at the frontier. Anthropic's Claude Opus 4.6 trails by 4 points, while Meta's newly released Muse Spark enters at a strong 52. Open-weight models such as GLM-5.1 and Grok 4.20 occupy the mid-tier, demonstrating that the gap between proprietary and open-weight models continues to narrow.

Rank	Model	Provider	Intelligence Index
1	Gemini 3.1 Pro Preview	Google	57
2	GPT-5.4 (xhigh)	OpenAI	57
3	Claude Opus 4.6 (max)	Anthropic	53
4	Muse Spark	Meta	52
5	Claude Sonnet 4.6 (max)	Anthropic	52
6	GLM-5.1	Z AI	51
7	Grok 4.20 0309 v2	xAI	49
8	Gemini 3 Flash	Google	46
9	DeepSeek V3.2	DeepSeek	42
10	NVIDIA Nemotron 3 Super	NVIDIA	36
11	gpt-oss-120B (high)	OpenAI	33

gpt-oss-120B Tops Output Speed at 218 Tokens/s -- More Than 4x Faster Than Claude Opus 4.6

The fastest model, gpt-oss-120B (high), outputs 218 tokens per second -- more than 4x the speed of Claude Opus 4.6 and nearly 5x that of DeepSeek V3.2. This highlights the trade-off between raw intelligence and throughput: the two intelligence leaders (Gemini 3.1 Pro Preview and GPT-5.4) sit in the middle of the speed table at 118 and 83 tokens/s respectively.

Rank	Model	Output Tokens/s
1	gpt-oss-120B (high)	218
2	Grok 4.20 0309 v2	188
3	NVIDIA Nemotron 3 Super	172
4	Gemini 3 Flash	153
5	Gemini 3.1 Pro Preview	118
6	GPT-5.4 (xhigh)	83
7	GLM-5.1	64
8	Claude Sonnet 4.6 (max)	63
9	Claude Opus 4.6 (max)	49
10	DeepSeek V3.2	46

Note: Muse Spark has no published speed data in this benchmark.

From \$0.30 to \$10 per Million Tokens: A 33x Pricing Gap Across Models


The cheapest options -- gpt-oss-120B and DeepSeek V3.2 at \$0.30 per million tokens -- are over 33x less expensive than Claude Opus 4.6 at \$10.00. For budget-sensitive workloads, this price differential is a decisive factor. Notably, Gemini 3 Flash offers a compelling mid-range value at \$1.10, pairing solid speed (153 tokens/s) with reasonable intelligence (46).

Rank	Model	USD per 1M Tokens
1	gpt-oss-120B (high)	\$0.30
2	DeepSeek V3.2	\$0.30
3	NVIDIA Nemotron 3 Super	\$0.40
4	Gemini 3 Flash	\$1.10
5	GLM-5.1	\$2.10
6	Grok 4.20 0309 v2	\$3.00
7	Gemini 3.1 Pro Preview	\$4.50
8	GPT-5.4 (xhigh)	\$5.60
9	Claude Sonnet 4.6 (max)	\$6.00
10	Claude Opus 4.6 (max)	\$10.00

Note: Muse Spark has no published pricing data in this benchmark.

No Model Wins on All Three Axes: Navigating the Intelligence-Efficiency-Cost Frontier

A scatter plot of Intelligence Index versus output tokens consumed during evaluation reveals which models occupy the "most attractive quadrant" -- high intelligence with efficient token usage.



Metric

Proprietary Leader

Score

Open-Weight Leader

Score

Open Weights Close the Gap: GLM-5.1 Reaches 51 Intelligence at \$2.10/M Tokens

Key findings and recommendations for AI engineering leads and product managers:

Two-way tie at the top: Gemini 3.1 Pro Preview and GPT-5.4 (xhigh) both score 57 on the Artificial Analysis Intelligence Index, making them the co-leaders as of April 2026.

Speed and intelligence are inversely correlated: The fastest model (gpt-oss-120B at 218 tokens/s) scores only 33 on intelligence, while the smartest models operate at 49-118 tokens/s.

Massive pricing spread: A 33x cost gap separates the cheapest (\$0.30/1M tokens for gpt-oss-120B and DeepSeek V3.2) from the most expensive (\$10.00/1M tokens for Claude Opus 4.6).

Open-weight models are closing the gap: GLM-5.1 at 51 intelligence is within 10% of the proprietary leaders, while also being significantly cheaper at \$2.10/1M tokens.

No single model wins on all three axes: Choosing a model requires explicit trade-offs among intelligence, speed, and cost depending on the use case.

Recommendation: Prioritize intelligence for research and complex reasoning tasks; prioritize speed for real-time applications; prioritize cost for high-volume, budget-sensitive deployments. Open-weight models like GLM-5.1 and gpt-oss-120B now offer compelling alternatives for specific workload profiles.