

The 2025 LLM Usage Landscape

Real-World Data from OpenRouter

SOURCE: OpenRouter Rankings · March 2025

Based on millions of live API calls — revealing who really leads the LLM market

26T+

Tokens processed
this week

10

Providers in
Top 10

54.3%

Chinese provider
combined share

This Week's Top 10 Models

OpenRouter Rankings · Week of March 2025 | by
Token Usage


Rank	Model	Provider	Weekly Tokens	WoW Change
1	MiMo-V2-Pro	Xiaomi	3.96T	+165%
2	Step 3.5 Flash (free)	Stepfun	1.49T	+1%
3	MiniMax M2.7	MiniMax	1.29T	+253%
4	DeepSeek V3.2	DeepSeek	1.24T	+9%
5	Claude Sonnet 4.6	Anthropic	1.04T	-0%
6	Claude Opus 4.6	Anthropic	995B	+2%
7	Gemini 3 Flash Preview	Google	972B	+3%
8	GLM 5 Turbo	Zhipu (z-ai)	968B	-6%
9	MiniMax M2.5	MiniMax	909B	-30%
KEY: Xiaomi MiMo-V2-Pro leads at 3.96T tokens (+165% WoW) · Chinese providers hold 6 of 10 spots · MiniMax M2.7 fastest-growing at +253%				

= Chinese provider

Provider Market Share Breakdown

By Weekly Token Volume · March 2025

Token Volume Share by Provider



25.9%

Xiaomi leads — more than double second-place Google

54.3%

Combined Chinese provider share exceeds half the market

7.9%

OpenAI ranks only 5th — below Xiaomi, Google, Anthropic, and MiniMax


No single player >30%

Market remains highly fragmented across 8+ providers

Note: Provider-level share covers all models & apps on OpenRouter, not just the top-10 model list.

Shifts in the Market Landscape

Key Competitive Dynamics · March 2025



Market grew 13× in one year: from ~2T tokens/week (Mar 2025) to ~26T tokens/week (Mar 2026)

1 Chinese Models Break Through Globally

On OpenRouter (an international platform), Chinese providers now hold >54% combined share. Xiaomi, DeepSeek, MiniMax, Stepfun, and Zhipu each hold top-10 positions.

2 Flash/Turbo Variants Outpace Flagships


Step 3.5 Flash, GLM 5 Turbo, Gemini 3 Flash — 'fast & affordable' models far outpace their flagship counterparts. Speed and cost beat peak benchmarks in production.

3 Anthropic Growth Has Stalled

Both Claude models (Sonnet 4.6 and Opus 4.6) appear in the top 10, but combined growth is ~0%, suggesting a plateau despite strong absolute volume (~2T tokens).

This Week's Rising Stars: Fastest-Growing Models

Fastest-Growing Models (WoW %)



MiniMax M2.7 +253%

New generation model. MiniMax's M-series has shown rapid iteration; M2.7 nearly quadruples usage in a single week. Strong cost-performance ratio driving adoption.


MiMo-V2-Pro +165%

Xiaomi's reasoning-focused model. Already the #1 model by volume, still growing aggressively — indicative of sustained enterprise on-ramp.

Cline +128%

An independent open-source VS Code coding agent. Its surge confirms that AI coding tools are the dominant LLM use case in developer workflows.

Top 10 Apps by Daily Token Consumption



AI Coding Tools Dominate

4 of the top 5 apps — OpenClaw, Kilo Code, Claude Code, and Cline — are AI coding assistants. Developer tooling has emerged as the single largest LLM consumption use case.

OpenClaw Is the Runaway Leader

OpenClaw consumes 829B tokens/day — nearly 3× second-place Kilo Code (271B). As a general AI productivity assistant, it demonstrates the massive scale of broad-use agents.

Entertainment Holds Steady

Janitor AI (29.2B) and ISEKAI ZERO (28.8B) hold the #5 and #6 spots, showing that roleplay and entertainment remain significant, if niche, LLM applications.

Other

Key Insights and Industry Implications

Chinese Models Now Dominate the Global Market

On the international OpenRouter platform, Chinese providers (Xiaomi, MiniMax, Stepfun, DeepSeek, Zhipu) hold a combined 54.3% of token volume. This is no longer an emerging story — it is the new reality. Product and investment teams ignoring Chinese LLMs face a significant competitive blind spot.

The Market Is Highly Fragmented — No Winner Yet

The top 10 models span 8 different providers. No provider exceeds 30% share. This level of fragmentation signals that the LLM stack is still unsettled — moats are thin and switching costs are low.

Implication for R&D: Abstraction layers and multi-model strategies reduce lock-in risk.

AI Coding Agents Are the Dominant Use Case

5 of the top 10 apps by daily token consumption are AI coding tools. OpenClaw leads at 829B tokens/day, nearly 3x second place. Developer tooling is where LLMs are converting into real revenue today.

Fast & Affordable Beats Peak Performance in Production

Flash/Turbo variants (Step 3.5 Flash, GLM 5 Turbo, Gemini 3 Flash) far outpace their flagship counterparts in actual usage. The market votes for speed and cost over leaderboard benchmark scores.

Implication for product: Optimize for latency and cost-efficiency, not just raw capability.