

The 2025 LLM Usage Landscape: Real-World Data from OpenRouter

OpenRouter Rankings • Latest week of March 2025

26T+ tokens processed this week

Real usage across millions of OpenRouter users; reflects production demand, not benchmarks.


based on real usage data from millions of users accessing models through OpenRouter.

Top Models

Weekly usage of models across OpenRouter


26T

19.5T



This Week's Top 10 Models

Ranked by weekly token usage



Highlight: Xiaomi MiMo-V2-Pro leads at 3.96T weekly tokens (+165% WoW).


Source: OpenRouter Rankings (latest week of March 2025)

Provider	Tokens	WoW
Xiaomi	3.96T	+165%
Stepfun	1.49T	+1%
MiniMax	1.29T	+253%
DeepSeek	1.24T	+9%
Anthropic	1.04T	-0%
Anthropic	995B	+2%
Google	972B	+3%
Zhipu (z-ai)	968B	-6%
MiniMax	909B	-30%
x-ai	599B	+32%

Tokens shown as billions (B)


Provider Market Share Breakdown

Share of weekly token volume across all usage



Xiaomi leads at 25.9% — more than 2× #2 Google (11.8%).
Combined Chinese provider share (Xiaomi + MiniMax + Stepfun + DeepSeek + Zhipu) ≈ 54.3%.
OpenAI ranks #5 at 7.9%; market remains fragmented (no provider >30%).
Note: provider share covers all models & apps (not just top-10 models).

Source: OpenRouter provider market share table



Shifts in the Market Landscape

Competitive shifts and usage-driven preferences

Chinese models now exceed 54% combined share on an international platform.
Anthropic maintains 2 top-10 models (~2T combined), but growth has stalled.
Flash/Turbo variants (fast + affordable) outpace flagships in real-world demand.

Signal: usage is scaling rapidly

Weekly tokens expanded from ~2T to ~26T.
Competitive advantage is translating into volume.

Based on real usage data from millions of users accessing models through OpenRouter.

Top Models

Weekly usage of models across OpenRouter


26T

19.5T



This Week's Rising Stars: Fastest-Growing Models

Fastest week-over-week growth




Also notable: Ampere (+576%), Codex (+112%), Hermes Agent (+92%).
Interpretation: growth concentrates in cost-performance leaders and coding-agent workflows.

1.	 MiMo-V2-Pro by xiaomi	3.96T tokens ↑ 165%
2.	 Step 3.5 Flash (free) by stepfun	1.49T tokens ↑ 1%
3.	 MiniMax M2.7 by minimax	1.29T tokens ↑ 253%
4.	 DeepSeek V3.2 by deepseek	1.24T tokens ↑ 9%
5.	 Claude Sonnet 4.6 by anthropic	1.04T tokens ↓ 0%
6.	 Claude Opus 4.6 by anthropic	995B tokens ↑ 2%
7.	 Gemini 3 Flash Preview by google	972B tokens ↑ 3%
8.	 GLM 5 Turbo by z.ai	968B tokens ↓ 6%

App Ecosystem: Who's Calling These Models

Top apps by daily token consumption



1.		OpenClaw > The AI that actually does things	829B tokens
2.		Kilo Code > AI coding agent for VS Code	271B tokens
3.		Claude Code > The AI for problem solvers	78B tokens
4.		Cline > Autonomous coding agent right in your IDE	69.1B tokens
5.		Janitor AI > Character chat and creation	29.2B tokens
6.		ISEKAI ZERO > AI adventures. Travel with your favorite characters	28.8B tokens
7.		Descript > AI Video & Podcast Editor	27B tokens
8.		Hermes Agent > new	19.2B tokens
9.		Roo Code > A whole dev team of AI agents in your editor	11.8B tokens
10.		Scripty > new	9.95B tokens

[Visit the new App & Agent Rankings →](#)

Top 4 apps are AI coding tools (OpenClaw, Kilo Code, Claude Code, Cline).
AI coding agents are the single largest LLM consumption scenario.
OpenClaw alone consumes 829B tokens/day (~3× #2).
Source: OpenRouter top apps ranking

Key Insights and Industry Implications

What this implies for R&D, product, and investment

Chinese models now hold a dominant global share on OpenRouter (54%+ combined providers).
AI coding agents are the largest consumption scenario; distribution is app-led.
The market remains fragmented: top-10 spans 8 providers; no provider >30%.
Fast & affordable variants (Flash/Turbo) outpace flagships in real usage.

What to monitor next (from observed patterns)

Cost/perf improvements that translate into sustained token share.
Coding-agent tooling that captures developer workflows (VS Code + CLI).
Whether stalled incumbents re-accelerate via faster/cheaper variants.