

INDUSTRY REPORT · 2025

The 2025 LLM Usage Landscape

Real-World Data from OpenRouter

Based on millions of API calls · Latest week of March 2025

HEADLINE METRIC — WEEKLY TOKEN VOLUME

26T+ tokens processed

Data Source: OpenRouter Rankings

Period: March 2025

Platform: openrouter.ai

26T

tokens / week

This Week's Top 10 Models

Xiaomi's MiMo-V2-Pro leads with 3.96T tokens — Chinese providers hold 6 of the top 10 spots on this international platform


RANK	MODEL	PROVIDER	WEEKLY TOKEN USAGE	VOLUME BAR	WoW CHANGE
1	MiMo-V2-Pro Xiaomi · Chinese	Xiaomi	3.96T		↗️ +165%
2	Step 3.5 Flash (free) Stepfun · Chinese	Stepfun	1.49T		+1%
3	MiniMax M2.7 MiniMax · Chinese	MiniMax	1.29T		↗️ +253%
4	DeepSeek V3.2 DeepSeek · Chinese	DeepSeek	1.24T		+9%
5	Claude Sonnet 4.6 Anthropic · US	Anthropic	1.04T		-0%
6	Claude Opus 4.6 Anthropic · US	Anthropic	995B		+2%
7	Gemini 3 Flash Preview Google · US	Google	972B		+3%
8	GLM 5 Turbo Zhipu (z-ai) · Chinese	Zhipu (z-ai)	968B		↘️ -6%
9	MiniMax M2.5 MiniMax · Chinese	MiniMax	909B		↘️ -30%
10	Grok 4.1 Fast x-ai · US	x-ai	599B		↗️ +32%

Chinese providers (Xiaomi, Stepfun, MiniMax ×2, DeepSeek, Zhipu) = 6 of 10 spots


Anthropic #5 + #6 combine for ~2T tokens but both show near-zero WoW growth

Provider Market Share Breakdown

Xiaomi dominates at 25.9%; Chinese providers combined (Xiaomi + MiniMax + Stepfun + DeepSeek + Zhipu) hold 54.3% of global token volume



Provider Rankings (Weekly Token Volume)



Chinese Providers Combined Share
54.3% (Xiaomi + MiniMax + Stepfun + DeepSeek + Zhipu)

Note: Provider share covers all models and apps on OpenRouter, not limited to the top-10 models ranking.

Shifts in the Market Landscape

Token volume grew 13× in one year (~2T to 26T weekly); Chinese models now define the global LLM market

MARKET GROWTH · MARCH 2025 vs MARCH 2024

~2T tokens/week → **26T+ tokens/week** = **13×** year-over-year market expansion

1

Chinese Models Dominate Globally

54.3%

combined share of all token volume

On OpenRouter — an international platform — Chinese providers now exceed the combined share of Google, Anthropic, and OpenAI.

Xiaomi

DeepSeek

MiniMax

Stepfun

Zhipu

Strategic implication:

Chinese model cost-performance advantage has translated into market dominance.

2

Flash/Turbo Variants Far Outpace Flagships

Top models by usage

Step 3.5 Flash (free)

#2 — 1.49T

GLM 5 Turbo

#8 — 968B

Gemini 3 Flash Preview

#7 — 972B

Users vote with their feet: in real deployments, speed and cost often matter more than peak benchmark performance.

Strategic implication:

Optimize for latency and cost, not just MMLU scores, to capture real usage.

3

Anthropic's Growth Has Stalled

WoW growth for Anthropic models

Claude Sonnet 4.6

#5 · 1.04T tokens/week

-0%

Claude Opus 4.6

#6 · 995B tokens/week

+2%

Despite 2 models in the top 10 with combined ~2T tokens, Anthropic sees near-zero growth while rivals post **+165% and +253% WoW.**

Strategic implication:

Premium positioning may be losing ground to cost competitive Chinese alternatives.

This Week's Rising Stars: Fastest-Growing Models

MiniMax M2.7 surges +253% WoW — agentic and coding-specialized models lead the growth wave this week

FASTEST GROWTH · #1

MiniMax M2.7
MiniMax · Chinese provider



+253%
week-over-week growth

1.29T tokens · #3 overall this week

FASTEST GROWTH · #2

MiMo-V2-Pro
Xiaomi · Chinese provider



+165%
week-over-week growth

3.96T tokens · #1 overall this week

FASTEST GROWTH · #3

Cline
Coding agent · Independent dev



+128%
week-over-week growth

VS Code autonomous coding agent

Additional Fast-Movers This Week

MODEL	PROVIDER	WoW GROWTH	DRIVER
Ampere	Ampere	+576% 	New entrant momentum — first-week surge
Codex	OpenAI	+112% 	Coding agent specialization driving adoption
Hermes Agent	Nous Research	+92% 	Agentic workflow adoption expanding


Growth Drivers This Week:
① Chinese model cost-performance advantage
② Coding & agentic specialization

③ New entrant momentum (Ampere: +576%)
④ Autonomous coding agent adoption curve accelerating

App Ecosystem: Who's Calling These Models

4 of the top 5 apps are AI coding tools — developer tooling is the #1 LLM consumption use case today

Top Apps by Daily Token Volume (B = billion)



Key Insights

AI Coding Agents Dominate

Top 4 apps are all AI coding tools:
OpenClaw · Kilo Code · Claude Code · Cline

829B tokens/day

OpenClaw alone — nearly 3× the second-place app
Developer tooling = the largest LLM ROI zone today

→ Investment signal: Coding infrastructure is where usage concentrates

Consumer Adoption Persists

Entertainment/roleplay apps maintain strong presence:
Janitor AI (29.2B/day) · ISEKAI ZERO (28.8B/day)

Consumer LLM adoption extends well beyond productivity.

→ Long-tail entertainment demand remains a significant volume driver

Descript AI video editor (27B/day) signals media production next

Market Structure: No Monopoly in Apps Either

Top 10 apps span coding, entertainment, media, and agents.
Demand is diversified — no single app category monopolizes LLM consumption.

→ Horizontal platform plays (OpenRouter-style) benefit most from fragmentation

Key Insights and Industry Implications

Four forces reshaping the LLM industry — implications for R&D, product, and investment decisions

1

Chinese Models Have Won the Global Usage Race

On OpenRouter — the largest neutral LLM API platform — Chinese providers hold

54.3% combined token share

This is not a forecast — it is current market reality.

Implication → Global players must reassess China-origin model partnerships

2

AI Coding Agents = The Primary Commercial Use Case

5 of the top 10 apps are coding tools.

OpenClaw alone: 829B tokens/day.

Developer tooling is where LLM ROI is being realized today.

Implication → Investment and R&D should follow actual usage concentration

3

No Winner-Takes-All — The Market Is Highly Fragmented

Top 10 models span 8 different providers.

Competition is fierce; switching costs are low.

No single model or provider has locked in the market.

Implication → Multi-model strategies outperform single-provider lock-in bets

4

Fast & Affordable Beats Flagship in Real Usage

Flash/Turbo variants dominate usage

over premium flagship models.

Benchmark superiority does not translate to adoption.

Implication → Optimize for speed and cost first; users will choose accordingly