

The 2025 LLM Usage Landscape: Real-World Data from OpenRouter

Real usage data from millions of users

Latest week of March 2025

>26T

tokens processed this week

OpenRouter Rankings aggregates LLM API usage across a broad international ecosystem, revealing adoption patterns beyond benchmarks.

What this report covers

- Model demand
Top 10 weekly token usage
- Provider power
Market share by token volume
- Growth signals
Fastest-growing models and apps
- Implications
R&D, product, investment priorities

This Week's Top 10 Models

Weekly token usage and week-over-week change

Rank	Model	Provider	Weekly Tokens	WoW
1	MiMo-V2-Pro	Xiaomi	3.96T	+165%
2	Step 3.5 Flash (free)	Stepfun	1.49T	+1%
3	MiniMax M2.7	MiniMax	1.29T	+253%
4	DeepSeek V3.2	DeepSeek	1.24T	+9%
5	Claude Sonnet 4.6	Anthropic	1.04T	-0%
6	Claude Opus 4.6	Anthropic	995B	+2%
7	Gemini 3 Flash Preview	Google	972B	+3%
8	GLM 5 Turbo	Zhipu (z-ai)	968B	-6%
9	MiniMax M2.5	MiniMax	909B	-30%
10	Grok 4.1 Fast	x-ai	599B	+32%

Key signals

3.96T

MiMo-V2-Pro leads by a wide margin in weekly tokens.

Chinese providers hold 6 of the top 10 spots

MiniMax M2.7 posts the strongest growth: +253%

Provider Market Share Breakdown

Provider-level share by token volume across all models and apps

Market share ranking



54.3%

Combined Chinese provider share
Xiaomi + MiniMax + Stepfun +
DeepSeek + Zhipu

Xiaomi leads at 25.9%
More than double #2 Google

Fragmented field
No single provider exceeds 30%
OpenAI ranks fifth at 7.9%


Note: Provider figures represent total provider-level market share, including usage outside the top-10 model table.

Shifts in the Market Landscape

Competitive dynamics visible in real-world API usage

Weekly usage trend

OpenRouter weekly token usage grew from ~2T to ~26T in the referenced March 2025–March 2026 trend chart.



Chinese models break through

Chinese provider models now account for over half of all usage on an international platform.



Fast beats flagship

Flash/Turbo variants see higher usage because speed and cost often matter more than peak performance.



No settled winner


Top 10 spans 8 providers; no single model or provider monopolizes the market.

Anthropic maintains two top-10 models, but source brief notes growth has stalled.

This Week's Rising Stars: Fastest-Growing Models

Week-over-week growth signals from the source brief

Fastest-growing leaderboard



Why they surge

- Cost-performance**
MiMo, DeepSeek, and MiniMax are cited for significant cost-performance advantage.
- Workflow pull**
Cline and Codex show agent/coding workflows can drive rapid token growth.
- Speed matters**
Fast and affordable variants often outpace flagship models in real usage.

Source: Fastest-Growing Models table and Industry Trend Observations in source_brief.md.

App Ecosystem: Who's Calling These Models

Top apps by daily token consumption

Rank	App	Description	Daily Tokens
1	OpenClaw	AI productivity assistant	829B
2	Kilo Code	VS Code AI coding assistant	271B
3	Claude Code	Anthropic CLI coding tool	78B
4	Cline	VS Code autonomous coding agent	69.1B
5	Janitor AI	Character roleplay chat	29.2B
6	ISEKAI ZERO	AI adventure game	28.8B
7	Descript	AI video/podcast editor	27B
8	Hermes Agent	Nous Research AI agent	19.2B
9	Roo Code	VS Code AI coding assistant	11.8B
10	Scripty	Newly launched tool	9.95B

829B

OpenClaw daily token volume
nearly 3× the #2 app

Top 4 apps are all
AI coding tools

Coding agents are cited as
the single largest LLM
consumption scenario.

Key Insights and Industry Implications

What the usage landscape means for builders and capital allocators

1

Chinese providers are now a major global force

Combined share for Xiaomi, MiniMax, Stepfun, DeepSeek, and Zhipu is approximately 54.3%.

R&D: benchmark against cost-performance leaders, not only frontier incumbents.

2

AI coding agents dominate consumption

The top four apps are coding tools; five of the top ten apps are coding tools.

Product: prioritize developer workflows where token intensity and willingness to adopt are highest.

3

Competition remains fragmented

No single provider exceeds 30%; the top 10 models span eight providers.

Investment: avoid winner-take-all assumptions; track distribution and switching costs.

4

Fast and affordable models win usage

Flash/Turbo variants far outpace flagship models in real-world applications.

Strategy: optimize latency, price, and reliability alongside peak model quality.