

Measuring AI Ability to Complete Long Software Tasks

Time Horizon Doubling Every ~7 Months

Headline finding: The 50% task-completion time horizon has doubled approximately every 207 days (~7 months) since 2019 ($R^2 = 0.97$).

Audience: AI safety researchers, frontier AI lab leadership, policy analysts

Source: METR study on autonomous software-task completion horizons

A New Metric: 50% Task-Completion Time Horizon

1) Build task suite

Assemble 170 software tasks from HCAST, RE-Bench, and SWAA spanning seconds to hours.

2) Measure outcomes

Collect human task durations and model success/failure across the same task set.

3) Fit logistic curve

Estimate success probability vs. task duration; read off duration at 50% success.

Family	Length	Description
find_shell_script	3 seconds	Multiple choice: “Which file is a shell script?” Choices: “run.sh”, “run.txt”, “run.py”, “run.md”
wikipedia_research	1 minute	Research simple factual information from Wikipedia and provide accurate answers to straightforward questions.
oxdna_simple	9 minutes	Detect and fix a bug in the input files for a molecular dynamics simulation using the oxDNA package.
munge_data	56 minutes	Write a Python script to transform JSON data from one format to another by inferring the conversion rules from provided example files.
cuda_backtesting	8 hours	Speed up a Python backtesting tool for trade executions by implementing custom CUDA kernels while preserving all functionality, aiming for a 30x performance improvement.

Interpretation

The metric converts broad benchmark outcomes into a single, intuitive unit: “how long a skilled human task AI can finish with 50% reliability.” This supports trend tracking across model generations.

170 Tasks Spanning 3 Seconds to 8 Hours of Human Effort


Source	Task count	Duration range	Example tasks
HCAST	97	3 sec to 1+ hour	find_shell_script (3s); json_nest_1 (4m)
RE-Bench	7	30 min to 8 hours	optimize_cython (30m); cuda_backtesting (8h)
SWAA	66	seconds to hours	debugging, scripting, end-to-end software tasks

Suite composition is deliberately heterogeneous to sample varied software workflows and avoid narrow benchmark overfitting.

Methodological note

Task timing is anchored to skilled human completion time, then linked to model success probabilities. This creates a shared “time” axis for comparing models of different vintages and architectures.

AI Task Horizon Grew from 2 Seconds (GPT-2) to ~1 Hour 50 Minutes (o3) in 6 Years




Trend statistics

- Doubling time: ~207 days (~7 months)
- Exponential fit: $R^2 = 0.97$
- GPT-2 (2019): ~2 seconds
- GPT-4 (0314): ~5 minutes
- o3 (2025): ~1 hour 50 minutes

Interpretation: model capability growth in autonomous coding appears persistent and strongly exponential over 2019–2025.

Success Probability Drops Sharply as Task Length Increases for All Models



Model comparisons

- Older models fail quickly as duration rises.
- Frontier models shift the curve rightward.
- Claude 3.7 Sonnet: ~54 min horizon.
- o3: ~1h50m horizon at 50% success.
- All models still degrade on long tasks.

The family of curves suggests broad reliability gains, not just isolated benchmark spikes.

Reasoning, Tool Use, and Error Recovery Drive Time Horizon Growth

Improved logical reasoning

Better decomposition and planning increase completion on multi-step coding tasks.

Stronger tool use

More effective use of development tools and environment feedback loops.

Reliability and self-monitoring

Greater consistency and self-awareness reduce brittle execution failures.

Adaptation to mistakes

Agents more often detect, diagnose, and recover from intermediate errors.

Limitation: performance is substantially weaker on less structured, “messier” tasks, indicating current autonomy remains context-dependent.

Extrapolation Predicts Month-Long Task Automation by 2028–2031

Forward projection (naive trend extension)

- Historical basis: ~207-day doubling at 50% horizon.
- 2023–2025 growth appears ~20% faster than 2019–2025 average.
- Crossing 167 work-hours (~1 month) estimated between mid-2028 and mid-2031.
- Implication: many month-scale software projects may become automatable.

These forecasts are scenario indicators, not physical limits.

Caveats

- External validity remains uncertain.
- Real-world task distributions may differ.
- Trend slope can slow, speed up, or break.

Policy relevance

Planning windows for governance, evaluation standards, and deployment safeguards may be measured in years, not decades.

Takeaways: Exponential Autonomous Capability Growth Demands Proactive Safety Governance

1) 50% time horizon is a robust, intuitive autonomy metric ($R^2 = 0.97$ over 6 years).

2) Capability doubled roughly every ~7 months, with signs of post-2024 acceleration.

3) Frontier systems now reach ~2 hours of skilled human-equivalent software work (o3).

4) External validity is unresolved: benchmark tasks may not fully mirror real-world labor.

5) Key limits persist on unstructured tasks and real-world generalization.

Governance implication: capability tracking should be paired with staged deployment controls, stress-tests on messy tasks, and continuous post-deployment monitoring.