

DEEPSEEK TECHNICAL BRIEFING

DeepSeek-V2: Efficient MoE Language Model with Multi-Head Latent Attention

A Strong, Economical, and Efficient MoE Language Model

236B / 21B

Total / Activated Parameters

128K

Context Length

MLA + DeepSeekMoE

Core Innovations

Audience: ML researchers, NLP engineers, graduate students • Based on DeepSeek-V2 paper benchmarks