

DeepSeek-V2: Efficient MoE Language Model with Multi-Head Latent Attention

236B total params | 21B activated params | 128K context length

A Strong, Economical, and Efficient MoE Language Model

Target audience: ML researchers, NLP engineers, graduate students

Based on DeepSeek-V2 technical report

The LLM Scaling Dilemma: Performance vs. Cost and Speed

Scaling Trend

Larger dense models improve reasoning and language quality
But pretraining cost grows rapidly with model size
Inference becomes memory-bound by KV cache at long context

Design Requirement

Reduce active compute without sacrificing capability
Compress attention state for fast long-context decoding
Maintain high benchmark performance
=> DeepSeek-V2 combines MLA + DeepSeekMoE

DeepSeek-V2 Architecture: MLA + DeepSeekMoE in Transformer Blocks





Figure reference: DeepSeek-V2 block design integrating Multi-Head Latent Attention and DeepSeekMoE

Multi-Head Latent Attention Compresses KV Cache by 93.3%



MLA mechanism

Low-rank joint compression for keys and values
Decoupled RoPE preserves positional expressiveness
Avoids quality loss seen in naive cache reduction
93.3% KV cache reduction vs DeepSeek 67B

DeepSeekMoE: Fine-Grained Experts with Shared Expert Isolation

Core design choices

Fine-grained expert segmentation improves specialization

Shared experts capture common knowledge; routed experts model diversity

Top-K routing with device-limited routing lowers communication overhead

Auxiliary load-balancing losses stabilize expert utilization

Node-limited routing during training controls cross-node traffic

Result: 236B total parameters with only 21B activated per token

42.5% Training Cost Reduction and 5.76× Generation Throughput vs. DeepSeek 67B

DeepSeek 67B


Dense baseline architecture
Higher training cost
Larger KV cache
Lower max generation throughput

DeepSeek-V2

42.5% lower training cost
93.3% smaller KV cache
5.76× higher max generation throughput

Relative improvements reported against DeepSeek 67B; absolute values omitted by design.

DeepSeek-V2 Tops MMLU at Only 21B Activated Parameters



Key points

DeepSeek-V2: 21B, 78.5

LLaMA 3 70B: 70B, 77.0

Mixtral 8x22B: 39B, 77.6

Qwen1.5 72B: 72B, 75.6

Pareto-optimal: higher MMLU at lower activated params

DeepSeek-V2 Chat (RL) Rivals Closed-Source Models on Chinese Benchmarks

Alignment pipeline: SFT on 1.5M sessions → GRPO-based RL

Benchmark	DeepSeek-V2 Chat (RL)
AlpacaEval 2.0 (length-controlled win rate)	38.9%
MT-Bench (overall)	8.97
AlignBench (overall, Chinese)	7.91

On AlignBench (Chinese), DeepSeek-V2 Chat (RL) outperforms all open-source and most closed-source models.

Key Takeaways: Architectural Innovation Unlocks Efficient Scaling

Summary

MLA is a superior attention design for inference efficiency, delivering 93.3% KV cache reduction.

DeepSeekMoE enables economical scaling: 236B total parameters with only 21B activated.

DeepSeek-V2 reduces training cost by 42.5% and improves generation throughput by 5.76× vs DeepSeek 67B.

MMLU 78.5 at 21B activated parameters places DeepSeek-V2 on the Pareto frontier among open-source LLMs.

Post-training alignment (SFT + GRPO RL) yields strong multilingual chat performance, especially in Chinese benchmarks.

DeepSeek-V2 demonstrates a practical path to high-performance, cost-efficient open-source LLM scaling.