

SCALE SEAL LEADERBOARDS

Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of 100+ AI models spanning agentic coding, reasoning, safety, and multimodal capabilities — evaluated by Scale Labs



OpenAI



Google



Anthropic



Muse Spark

20+ Benchmarks Spanning Four Capability Domains

Evaluating 100+ models from OpenAI, Google, Anthropic, Meta, and open-source contributors

AGENTIC

Coding & Tool Use

- SWE Atlas (Test Writing)
- SWE Atlas (Codebase QnA)
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Pub/Priv)

FRONTIER

Reasoning & Knowledge

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

SAFETY

Alignment & Honesty

- MASK (Honesty)
- PropensityBench
- Fortress (Nat'l Security)
- Professional Reasoning

MULTIMODAL

Audio & Vision

- AudioMultiChallenge
- VISTA (Vision-Language)
- VisualToolBench

Private datasets prevent overfitting | All scores include confidence intervals

GPT-5.4 Dominates Agentic Coding

Top scores on SWE Atlas and SWE-Bench Pro | Claude Opus 4.6 leads on private repositories


Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

Key Insight: Claude Opus 4.6 leads SWE-Bench Pro (Private) at 47.10, outperforming GPT-5.4 on private repos.

New Entrant: Muse Spark ranks #1 on MCP Atlas tool use (78.30) and top-3 on all SWE-Bench Pro subtasks.

Frontier Reasoning: Scores Remain Below 50%

Even top-performing models are well below saturation on the hardest benchmarks



44.3%

HLE Top Score — GPT-5.4-pro

23.8%

EnigmaEval Top — GPT-5.4-pro

75.5%

MultiChallenge #1 — Muse Spark

25.3%

SciPredict #1 — Gemini-3-pro

Gemini Sweeps Audio and Vision-Language Benchmarks

Google holds top positions across AudioMultiChallenge and VISTA

AudioMultiChallenge

Text Output

Gemini-3-pro-preview



Gemini-2.5-pro



Gemini-2.5-flash



Audio Output

Gemini-3.1-flash-live



GPT-realtime-1.5



VISTA (Vision-Language)

54.65

Gemini 2.5 Pro Exp

53.89

GPT-5.4-pro

VisualToolBench

GPT-5.4 29.17 vs Gemini-3.1-pro 28.97

Safety Alignment: Near-Ceiling Honesty, Critically Low Security

Claude dominates safety benchmarks, but Fortress scores reveal persistent gaps across all models

MASK (Honesty)

NEAR CEILING

Claude Opus 4.6



Claude Sonnet 4



Fortress (National Security)

CRITICALLY LOW

Top Model



Claude Sonnet 4



PropensityBench

10.50

o3 (lower=better)

Prof. Reasoning — Legal

52.27

Claude Opus 4.6 #1

Prof. Reasoning — Finance

53.28

Claude Opus 4.6 #1

Remote Labor Index

4.17

Claude Opus 4.6 (CoWo)

Competitive Landscape: No Single Lab Wins Everywhere

Heatmap of #1-ranked lab/model per benchmark across all four domains

Benchmark	Domain	#1 Lab	#1 Model	Score
SWE Atlas – Test Writing	Agentic	OpenAI	GPT-5.4-xHigh	44.36
SWE Atlas – Codebase QnA	Agentic	OpenAI	GPT 5.4 xHigh	40.80
MCP Atlas (Tool Use)	Agentic	Muse Spark	Muse Spark	78.30
SWE-Bench Pro (Public)	Agentic	OpenAI	GPT-5.4-pro	59.10
SWE-Bench Pro (Private)	Agentic	Anthropic	Claude Opus 4.6	47.10
Humanity's Last Exam	Frontier	OpenAI	GPT-5.4-pro	44.32
EnigmaEval	Frontier	OpenAI	GPT-5.4-pro	23.82
MultiChallenge	Frontier	Muse Spark	Muse Spark	75.52
MultiNRC	Frontier	OpenAI	GPT-5-pro	65.20
SciPredict	Frontier	Google	Gemini-3-pro	25.27
AudioMultiChallenge (Text)	Multimodal	Google	Gemini-3-pro	54.65
AudioMultiChallenge (Audio)	Multimodal	Google	Gemini-3.1-flash	36.06
VISTA	Multimodal	Google	Gemini 2.5 Pro	54.65
VisualToolBench	Multimodal	OpenAI	GPT-5.4	29.17
MASK (Honesty)	Safety	Anthropic	Claude Opus 4.6	96.28
Prof. Reasoning – Legal	Safety	Anthropic	Claude Opus 4.6	52.27
Prof. Reasoning – Finance	Safety	Anthropic	Claude Opus 4.6	53.28
Remote Labor Index	Safety	Anthropic	Claude Opus 4.6	4.17

OpenAI: 7 wins

Google: 4 wins

Anthropic: 5 wins

Muse Spark: 2 wins

Key Takeaways

No Single Model Wins Everything — Specialization Is the New Frontier

1

GPT-5.4 leads agentic coding and frontier reasoning

Dominates SWE Atlas, SWE-Bench Pro (Public), HLE, and EnigmaEval — but does not lead multimodal or safety.

2

Gemini sweeps audio and vision-language

All top-3 positions on AudioMultiChallenge (Text Output) are Gemini variants. Leads VISTA vision-language.

3

Claude leads safety, honesty, and professional reasoning

96%+ on MASK honesty. #1 on both Legal and Finance professional reasoning benchmarks.

4

Benchmarks remain far from saturation

HLE (44%), EnigmaEval (24%), SciPredict (25%) — substantial room for improvement across all labs.

5

Muse Spark emerges as a strong new competitor

Top 3 on 6+ benchmarks including #1 on MultiChallenge (75.5%) and MCP Atlas tool use (78.3).