

SCALE SEAL LEADERBOARDS

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ AI Benchmarks

Comparative analysis of 100+ models from OpenAI, Google, Anthropic, Meta, and new entrants — Agentic, Frontier, Safety, and Multimodal benchmarks.

April 2026

100+ Models Evaluated

OpenAI

Google

Anthropic

Muse Spark

Meta & OSS

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

AGENTIC	FRONTIER	SAFETY	MULTIMODAL
SWE Atlas – Test Writing	Humanity's Last Exam	MASK (Honesty)	AudioMultiChallenge (Text Out)
SWE Atlas – Codebase QnA	HLE (Text Only)	PropensityBench	AudioMultiChallenge (Audio Out)
SWE-Bench Pro (Public)	EnigmaEval	Fortress (Nat'l Security)	VISTA Vision-Language
SWE-Bench Pro (Private)	MultiChallenge	Professional Reasoning – Legal	VisualToolBench
MCP Atlas (Tool Use)	MultiNRC	Professional Reasoning – Finance	
	SciPredict	Remote Labor Index	

100+ Models

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model
SWE Atlas Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex Xhigh
SWE Atlas Codebase QnA	GPT-5.4-xHigh (Codex)	40.80	GPT-5.4-xHigh (Mini-SWE)	36.30	Claude Opus 4.6 (Claude Code)
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)

OpenAI

Anthropic


Google

Muse Spark


GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Top models on key Frontier benchmarks. No model exceeds 50% — substantial room for improvement remains.


Humanity's Last Exam (HLE) — Top 3



EnigmaEval — Top 3 (Max ~24%)




MultiChallenge — Muse Spark Leads at 75.5%



Gemini Models Sweep Audio and Vision-Language Benchmarks



Google's Gemini variants hold top positions across all three multimodal leaderboards.



Gemini Sweep: All 3 top spots on AudioMultiChallenge (Text Output) held by Gemini variants. Gemini 2.5 Pro Experimental leads VISTA at 54.65 vs. GPT-5.4-pro at 53.89 — a razor-thin margin. VisualToolBench: GPT-5.4 edges Gemini-3.1-pro (29.17 vs. 28.97), the one benchmark where OpenAI takes the visual crown.

Claude Models Score 95–96% on MASK Honesty; Fortress National Security Scores Stay Low

Safety alignment varies dramatically by task type — near-ceiling honesty vs. critically low national security scores.



PropensityBench: o3-2025-04-16 achieves lowest (best) score of 10.50 ± 0.60 (lower = better)

Professional Reasoning: Claude Opus 4.6 leads both Legal (52.27) and Finance (53.28)

Contrast: Claude scores ~96% on MASK (honesty) but Fortress scores remain 8–13 — illustrating that safety alignment is task-type specific, not universal.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Which lab's model ranks #1 on each benchmark? Heatmap view — no single lab wins everywhere.

Domain	Benchmark	#1 Lab	#1 Model	Score
AGENTIC	SWE Atlas – Test Writing	OpenAI	GPT-5.4-xHigh (Codex CLI)	44.36
AGENTIC	SWE Atlas – Codebase QnA	OpenAI	GPT-5.4-xHigh (Codex)	40.80
AGENTIC	MCP Atlas (Tool Use)	Muse Spark	Muse Spark	78.30
AGENTIC	SWE-Bench Pro (Public)	OpenAI	GPT-5.4-pro (xHigh)	59.10
AGENTIC	SWE-Bench Pro (Private)	Anthropic	Claude Opus 4.6	47.10
FRONTIER	Humanity's Last Exam	OpenAI	GPT-5.4-pro	44.32
FRONTIER	EnigmaEval	OpenAI	GPT-5.4-pro	23.82
FRONTIER	MultiChallenge	Muse Spark	Muse Spark	75.52
FRONTIER	SciPredict	Google	Gemini-3-pro-preview	25.27
SAFETY	MASK (Honesty)	Anthropic	Claude Opus 4.6	96.28
SAFETY	Professional Reasoning – Legal	Anthropic	Claude Opus 4.6	52.27
SAFETY	Professional Reasoning – Finance	Anthropic	Claude Opus 4.6	53.28
MULTIMODAL	AudioMultiChallenge (Text Output)	Google	Gemini-3-pro-preview	54.65
MULTIMODAL	AudioMultiChallenge (Audio Output)	Google	Gemini-3.1-flash-live	36.06
MULTIMODAL	VISTA Vision-Language	Google	Gemini 2.5 Pro Exp	54.65
MULTIMODAL	Vista Tool Use	OpenAI	GPT-5.4	28.17

KEY TAKEAWAYS

No Single Model Wins Everything — Specialization Is the New Frontier

01

GPT-5.4 leads agentic coding (SWE Atlas, SWE-Bench Pro Public) and frontier reasoning (HLE at 44.3%, EnigmaEval at 23.8%) — but not multimodal.

02

Gemini sweeps audio benchmarks — all top-3 on AudioMultiChallenge (Text Output) are Gemini variants. Leads VISTA vision-language at 54.65.

03

Claude leads safety & honesty: 96.28% on MASK (#1), plus both Professional Reasoning benchmarks (Legal 52.27, Finance 53.28).

04

Frontier benchmarks remain wide open: HLE 44%, EnigmaEval 24%, SciPredict 25% — significant room for improvement across all labs.

05

New entrant Muse Spark ranks top-3 on 6+ benchmarks: #1 MultiChallenge (75.5%), #1 MCP Atlas (78.3%), top-3 in SWE-Bench Pro and HLE.

20+

Benchmarks

100+

Models

4

Domains

6+

Muse Spark Top-3

scale.com/leaderboard