



Scale SEAL Leaderboards

Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated | Agentic | Frontier | Safety | Multimodal

20+ Benchmarks Spanning Four Capability Domains

Evaluate 100+ Models

AGENTIC

GPT-5.4 leads

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)

FRONTIER

GPT-5.4 / Muse Spark

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

SAFETY

Claude dominates

- MASK (Honesty)
- PropensityBench
- Fortress (Nat'l Security)
- Prof. Reasoning – Legal
- Prof. Reasoning – Finance

MULTIMODAL

Gemini sweeps

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA (Vision-Language)
- VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

OpenAI

Anthropic

Muse Spark


Google

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

Key Insight: GPT-5.4 leads 4 of 5 agentic benchmarks, but Claude Opus 4.6 takes #1 on SWE-Bench Pro (Private). Muse Spark appears in top 3 on 3 of 5 benchmarks.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%


Humanity's Last Exam




EnigmaEval



MultiChallenge




SciPredict




Gemini Models Sweep Audio and Vision-Language Benchmarks


AudioMultiCh. (Text)



AudioMultiCh. (Audio)



VISTA (Vision-Language)



VisualToolBench: GPT-5.4 edges Gemini-3.1-pro (29.17 vs 28.97) — the only multimodal benchmark not led by Gemini. All top-3 AudioMultiChallenge (Text Output) positions go to Gemini.

Claude Models Score 95-96% on MASK; Fortress Scores Stay Low

MASK Honesty — Near Ceiling

96.28% Claude Opus 4.6

96.13% Claude Sonnet 4

Fortress (Nat'l Security) — Critically Low

8.24 #1 Model (undisclosed)

12.80 Claude Sonnet 4

PropensityBench (lower = better)

10.50 o3-2025-04-16

Professional Reasoning

52.27 Legal – Claude Opus 4.6

53.28 Finance – Claude Opus 4.6

Contrast: MASK scores near 96% ceiling vs. Fortress at 8–13% floor — safety alignment varies dramatically by task type.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	OpenAI	Google	Anthropic	Muse Spark
• SWE Atlas – Test Writing	#1			
• SWE Atlas – Codebase QnA	#1		#3	
• MCP Atlas (Tool Use)		#3	#2	#1
• SWE-Bench Pro (Public)	#1		#3	#2
• SWE-Bench Pro (Private)	#3		#1	#2
• Humanity's Last Exam	#1	#3		#2
• EnigmaEval	#1	#2		
• MultiChallenge	#3	#2		#1
• SciPredict		#1	#2	
• AudioMultiCh. (Text)		#1		
• AudioMultiCh. (Audio)	#2	#1		
• VISTA	#2	#1		
• MASK (Honesty)			#1	
• Fortress			#2	
• Prof. Reasoning (Legal)			#1	
• Prof. Reasoning (Finance)			#1	#2

Key Takeaways

No Single Model Wins Everything — Specialization Is the New Frontier

- 01** GPT-5.4 leads agentic coding and frontier reasoning, but does not compete in multimodal or safety domains.
- 02** Gemini dominates audio and vision-language benchmarks, sweeping all top-3 on AudioMultiChallenge (Text).
- 03** Claude leads safety (96%+ MASK) and professional reasoning, but Fortress scores remain critically low across all models.
- 04** Key benchmarks unsaturated: HLE (44%), EnigmaEval (24%), SciPredict (25%) — substantial room for improvement.
- 05** New entrant Muse Spark: top 3 on 6+ benchmarks — #1 on MultiChallenge at 75.5%, spanning coding and reasoning.