

SCALE LABS

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ Models | 20+ Benchmarks | 4 Capability Domains

Agentic

Frontier

Safety

Multimodal

Based on public Scale Labs Leaderboard data | April 2026

20+ Benchmarks Spanning Four Capability Domains

Evaluating 100+ frontier models from leading AI labs

Agentic

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)
- MCP Atlas (Tool Use)

Frontier

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

Safety

- MASK (Honesty)
- PropensityBench
- Fortress
- Professional Reasoning – Legal
- Professional Reasoning – Finance

Multimodal

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA Vision-Language
- VisualToolBench

Benchmarks combine private evaluation sets to prevent overfitting with open-source datasets for comparability

GPT-5.4 Dominates Agentic Coding

Top model scores across SWE Atlas and SWE-Bench Pro benchmarks

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh	44.36	GPT-5.4-xHigh (Mini)	40.00	GPT-5.3-Codex	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh	40.80	GPT 5.4 xHigh (Mini)	36.30	Opus 4.6	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro	43.40

Muse Spark is a notable new entrant, appearing in top 3 across MCP Atlas, SWE-Bench Pro (Public & Private), and multiple other benchmarks.


● OpenAI ● Anthropic ● Google ● Muse

Scores shown with $\pm$ confidence intervals. Private benchmarks use held-out repositories to prevent data contamination.


Frontier Reasoning: Scores Remain Below 50%

Top scores on Humanity's Last Exam, EnigmaEval, and MultiChallenge benchmarks

Humanity's Last Exam



MultiChallenge



EnigmaEval

23.8%

GPT-5.4-pro leads; scores under 24%

SciPredict

25.3%

Gemini-3-pro-preview tops; forecasting remains hard


Key Insight

All frontier benchmarks remain far from saturation — significant room for improvement


Gemini Sweeps Audio and Vision-Language

Google's Gemini variants hold top positions across all AudioMultiChallenge and VISTA leaderboards

AudioMultiChallenge (Text Output)



VISTA Vision-Language



AudioMultiChallenge (Audio Output)

Gemini-3.1-flash-live-preview: 36.06

Gemini leads by 1.33 points with tight confidence intervals (± 4.41)

vs GPT-realtime-1.5 at 34.73

VisualToolBench: GPT-5.4 edges Gemini-3.1-pro (29.17 vs 28.97) — the one multimodal area where OpenAI leads

Safety Alignment: High MASK, Low Fortress

Anthropic's Claude models lead honesty benchmarks but safety performance varies dramatically by task type

MASK Honesty Benchmark

Near-ceiling scores for Claude models




± Confidence: $\pm 0.41 - \pm 0.57$

Claude dominates with >96% honesty, indicating strong alignment properties.

Fortress (National Security)

Critically low scores across all models



± Confidence: $\pm 1.93 - \pm 2.36$

National security tasks remain extremely challenging. Scores below 13 indicate substantial room for safety improvements.

PropensityBench: o3-2025-04-16 leads at 10.50 (lower is better)

Professional Reasoning: Claude Opus 4.6 leads Legal (52.27) and Finance (53.28)

Safety benchmarks reveal significant variance: near-perfect honesty (96%) coexists with critical weakness in national security scenarios (8–13%)

No Single Lab Wins Everywhere

Competitive landscape mapping #1-ranked lab across each benchmark domain

Benchmark	OpenAI	Google	Anthropic	Muse
SWE Atlas – Test Writing	●			
SWE Atlas – Codebase QnA	●			
MCP Atlas (Tool Use)				●
SWE-Bench Pro (Public)	●			
SWE-Bench Pro (Private)			●	
Humanity's Last Exam	●			

OpenAI leads coding & reasoning | Google leads multimodal | Anthropic leads safety | Muse emerging

Key Takeaways

Specialization is the defining trend in frontier AI

1

GPT-5.4 Leads Agentic & Reasoning

Dominates SWE Atlas, SWE-Bench Pro, Humanity's Last Exam, and EnigmaEval. OpenAI's coding and reasoning lead is clear.

2

Gemini Dominates Audio/Vision

Sweeps all AudioMultiChallenge benchmarks (Text and Audio output) and VISTA vision-language. Google owns multimodal.

3

Claude Leads Safety & Honesty

96%+ on MASK honesty, #1 on Professional Reasoning (Legal & Finance). But Fortress scores (8–13) reveal critical vulnerability.

4

Muse Spark: Strong New Entrant

Ranks top 3 on 6+ benchmarks including #1 on MultiChallenge (75.5%) and MCP Atlas Tool Use (78.3%). One to watch.

Benchmarks remain unsaturated: HLE (44%), EnigmaEval (24%), SciPredict (25%) — substantial room for improvement across all frontier models