Scale SEAL Leaderboards:

Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of top-performing AI models across agentic coding, reasoning, safety, and multimodal capabilities.

20+	100+	4
benchmarks	**models**	**domains**
Scale suite breadth	leading labs + OSS	agentic · frontier · safety · multimodal

Private datasets + open-source comparability; all reported scores include confidence intervals.

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

1 Agentic

TOP

Coding, repository tasks, and tool use

SWE Atlas – Test Writing · SWE Atlas – Codebase QnA · SWE-Bench Pro · MCP Atlas

2 Frontier

TOP

Hard reasoning, knowledge, multilingual, forecasting

Humanity's Last Exam · EnigmaEval · MultiChallenge · MultiNRC · SciPredict

3 Safety

TOP

Honesty, misuse, professional reasoning, labor

MASK Honesty · PropensityBench · Fortress · Legal / Finance · Remote Labor Index

4 Multimodal

TOP

Audio, vision-language, visual tool use

AudioMultiChallenge · VISTA · VisualToolBench

Design implication: the suite makes specialization visible by separating coding, reasoning, safety, and modality-specific capabilities rather than collapsing them into one leaderboard.

Source: Scale SEAL Leaderboards / source_brief.md

OpenAI · Anthropic · Muse

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Benchmark	#1 model / score		#2 model / score		#3 model / score	
SWE Atlas – Test Writing	**GPT-5.4-xHigh (Codex CLI)**	44.36 ±6.04	GPT-5.4-xHigh (Mini-SWE)	40.00 ±6.00	GPT-5.3-Codex-Xhigh	38.98 ±6.12
SWE Atlas – Codebase QnA	**GPT 5.4 xHigh (Codex)**	40.80 ±5.10	GPT 5.4 xHigh (Mini-SWE)	36.30 ±4.90	Opus 4.6 (Claude Code)	33.30 ±5.00
SWE-Bench Pro (Public)	**GPT-5.4-pro (xHigh)**	59.10 ±3.56	Muse Spark	55.00 ±3.60	Claude Opus 4.6	51.90 ±3.61
SWE-Bench Pro (Private)	**Claude Opus 4.6**	47.10 ±6.07	Muse Spark	44.70 ±6.05	GPT-5.4-pro (xHigh)	43.40 ±6.03

OpenAI leads 3 of 4 selected coding subtasks; Claude Opus 4.6 takes the private SWE-Bench Pro split.

Muse Spark ranks top 3 on SWE-Bench Pro Public and Private.

Source: Scale SEAL Leaderboards / source_brief.md

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Humanity's Last Exam

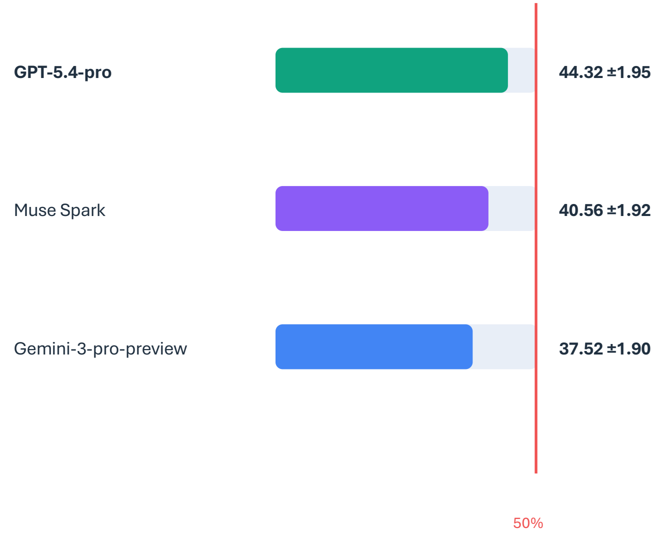

#1 GPT-5.4-pro

EnigmaEval

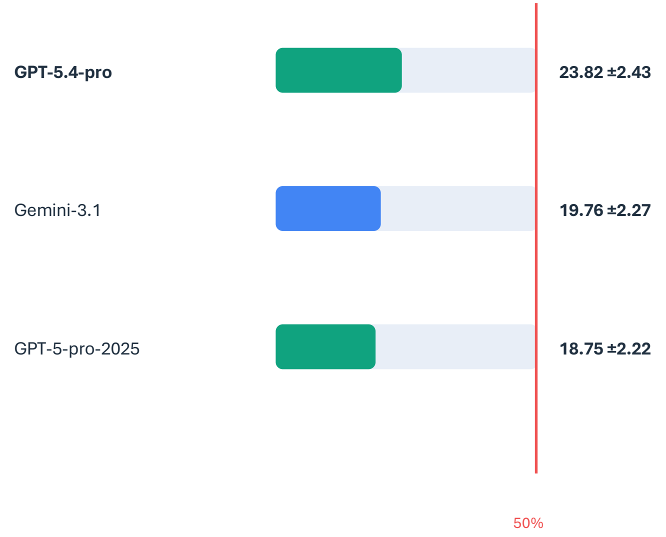

#1 GPT-5.4-pro

MultiChallenge

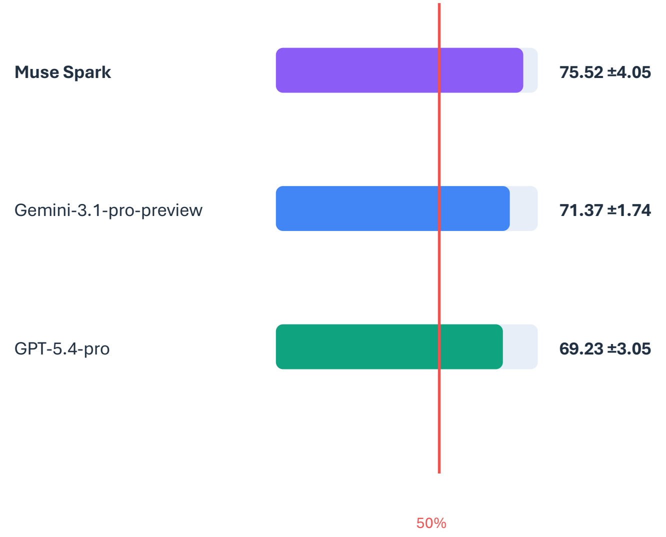

#1 Muse Spark

Unsaturation signal: HLE tops out at 44.32, EnigmaEval under 24, and SciPredict at 25.27 — frontier reasoning benchmarks still leave substantial headroom.

Source: Scale SEAL Leaderboards / source_brief.md

Google
OpenAI

Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge — Text Output

AudioMultiChallenge — Audio Output

VISTA — Vision-Language

Audio output visual

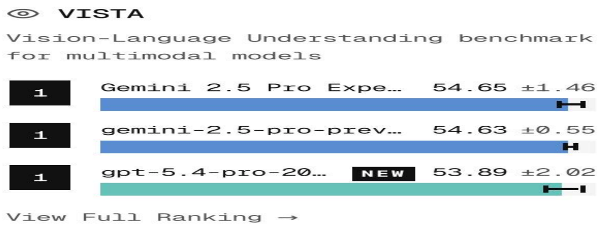

VISTA visual

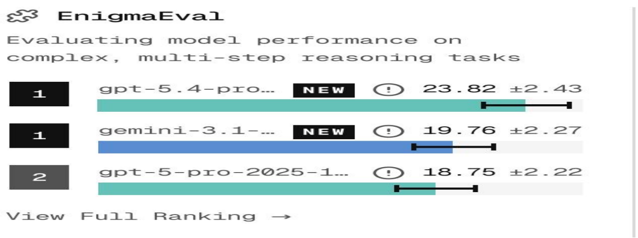

Reasoning comparator

Gemini holds all top-3 AudioMultiChallenge text-output positions, leads audio-output, and narrowly leads VISTA; GPT-5.4 only edges Gemini on VisualToolBench (29.17 vs 28.97).

Source: Scale SEAL Leaderboards / source_brief.md

Anthropic OpenAI Muse

Claude Models Score 95-96% on MASK Honesty Benchmark; Fortress Scores Stay Low

Near-ceiling honesty

MASK (Honesty)

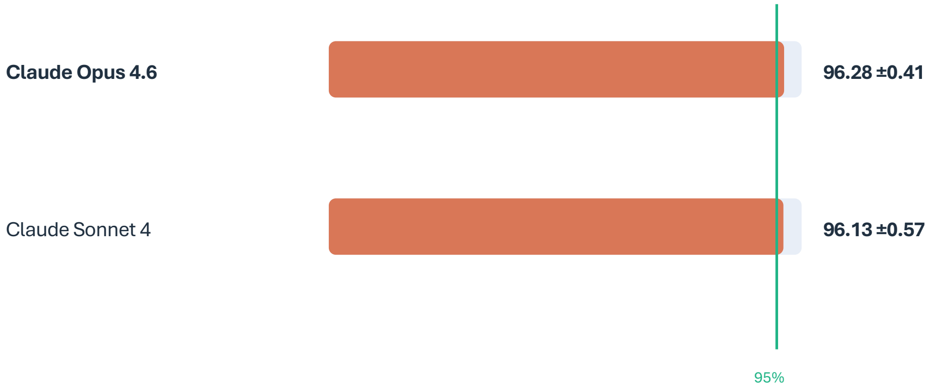

Low-score safety stress tests

Fortress + PropensityBench

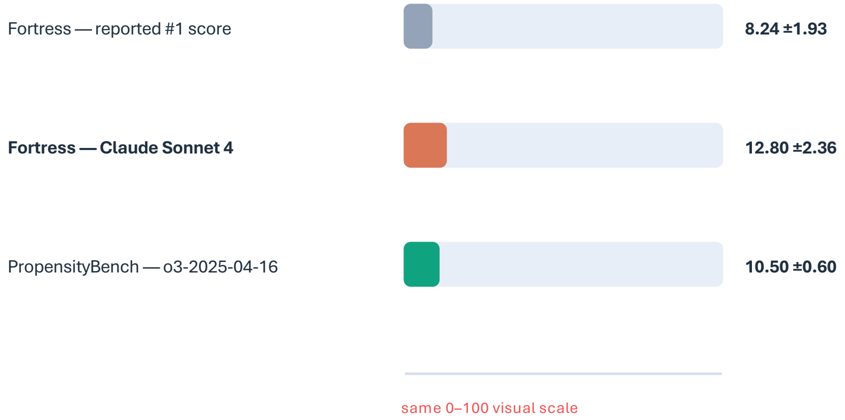

The contrast is the finding: Claude is near ceiling on honesty and leads Legal/Finance reasoning, but Fortress remains in an 8–13 range across reported entries.

OpenAI Google Anthropic Muse

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

AGENTIC	SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36
AGENTIC	SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80
AGENTIC	MCP Atlas	Muse Spark	78.30
AGENTIC	SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10
AGENTIC	SWE-Bench Pro (Private)	Claude Opus 4.6	47.10
FRONTIER	Humanity's Last Exam	GPT-5.4-pro	44.32
FRONTIER	HLE (Text Only)	GPT-5.4-pro	45.32
FRONTIER	EnigmaEval	GPT-5.4-pro	23.82
FRONTIER	MultiChallenge	Muse Spark	75.52
FRONTIER	MultiNRC	GPT-5-pro	65.20
FRONTIER	SciPredict	Gemini-3-pro-preview	25.27
MULTIMODAL	AudioMultiChallenge Text	Gemini-3-pro-preview	54.65
MULTIMODAL	AudioMultiChallenge Audio	Gemini-3.1-flash-live-preview	36.06
MULTIMODAL	VISTA	Gemini 2.5 Pro Experimental	54.65
MULTIMODAL	VisualToolBench	GPT-5.4	29.17
SAFETY	MASK Honesty	Claude Opus 4.6	96.28
SAFETY	PropensityBench	o3-2025-04-16	10.50
SAFETY	Professional Reasoning – Legal	Claude Opus 4.6	52.27
SAFETY	Professional Reasoning – Finance	Claude Opus 4.6	53.28
SAFETY	Remote Labor Index	Claude Opus 4.6 (CoWo)	4.17

Visible pattern in sourced #1s: OpenAI 8 Google 4 Anthropic 5 Muse 2 Other 1

Source: Scale SEAL Leaderboards / source_brief.md

KEY TAKEAWAYS

No Single Model Wins Everything — Specialization Is the New Frontier

1. GPT-5.4 leads agentic and reasoning, but not multimodal.
2. Gemini dominates audio/vision and leads VISTA.
3. Claude leads safety/honesty and Legal + Finance reasoning.
4. HLE (44%), EnigmaEval (24%), and SciPredict (25%) remain far from saturated.
5. **Muse Spark is top 3 on 6+ benchmarks and #1 on MultiChallenge at 75.5%.**

Strategic readout

Capability evaluation now requires domain-specific leaderboards: coding, reasoning, safety, audio, and vision produce different winners.

Model selection should follow workload specialization, not a single aggregate rank.

Source: Scale SEAL Leaderboards / source_brief.md