

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated across Agentic, Frontier, Safety, and Multimodal domains

scale.com/leaderboard

For ML researchers, AI lab leadership, and engineering managers

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

AGENTIC

Coding & Tool Use

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- SWE-Bench Pro (Public/Private)
- MCP Atlas (Tool Use)

FRONTIER

Reasoning & Knowledge

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC (Multilingual)
- SciPredict

SAFETY

Alignment & Honesty


- MASK (Honesty)
- PropensityBench
- Fortress (National Security)
- Professional Reasoning – Legal / Finance
- Remote Labor Index

MULTIMODAL

Audio & Vision-Language


- AudioMultiChallenge (Text/Audio Output)
- VISTA (Vision-Language)
- VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro



Insight: GPT-5.4 leads on 3 of 4 benchmarks, but Claude Opus 4.6 takes #1 on private repos (47.1%). Muse Spark is competitive across all four.



GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%



Benchmarks remain far from saturation. Even top scores on HLE (44%) and EnigmaEval (24%) show substantial room for improvement.

SciPredict: Gemini-3-pro-preview leads at 25.3%, followed by Claude Opus 4.5 (23.1%) and Claude Opus 4.1 (22.2%)

Gemini Models Sweep Audio and Vision-Language Benchmarks





VISTA (Vision-Language)

- #1 Gemini 2.5 Pro Experimental
54.65 $\pm$ 1.46
- #2 GPT-5.4-pro
53.89 $\pm$ 2.02

VisualToolBench

- #1 GPT-5.4
29.17
- #2 Gemini-3.1-pro
28.97

Claude Models Score 95–96% on MASK Honest Benchmark; Fortress Scores Stay Low



Professional Reasoning


- Legal: Claude Opus 4.6 — 52.27 ± 0.66
- Finance: Claude Opus 4.6 — 53.28 ± 0.18
- Remote Labor Index: Claude Opus 4.6 (CoWo) — 4.17

PropensityBench (lower = better)

- #1 o3-2025-04-16 — 10.50 ± 0.60
- #2 Claude Sonnet 4 — (score not reported)

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	#1 Lab	#1 Model	Score
SWE Atlas – Test Writing	OpenAI	GPT-5.4-xHigh (Codex CLI)	44.36
SWE Atlas – Codebase QnA	OpenAI	GPT 5.4 xHigh (Codex)	40.80
MCP Atlas (Tool Use)	Muse Spark	Muse Spark	78.30
SWE-Bench Pro (Public)	OpenAI	GPT-5.4-pro (xHigh)	59.10
SWE-Bench Pro (Private)	Anthropic	Claude Opus 4.6	47.10
Humanity's Last Exam	OpenAI	GPT-5.4-pro	44.32
EnigmaEval	OpenAI	GPT-5.4-pro	23.82
MultiChallenge	Muse Spark	Muse Spark	75.52
MultiNRC (Multilingual)	OpenAI	GPT-5-pro	65.20
SciPredict	Google	Gemini-3-pro-preview	25.27
AudioMultiChallenge (Text)	Google	Gemini-3-pro-preview	54.65
AudioMultiChallenge (Audio)	Google	Gemini-3.1-flash-live	36.06
VISTA (Vision-Language)	Google	Gemini 2.5 Pro Experimental	54.65
VisualToolBench	OpenAI	GPT-5.4	29.17
MASK (Honesty)	Anthropic	Claude Opus 4.6	96.28
PropensityBench	OpenAI	o3-2025-04-16	10.50
Fortress (National Security)	—	Leader (unspecified)	8.24
Professional Reasoning – Legal	Anthropic	Claude Opus 4.6	52.27
Professional Reasoning – Finance	Anthropic	Claude Opus 4.6	53.28



Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

01 GPT-5.4 leads agentic coding and frontier reasoning, but does not dominate multimodal or safety benchmarks.

02 Gemini sweeps audio and vision-language — all top-3 positions on AudioMultiChallenge (Text Output) are Gemini variants.

03 Claude leads safety and honesty with 96%+ on MASK, plus #1 on both professional reasoning benchmarks (Legal and Finance).

04 Key benchmarks remain far from saturation: HLE at 44%, EnigmaEval at 24%, and SciPredict at 25% show substantial room for improvement.

05 New entrant Muse Spark is a strong competitor, ranking top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5%.

The competitive landscape is fragmenting. Model specialization — not general supremacy — is the defining trend of 2025.