

SCALE SEAL LEADERBOARDS

Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of 100+ models from OpenAI, Google, Anthropic, and emerging labs evaluated across agentic coding, reasoning, safety, and multimodal capabilities.

MODELS	BENCHMARKS	DOMAINS	SOURCE
100+	20+	4	scale.com/leaderboard



20+ Benchmarks Spanning Four Capability Domains

Private datasets prevent overfitting; all scores include confidence intervals

Agentic

- SWE Atlas - Test Writing
- SWE Atlas - Codebase QnA
- MCP Atlas (Tool Use)
- SWE Bench Pro (Pub/Priv)

Leader: GPT-5.4

Frontier

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC - SciPredict

Leader: GPT-5.4-pro

Safety


- MASK (Honesty)
- PropensityBench
- Fortress (Nat'l Security)
- Prof. Reasoning (Legal/Fin)

Leader: Claude Opus 4.6

Multimodal

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA (Vision-Language)
- VisualToolBench

Leader: Gemini




GPT-5.4 Dominates Agentic Coding

Top 3 models on SWE Atlas and SWE-Bench Pro subtasks

BENCHMARK	#1 MODEL	SCORE	#2 MODEL	SCORE	#3 MODEL	SCORE
SWE Atlas - Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas - Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40


● OpenAI ● Anthropic ● Muse Spark




GPT-5.4-Pro Leads Humanity's Last Exam at 44.3%

Scores remain well below saturation across all frontier reasoning benchmarks

HUMANITY'S LAST EXAM



ENIGMAEVAL




MULTICHALLENGE



Gemini Sweeps Audio and Vision-Language Benchmarks

Google holds top positions across all three multimodal leaderboards




Google (all top-3 on Audio Text Output) OpenAI (close #2 on VISTA, Audio Out)




Claude Scores 96% on MASK Honesty; Fortress Stays Low

Safety alignment varies dramatically by task type




PropensityBench (lower = better)
o3: 10.50

Prof. Reasoning - Legal
Opus 4.6: 52.27



Prof. Reasoning - Finance
Opus 4.6: 53.28

Remote Labor Index
Opus 4.6: 4.17



OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

No single lab dominates all benchmarks — Muse Spark emerges as a strong new entrant

Benchmark	OpenAI	Google	Anthropic	Muse Spark
SWE Atlas - Test Writing	#1	—	—	—
SWE Atlas - Codebase QnA	#1	—	#3	—
MCP Atlas (Tool Use)	—	#3	#2	#1
SWE-Bench Pro (Public)	#1	—	#3	#2
SWE-Bench Pro (Private)	#3	—	#1	#2
Humanity's Last Exam	#1	#3	—	#2
EnigmaEval	#1	#2	—	—
MultiChallenge	#3	#2	—	#1
SciPredict	—	#1	#2	—
AudioMultiChallenge (Text)	—	#1	—	—
AudioMultiChallenge (Audio)	#2	#1	—	—
VISTA (Vision-Language)	#2	#1	—	—
MASK (Honesty)	—	—	#1	—
Prof. Reasoning (Legal)	—	—	#1	—
Prof. Reasoning (Finance)	—	—	#1	#2



No Single Model Wins Everything — Specialization Is the New Frontier

01 **GPT-5.4 leads agentic coding and frontier reasoning** but does not dominate multimodal or safety benchmarks. Top SWE-Bench Pro (Public) score: 59.10.

02 **Gemini dominates audio and vision-language**, holding all top-3 positions on AudioMultiChallenge (Text Output) and leading VISTA at 54.65.

03 **Claude leads safety and professional reasoning** with 96.28% on MASK honesty and #1 on both Legal (52.27) and Finance (53.28) benchmarks.

04 **Key benchmarks remain far from saturation:** HLE (44%), EnigmaEval (24%), and SciPredict (25%) show substantial room for improvement across all labs.

