

SCALE SEAL LEADERBOARDS

Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across agentic coding, frontier reasoning,
safety alignment, and multimodal understanding

Agentic

Frontier

Safety

Multimodal

20+ Benchmarks Spanning Four Capability Domains

Scale's benchmark suite evaluates 100+ models across coding, reasoning, safety, and perception



Agentio

SWE Atlas, SWE-Bench Pro, MCP Atlas

5 benchmarks



Frontier

HLE, EnigmaEval, MultiChallenge, MultiNRC, SciPredict

6 benchmarks



Safety

MASK, PropensityBench, Fortress, Prof. Reasoning, RLI


6 benchmarks



Multimodal


AudioMultiChallenge, VISTA, VisualToolBench

4 benchmarks




GPT-5.4 Dominates Agentic Coding

Top scores on SWE Atlas and SWE-Bench Pro — Claude Opus 4.6 leads private repos





GPT-5.4-Pro Leads HLE at 44.3% – Scores Remain Below 50%

Frontier reasoning benchmarks remain far from saturation, with top scores under 45%




Gemini Models Sweep Audio and Vision-Language Benchmarks

All top-3 positions on AudioMultiChallenge (Text) held by Gemini variants



Claude Scores 95–96% on MASK; Fortress Scores Stay Below 13

Safety alignment varies dramatically by task type — honesty near ceiling, security far from solved



OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

No single lab wins everywhere — specialization is the defining trend

BENCHMARK	OPENAI	AGENTIC	GOOGLE	ANTHROPIC	MUSE SPARK
SWE Atlas Test Writing	#1		—	—	—
SWE Atlas Codebase QnA	#1		—	—	—
MCP Atlas (Tool Use)	—		—	—	#1
SWE-Bench Pro (Public)	#1		—	—	—
SWE-Bench Pro (Private)	—		—	#1	—
FRONTIER REASONING					
Humanity's Last Exam	#1		—	—	—
EnigmaEval	#1		—	—	—
MultiChallenge	—		—	—	#1
MultiNRC	#1		—	—	—
SciPredict	—		#1	—	—
MULTIMODAL					
AudioMulti (Text)	—		#1	—	—
AudioMulti (Audio)	—		#1	—	—
VISTA	—		#1	—	—
VisualToolBench	#1		—	—	—
SAFETY & ALIGNMENT					
MASK (Honesty)	—		—	#1	—
PropensityBench	#1		—	—	—
Prof. Reasoning (Legal)	—		—	#1	—
Prof. Reasoning (Finance)	—		—	#1	—

Key Takeaways: Specialization Is the New Frontier

No single model wins everywhere — the competitive landscape rewards domain strength

- 01** GPT-5.4 leads agentic coding and frontier reasoning but does not win on multimodal or safety benchmarks — strength is concentrated, not universal.
- 02** Gemini dominates audio and vision-language — all top-3 on AudioMultiChallenge (Text Output) are Gemini variants, and it leads VISTA.
- 03** Claude leads safety and professional reasoning — 96%+ MASK honesty, #1 on Legal and Finance, but Fortress scores remain critically low across all models.
- 04** Key benchmarks remain far from saturation: HLE (44%), EnigmaEval (24%), SciPredict (25%) — substantial room for improvement across all labs.
- 05** Muse Spark emerges as a strong new competitor, ranking top 3 on 6+ benchmarks including #1 on MultiChallenge (75.5%), top 2 on HLE, and top 3 on SWE-Bench Pro.