

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ Benchmarks

20+


Benchmarks

100+

Models Evaluated

4

Capability Domains



20+ Benchmarks Spanning Four Capability Domains



Agentic

SWE Atlas — Test Writing
SWE Atlas — Codebase QnA
MCP Atlas (Tool Use)
SWE-Bench Pro



Frontier

Humanity's Last Exam
EnigmaEval
MultiChallenge
SciPredict




Safety

MASK (Honesty)
Fortress
PropensityBench
Professional Reasoning



Multimodal

AudioMultiChallenge
VISTA (Vision-Language)
VisualToolBench




GPT-5.4 Dominates Agentic Coding

Top scores across SWE Atlas and SWE-Bench Pro benchmarks

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas — Test Writing	GPT-5.4-xHigh	44.36	GPT-5.4-xHigh (Mini)	40.00	GPT-5.3-Codex	38.98
SWE Atlas — Codebase QnA	GPT-5.4-xHigh	40.80	GPT-5.4-xHigh (Mini)	36.30	Claude Opus 4.6	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro	43.40

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3%

Scores remain below 50% — benchmarks still challenge frontier models



Key insight: EnigmaEval (24%), SciPredict (25%) remain far from saturation — substantial room for improvement.

Gemini Models Sweep Audio and Vision Benchmarks

AudioMultiChallenge (Text Output)



AudioMultiChallenge (Audio Output)



VISTA (Vision-Language)



VisualToolBench



Claude Scores 96% on MASK — But Fortress Scores Stay Low

Safety alignment varies dramatically by task type

MASK Honesty

96.28%

Claude Opus 4.6 — Near-ceiling performance

Claude Sonnet 4: 96.13%

Fortress (National Security)

8.24

Claude Opus 4.6 — Critical gap remains

Claude Sonnet 4: 12.80

Additional Safety Metrics

PropensityBench (lower=better)
o3: 10.5 vs Claude Sonnet 4

Professional Reasoning — Legal
Claude Opus 4.6: 52.27%

Professional Reasoning — Finance
Claude Opus 4.6: 53.28% | Muse Spark: 52.44%

Each Lab Claims #1 in Different Domains

Benchmark leader by domain — no single model wins everywhere

Benchmark	OpenAI	Google	Anthropic	Muse
SWE Atlas — Test Writing	GPT-5.4-xHigh 44.36			
MCP Atlas (Tool Use)				Muse Spark 78.30
SWE-Bench Pro (Public)	GPT-5.4-pro 59.10			
SWE-Bench Pro (Private)			Claude Opus 4.6 47.10	
Humanity's Last Exam	GPT-5.4-pro 44.32			
MultiChallenge				Muse Spark 75.52
AudioMultiChallenge		Gemini-3-pro 54.65		
VISTA Vision-Language		Gemini 2.5 Pro 54.65		
MASK (Honesty)			Claude Opus 4.6 96.28	
Professional Reasoning			Claude Opus 4.6 52-53%	

Notable: Muse Spark is a strong new entrant, ranking top 3 on 6+ benchmarks including #1 on MultiChallenge (75.52%) and MCP Atlas (78.30%).

Key Takeaways: Specialization Is the New Frontier

- 1** GPT-5.4 leads **agentic coding** with dominant scores on SWE Atlas and SWE-Bench Pro, but does not lead multimodal benchmarks.
- 2** Gemini **dominates audio/vision** — holding all top-3 positions on AudioMultiChallenge (Text Output) and leading VISTA vision-language understanding.
- 3** Claude **leads safety and honesty** with 96%+ on MASK, but Fortress national security scores remain critically low (8—13) across all models.
- 4** **Key benchmarks remain unsaturated:** HLE (44%), EnigmaEval (24%), SciPredict (25%) — substantial room for improvement on frontier reasoning.
- 5** Muse **Spark** is a strong new entrant, ranking #1 on MultiChallenge (75.5%) and MCP Atlas (78.3%), and top 3 on 6+ benchmarks spanning coding, tool use, and reasoning.