

SCALE LABS · SEAL LEADERBOARDS

Frontier model performance across 20+ AI benchmarks

A comparative readout of 100+ models across Agentic, Frontier, Safety, and Multimodal evaluations.

● OpenAI ● Google ● Anthropic ● Muse Spark

BENCHMARK LANDSCAPE

20+ benchmarks spanning four capability domains evaluate 100+ models

AGENTIC

Coding & tool use

- SWE Atlas Test Writing
- SWE Atlas Codebase QnA
- MCP Atlas
- SWE-Bench Pro

Private + open datasets

FRONTIER

Reasoning & knowledge

- Humanity's Last Exam
- HLE Text Only
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

Room below saturation

SAFETY

Alignment & honesty

- MASK Honesty
- PropensityBench
- Fortress
- Professional Reasoning
- Remote Labor Index

Task-type variance

MULTIMODAL

Audio & vision-language


- AudioMultiChallenge
- Audio Output
- VISTA
- VisualToolBench

Gemini-heavy top ranks

AGENTIC CODING LEADERS

GPT-5.4 dominates SWE Atlas and public SWE-Bench; Claude leads private SWE-Bench

● OpenAI ● Anthropic ● Muse Spark




FRONTIER REASONING

GPT-5.4-Pro leads HLE at 44.3%; hard reasoning remains unsaturated

Humanity's Last Exam



EnigmaEval



MultiChallenge



23.82%

Top EnigmaEval score: GPT-5.4-pro. The brief notes scores under 24% on this benchmark.

25.27%

Top SciPredict score: Gemini-3-pro-preview; scientific forecasting remains under 26%.

MULTIMODAL AUDIO & VISION

Gemini models sweep audio and lead VISTA; GPT-5.4 edges VisualToolBench

AudioMultiChallenge — Text Output



Audio Output



VISTA Vision-Language




VisualToolBench



Provided leaderboard snapshots


Evaluating spoken dialogue systems in multi-turn interaction



Evaluating model performance on complex, multi-step reasoning tasks



Vision-Language Understanding benchmark for multimodal models




Audio Output: Gemini 3.1 Flash Live at 36.06 vs GPT-realtime-1.5 at 34.73. VISTA: Gemini 2.5 Pro Experimental at 54.65 vs GPT-5.4-pro at 53.89.

SAFETY & ALIGNMENT

Claude reaches near-ceiling MASK honesty, while Fortress remains low

Same safety domain, very different regimes



96%

Claude Opus 4.6 scores 96.28 ± 0.41 on MASK Honesty; Claude Sonnet 4 follows at 96.13 ± 0.57 .

8–13

Fortress national-security scores in the brief sit at 8.24 ± 1.93 and 12.80 ± 2.36 , indicating a much harder alignment setting.

COMPETITIVE LANDSCAPE

OpenAI, Google, Anthropic, and Muse Spark each claim #1 territory

● OpenAI ● Google ● Anthropic ● Muse Spark

Agentic SWE Atlas – Test Writing OpenAI - 44.36 ±6.04	Agentic SWE Atlas – Codebase QnA OpenAI - 40.80 ±5.10	Agentic MCP Atlas (Tool Use) Muse Spark - 78.30 ±2.40	Agentic SWE-Bench Pro (Public) OpenAI - 59.10 ±3.56	Agentic SWE-Bench Pro (Private) Anthropic - 47.10 ±6.07
Frontier Humanity's Last Exam OpenAI - 44.32 ±1.95	Frontier HLE (Text Only) OpenAI - 45.32 ±2.10	Frontier EnigmaEval OpenAI - 23.82 ±2.43	Frontier MultiChallenge Muse Spark - 75.52 ±4.05	Frontier MultiNRC OpenAI - 65.20 ±1.24
Frontier SciPredict Google - 25.27 ±1.92	Multimodal AudioMultiChallenge Text Google - 54.65 ±4.57	Multimodal Audio Output Google - 36.06 ±4.41	Multimodal VISTA Google - 54.65 ±1.46	Multimodal VisualToolBench OpenAI - 29.17
Safety MASK Honesty Anthropic - 96.28 ±0.41	Safety PropensityBench OpenAI - 10.50 ±0.60	Safety Professional Reasoning – Legal Anthropic - 52.27 ±0.66	Safety Professional Reasoning – Finance Anthropic - 53.28 ±0.18	Safety Remote Labor Index Anthropic - 4.17

IMPLICATIONS

No single model wins everything — specialization is the new frontier

GPT-5.4 leads agentic/reasoning

OpenAI tops SWE Atlas tasks, SWE-Bench Pro Public, HLE, HLE Text Only, EnigmaEval, MultiNRC, and VisualToolBench.

Gemini dominates audio/vision

Gemini variants hold all top-three AudioMultiChallenge Text positions and lead Audio Output and VISTA.

Claude leads safety and professions

Claude Opus 4.6 is #1 on MASK Honesty, Professional Reasoning Legal and Finance, and RLI.

Hard benchmarks remain open

Top scores: HLE 44%, EnigmaEval 24%, SciPredict 25%; the brief frames these as far from saturation.

Muse Spark is a credible new entrant

Top-three on at least 6 benchmarks, including #1 MultiChallenge at 75.52 $\pm$ 4.05 and #1 MCP Atlas at 78.30 $\pm$ 2.40.

Evaluation strategy shifts

Model selection is increasingly domain-specific: coding, multimodal, safety, and reasoning now point to different leaders.