


Scale SEAL Leaderboards

Frontier Model Performance Across 20+
AI Benchmarks

Comparing 100+ models across Agentic, Frontier, Safety &
Multimodal capabilities

Scale Labs · April 2026



20+ Benchmarks Spanning Four Capability Domains

Evaluate 100+ Models

Agentic

- SWE Atlas — Test Writing, Codebase QnA
- SWE-Bench Pro — Public & Private
- MCP Atlas — Tool Use

Frontier

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict


Safety

- MASK Honesty
- Fortress National Security
- PropensityBench
- Professional Reasoning
- RLI

Multimodal


- AudioMultiChallenge — Text/Audio
- VISTA Vision-Language
- VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro




Key Findings

- GPT-5.4 leads public repo coding with 59.1% on SWE-Bench Pro Public
- Claude Opus 4.6 edges ahead on Private repos (47.1% vs 43.4%)
- Muse Spark: strong new entrant, top-3 on both SWE-Bench Pro variants
- MCP Atlas (Tool Use): Muse Spark #1 at 78.3%, Claude #2 at 75.8%



GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%




44.3%

Best on HLE


- No model exceeds 50% on Humanity's Last Exam
- EnigmaEval: best score is 23.8% — cryptic multi-step puzzles remain hard
- MultiChallenge: Muse Spark leads at 75.5%, above GPT-5.4-pro (69.2%)
- SciPredict: Gemini-3-pro leads at 25.3%

Gemini Models Sweep Audio and Vision-Language Benchmarks


Audio: Text Output



Audio: Audio Output



VISTA: Vision-Language



Gemini variants hold all top-3 positions in AudioMultiChallenge Text Output (54.7 → 40.0%). In VISTA vision-language, Gemini 2.5 Pro edges ahead at 54.65 vs GPT-5.4-pro at 53.89.

Claude Models Score 96%+ on MASK Honesty; Fortress National Security Scores Stay Low


MASK Honesty Benchmark



PropensityBench (lower=better): o3-2025 leads at 10.5 — best refusal rate

Key insight: MASK honesty (96%) vs Fortress security (8%) — alignment performance varies by >80 percentage points across safety task types

Fortress: National Security



OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	Domain	#1 Model	Score	Lab
SWE Atlas Test Writing	Agentic	GPT-5.4-xHigh (Codex)	44.36	OAI
SWE Atlas Codebase QnA	Agentic	GPT-5.4-xHigh (Codex)	40.80	OAI
MCP Atlas (Tool Use)	Agentic	Muse Spark	78.30	MSK
SWE-Bench Pro (Public)	Agentic	GPT-5.4-pro (xHigh)	59.10	OAI
SWE-Bench Pro (Private)	Agentic	Claude Opus 4.6	47.10	ANT
Humanity's Last Exam	Frontier	GPT-5.4-pro	44.32	OAI
EnigmaEval	Frontier	GPT-5.4-pro	23.82	OAI
MultiChallenge	Frontier	Muse Spark	75.52	MSK
MASK Honesty	Safety	Claude Opus 4.6	96.28	ANT
Fortress Security	Safety	(All models low)	8.24	—
Prof. Reasoning (Legal)	Safety	Claude Opus 4.6	52.27	ANT
AudioMultiChallenge	Multimodal	Gemini-3-pro-preview	54.65	GGL
VISTA Vision-Language	Multimodal	Gemini 2.5 Pro Exp.	54.65	GGL



OpenAI



Google



Anthropic



Muse Spark

KEY TAKEAWAY S

Scale SEAL
Leaderboards
2026

1 GPT-5.4 Leads Agentic Coding & Frontier Reasoning

Top scores on SWE-Bench Pro Public (59.1%), Humanity's Last Exam (44.3%), and EnigmaEval (23.8%)

2 Gemini Dominates Audio & Vision-Language

Sweeps AudioMultiChallenge Text Output top-3 (54.7→40.0%) and leads VISTA at 54.65%

3 Claude Leads Safety & Professional Reasoning

96.3% on MASK Honesty, #1 on Legal & Finance reasoning — but Fortress scores remain critically low (8–13%)

4 Muse Spark: New Entrant, Top-3 on 6+ Benchmarks

#1 on MultiChallenge (75.5%) and MCP Atlas (78.3%), top-3 on SWE-Bench Pro and HLE

5 Frontier Benchmarks Remain Unsaturated

HLE best: 44.3% | EnigmaEval: 23.8% | SciPredict: 25.3% — substantial room for improvement