

Scale SEAL Leaderboards

Frontier Model Performance

Across 20+ Benchmarks

100+ models evaluated across agentic, frontier, safety, and multimodal capabilities


Source: scale.com/leaderboard

OpenAI

Google

Anthropic

Muse Spark



20+ Benchmarks Spanning Four Capability Domains

Evaluate 100+ models across agentic coding, frontier reasoning, safety alignment, and multimodal understanding

>{} AGENTIC

- SWE Atlas - Test Writing
- SWE Atlas - Codebase QnA
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)

◆ FRONTIER

- Humanity's Last Exam
- HLE (Text Only)
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

🕒 SAFETY

- MASK (Honesty)
- PropensityBench
- Fortress (NatSec)
- Prof. Reasoning - Legal
- Prof. Reasoning - Finance

◇ MULTIMODAL

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA
- VisualToolBench


All scores in Remote Labor Index. Private datasets prevent overfitting; open-source datasets enable comparability.

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

OpenAI Google Anthropic Muse Spark

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

GPT-5.4 variants hold #1 on 3 of 5 agentic benchmarks. Claude Opus 4.6 leads SWE-Bench Pro (Private) on real-world repos. Muse Spark claims #1 on MCP Atlas tool use at 78.30.



GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Even top scores are well below saturation, indicating these benchmarks still challenge frontier models.

Humanity's Last Exam

GPT-5.4-pro

Muse Spark

Gemini-3-pro-preview

44.3%

Top HLE score (GPT-5.4-pro)
Well below 100% saturation

EnigmaEval

GPT-5.4-pro

Gemini-3.1

GPT-5-pro-2025

23.8%

Top EnigmaEval score (GPT-5.4-pro)
Complex multi-step reasoning unsolved

MultiChallenge

Muse Spark

Gemini-3.1-pro-preview

GPT-5.4-pro

25.3%

Top SciPredict score (Gemini)
Scientific forecasting remains hard

Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge

Text Output




Audio Output



Gemini holds all top-3 positions on AudioMultiChallenge Text Output, demonstrating audio understanding dominance.

VISTA & VisualToolBench

VISTA (Vision-Language)




VisualToolBench



VISTA: Gemini 2.5 Pro edges GPT-5.4-pro by just 0.76 points — the closest race in the multimodal domain. VisualToolBench is equally tight at 0.20 points.

Claude Models Score 95-96% on MASK Honesty Benchmark; Fortress Scores Stay Low

Safety alignment varies dramatically by task type — near-ceiling honesty vs. critically low national security scores.



Remote Labor Index: Claude Opus 4.6 (CoWo) 4.17 | Claude Opus 4.5 3.75

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	Domain	OpenAI	Google	Anthropic	Muse Spark	Score
SWE Atlas – Test Writing	Agentic	#1				44.36
SWE Atlas – Codebase QnA	Agentic	#1				40.80
MCP Atlas (Tool Use)	Agentic				#1	78.30
SWE-Bench Pro (Public)	Agentic	#1				59.10
SWE-Bench Pro (Private)	Agentic			#1		47.10
Humanity's Last Exam	Frontier	#1				44.32
HLE (Text Only)	Frontier	#1				45.32
EnigmaEval	Frontier	#1				23.82
MultiChallenge	Frontier				#1	75.52
MultiNRC	Frontier	#1				65.20
SciPredict	Frontier		#1			25.27
MASK (Honesty)	Safety			#1		96.28
PropensityBench	Safety	#1				10.50
Prof. Reasoning – Legal	Safety			#1		52.27
Prof. Reasoning – Finance	Safety			#1		53.28
AudioMultiCh. (Text)	Multi.		#1			54.65
AudioMultiCh. (Audio)	Multi.		#1			36.06
VISTA	Multi.		#1			54.65
Overall Avg		#1				29.17

OpenAI 9 #1s | Google 4 #1s | Anthropic 4 #1s | Muse Spark 2 #1s

Key Takeaways

No Single Model Wins Everything — Specialization Is the New Frontier

1

GPT-5.4 leads agentic coding and frontier reasoning, but not multimodal

Dominates SWE Atlas, HLE, and EnigmaEval — yet loses audio/vision to Gemini.

2

Gemini dominates audio and vision-language benchmarks

Sweeps all top-3 on AudioMultiChallenge (Text), leads VISTA by a narrow margin.

3

Claude leads safety alignment and professional reasoning

96%+ MASK honesty, #1 on Legal and Finance reasoning — but Fortress remains critically low.

4

Key benchmarks remain far from saturation

HLE at 44%, EnigmaEval at 24%, SciPredict at 25% — substantial room for improvement.

5

Muse Spark emerges as a strong new competitor

Top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — #1 on MultiChallenge at 75.5%.

Source: Scale SEAL Leaderboards — scale.com/leaderboard