

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated • Agentic • Frontier • Safety • Multimodal

20+ Benchmarks Spanning Four Capability Domains

Evaluating 100+ models from leading AI labs

Agentic

- SWE Atlas
- SWE-Bench Pro
- MCP Atlas

Frontier

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge

Safety


- MASK Honesty
- Fortress
- PropensityBench

Multimodal

- AudioMultiChallenge
- VISTA
- VisualToolBench

GPT-5.4 Dominates Agentic Coding

Top Scores on SWE Atlas and SWE-Bench Pro



Key Insight


Claude Opus 4.6 leads SWE-Bench Pro Private at 47.1%

Muse Spark is a strong #2 on both SWE-Bench Pro benchmarks


GPT-5.4-Pro Leads Humanity's Last Exam at 44.3%

Scores Remain Below 50% — Substantial Room for Improvement


Humanity's Last Exam



EnigmaEval



MultiChallenge




Even top scores remain below 50% — all frontier reasoning benchmarks remain unsaturated

Source: Scale Labs SEAL Leaderboard


Gemini Models Sweep Audio and Vision-Language

Gemini variants hold top positions across all three leaderboards

AudioMultiChallenge (Text Output)



VISTA Vision-Language



Audio Output: Gemini-3.1-flash-live (36.06%) leads GPT-realtime-1.5 (34.73%)

Claude Models: Near-Ceiling MASK, Low Fortress

Safety Alignment Varies Dramatically by Task Type

MASK Honesty (Higher = Better)

96%

Claude Opus 4.6
±0.41

Claude Sonnet 4: 96.13%

Near-ceiling performance on honesty
benchmark

Fortress Security (Higher = Better)

8%

Best score
±1.93

All models: 8–13% range — substantial safety
gaps

PropensityBench: o3 scores 10.5 (lower=better)

Claude dominates safety benchmarks but Fortress reveals critical vulnerabilities across all models

OpenAI, Google, and Anthropic Claim #1 Across Domains

Heatmap: #1 Ranked Lab/Model by Benchmark

Domain	Benchmark	#1 Lab	#1 Model
Agentic	SWE Atlas Test Writing	OpenAI	GPT-5.4-xHigh
Agentic	SWE Atlas Codebase QnA	OpenAI	GPT-5.4-xHigh
Agentic	MCP Atlas (Tool Use)	Muse Spark	Muse Spark
Agentic	SWE-Bench Pro (Public)	OpenAI	GPT-5.4-pro
Agentic	SWE-Bench Pro (Private)	Anthropic	Claude Opus 4.6
Frontier	Humanity's Last Exam	OpenAI	GPT-5.4-pro
Frontier	EnigmaEval	OpenAI	GPT-5.4-pro
Frontier	MultiChallenge	Muse Spark	Muse Spark
Multimodal	AudioMultiChallenge	Google	Gemini-3-pro
Multimodal	VISTA	Google	Gemini 2.5 Pro
Labs OpenAI Google Anthropic Muse Spark			
Muse Spark is a notable new entrant — ranks #1 on MultiChallenge (16.0%) and MCP Atlas (16.5%)			

Key Takeaways

Specialization Is the New Frontier

- 1 GPT-5.4 leads agentic coding and frontier reasoning — OpenAI dominates what it was built for
- 2 Gemini sweeps all audio and vision-language benchmarks — Google's multimodal strategy excels
- 3 Claude leads safety and professional reasoning with 96%+ on MASK honesty benchmark
- 4 Key benchmarks (HLE 44%, EnigmaEval 24%, SciPredict 25%) remain far from saturation
- 5 New entrant Muse Spark ranks top 3 on 6+ benchmarks including #1 on MultiChallenge at 75.5%

No single model wins everywhere — specialization defines the frontier landscape