

Scale SEAL Leaderboards

How frontier models stack up across 20+ AI benchmarks

100+ models evaluated

4 capability domains

Private + public datasets

Source: Scale Labs SEAL Leaderboards

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic (coding/tool-use), Frontier (reasoning/knowledge), Safety (alignment/honesty), Multimodal (audio/vision-language)

Leading lab colors

Domain

- SWE Atlas (Test Writing)
- SWE Atlas (Codebase QnA)
- SWE-Bench Pro (Public/Private)
- MCP Atlas (Tool Use)

Leading lab colors

Domain

- Humanity's Last Exam (HLE)
- EnigmaEval
- MultiChallenge
- SciPredict
- MultiNRC (Multilingual)

Leading lab colors

Domain

- MASK (Honesty)
- Fortress (National Security)
- PropensityBench (lower=better)
- Professional Reasoning (Legal/Finance)
- Remote Labor Index (RLI)

Leading lab colors

Domain

- AudioMultiChallenge (Text Output)
- AudioMultiChallenge (Audio Output)
- VISTA (Vision-Language)
- VisualToolBench

Lab color coding:



OpenAI



Google



Anthropic



Muse Spark

Source: Scale Labs SEAL leaderboards (provided brief)

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Top-3 leaderboard snapshots (score $\pm$ CI). Note: Claude leads SWE-Bench Pro (Private).

SWE Atlas — Test Writing

Model	Score ($\pm$ CI)
OpenAI GPT-5.4-xHigh (Codex CLI)	44.36 $\pm$ 6.04
OpenAI GPT-5.4-xHigh (Mini-SWE)	40.00 $\pm$ 6.00

SWE-Bench Pro (Public)

Model	Score ($\pm$ CI)
OpenAI GPT-5.4-pro (xHigh)	59.10 $\pm$ 3.56
Muse Spark	55.00 $\pm$ 3.60
Anthropic Claude Opus 4.6	51.90 $\pm$ 3.61

SWE Atlas — Codebase QnA

Model	Score ($\pm$ CI)
OpenAI GPT 5.4 xHigh (Codex)	40.80 $\pm$ 5.10
OpenAI GPT 5.4 xHigh (Mini-SWE)	36.30 $\pm$ 4.90

SWE-Bench Pro (Private)

Model	Score ($\pm$ CI)
Anthropic Claude Opus 4.6	47.10 $\pm$ 6.07
Muse Spark	44.70 $\pm$ 6.05
OpenAI GPT-5.4-pro (xHigh)	43.40 $\pm$ 6.03

Interpretation: OpenAI leads both SWE Atlas leaderboards and SWE-Bench Pro (Public), while Anthropic's Claude Opus 4.6 takes #1 on SWE-Bench Pro (Private) — suggesting stronger performance on private repos vs public coding tasks.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Frontier benchmarks remain unsaturated (top HLE < 50%; EnigmaEval < 24%).


Humanity's Last Exam (HLE)

Top-3 scores (%)




EnigmaEval

Top-3 scores (%)



MultiChallenge

Top-3 scores (%)




Takeaway: Even leaderboard-topping models show large headroom on hard reasoning tasks (e.g., EnigmaEval

Source: Frontier Reasoning & Knowledge Benchmarks table (provided brief)

Gemini Models Sweep Audio and Vision-Language Benchmarks

Google holds #1 positions across audio and vision-language; OpenAI remains close on VISTA and VisualToolBench.

AudioMultiChallenge (Text Output)




AudioMultiChallenge (Audio Output)



AudioMultiChallenge - Audio Output

Evaluating spoken dialogue systems in multi-turn interaction




[View Full Ranking →](#)

Figure: AudioMultiChallenge (Audio Output)

VISTA

Vision-Language Understanding benchmark for multimodal models



[View Full Ranking →](#)

VISTA: Gemini 2.5 Pro Exp 54.65 ±1.46 vs GPT-5.4-pro 53.89 ±2.02 • VisualToolBench: GPT-5.4 29.17 vs Gemini-3.1-pro 28.97


Source: Multimodal benchmarks + figures (provided brief)

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low

Contrasting visualization: near-ceiling honesty vs low national-security robustness illustrates task-specific variance in alignment performance.


MASK (Honesty)

Higher is better



Fortress (National Security)

Scores remain very low (8–13 range)



OpenAI

PropensityBench (lower=better): o3-2025-04-16 = 10.50 ± 0.60

Source: Safety & Alignment Benchmarks table (provided brief)

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Heatmap view: which lab/model tops each leaderboard (per provided brief).

Domain	Benchmark	#1 lab / model	Score
Agentic	SWE Atlas — Test Writing	OpenAI — GPT-5.4-xHigh (Codex CLI)	44.36 ±6.04
Agentic	SWE Atlas — Codebase QnA	OpenAI — GPT 5.4 xHigh (Codex)	40.80 ±5.10
Agentic	MCP Atlas (Tool Use)	Muse Spark — Muse Spark	78.30 ±2.40
Agentic	SWE-Bench Pro (Public)	OpenAI — GPT-5.4-pro (xHigh)	59.10 ±3.56
Agentic	SWE-Bench Pro (Private)	Anthropic — Claude Opus 4.6	47.10 ±6.07
Frontier	Humanity's Last Exam	OpenAI — GPT-5.4-pro	44.32 ±1.95
Frontier	HLE (Text Only)	OpenAI — GPT-5.4-pro	45.32 ±2.10
Frontier	EnigmaEval	OpenAI — GPT-5.4-pro	23.82 ±2.43
Frontier	MultiChallenge	Muse Spark — Muse Spark	75.52 ±4.05
Frontier	MultiNRC (Multilingual)	OpenAI — GPT-5-pro	65.20 ±1.24
Frontier	SciPredict	Google — Gemini-3-pro-preview	25.27 ±1.92
Multimodal	AudioMultiChallenge (Text Output)	Google — Gemini-3-pro-preview	54.65 ±4.57
Multimodal	AudioMultiChallenge (Audio Output)	Google — Gemini-3.1-flash-live-preview	36.06 ±4.41
Multimodal	VISTA (Vision-Language)	Google — Gemini 2.5 Pro Experimental	54.65 ±1.46
Multimodal	VisualToolBench	OpenAI — GPT-5.4	29.17
Safety	PropensityBench (lower=better)	Anthropic — Claude Opus 4.6	96.28 ±0.4
Safety		OpenAI — o3-2025-04-16	8.24 ±1.93

Competitive landscape: OpenAI leads coding + core reasoning, Google leads audio/vision, Anthropic leads safety + professional reasoning, Muse Spark breaks into #1 on tool use + MultiChallenge.

Source: All #1 entries extracted from provided brief tables + notes

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

Implications for evaluation: choose benchmarks by product surface, not a single aggregate score.

1) OpenAI leads agentic + core reasoning

- SWE Atlas #1 (Test Writing 44.36, Codebase QnA 40.80)
- HLE #1 at 44.32% (and HLE Text Only 45.32%)
- EnigmaEval #1 at 23.82%

2) Google dominates multimodal audio/vision

- AudioMultiChallenge (Text Output) top-3 all Gemini variants
- Audio Output #1: Gemini-3.1-flash-live-preview 36.06
- VISTA #1: Gemini 2.5 Pro Experimental 54.65

3) Anthropic leads safety + professional reasoning

- MASK honesty near ceiling: Claude Opus 4.6 = 96.28%
- Professional Reasoning #1 (Legal 52.27; Finance 53.28)
- Fortress scores remain very low (8–13)

4) Benchmarks remain far from saturation

- HLE top score: 44%
- EnigmaEval top score: 24%
- SciPredict top score: 25%
- Frontier capability headroom is still large