


100+
models

20+
benchmarks

Agentic · Frontier · Safety ·
Multimodal

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of top-performing AI models across Scale Labs' benchmark suite spanning agentic coding, reasoning, safety, and multimodal capabilities.



20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)

5

Frontier

- Humanity's Last Exam
- HLE (Text Only)
- EnigmaEval
- MultiChallenge
- MultiNRC (Multilingual)
- SciPredict

6

Safety

- MASK (Honesty)
- PropensityBench
- Fortress (National Security)
- Professional Reasoning – Legal
- Professional Reasoning – Finance
- Remote Labor Index

6

Multimodal

- AudioMultiChallenge (Text Output)
- AudioMultiChallenge (Audio Output)
- VISTA (Vision-Language)
- VisualToolBench

4

Private datasets reduce overfitting risk; open-source datasets preserve comparability. Reported scores include confidence intervals.

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

SWE Atlas – Test Writing



SWE Atlas – Codebase QnA



SWE-Bench Pro (Public)



SWE-Bench Pro (Private)



Pattern: GPT-5.4 variants hold #1 on both SWE Atlas subtasks and SWE-Bench Pro Public; Claude Opus 4.6 leads Private repositories.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Humanity's Last Exam

Top 3 point estimates

#1 GPT-5.4-pro



#2 Muse Spark



#3 Gemini-3-pro-preview




#1 remains under 45%


EnigmaEval

Top 3 point estimates

#1 GPT-5.4-pro



#2 Gemini-3.1



#3 GPT-5-pro-2025



#1 remains under 24%

MultiChallenge

Top 3 point estimates

#1 Muse Spark



#2 Gemini-3.1-pro-preview



#3 GPT-5.4-pro



#1 reaches 75.5%

HLE, EnigmaEval, and SciPredict remain far from saturation, showing substantial room for improvement on hard reasoning and forecasting.

Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge Text Output

#1 Gemini-3-pro-preview



#2 Gemini-2.5-pro



#3 Gemini-2.5-flash



AudioMultiChallenge Audio Output

#1 Gemini-3.1-flash-live-preview



#2 GPT-realtime-1.5



VISTA Vision-Language

#1 Gemini 2.5 Pro Experimental



#2 GPT-5.4-pro



AudioMultiChallenge - Audio
Output
Evaluating spoken dialogue systems in
multiturn interaction

Rank	Model	Score
#1	Gemini-3.1-flash-live...	36.06 ± 4.41
#2	GPT-realtime-1.5	34.73 ± 4.36
#3	Gemini-2.5-flash...	40.04 ± 4.00

View Full Ranking →

Source figure: Audio output leaderboard shows Gemini 3.1 Flash Live at 36.06 vs GPT-Realtime-1.5 at 34.73.

VISTA
Vision-Language Understanding benchmark
for multimodal models

Rank	Model	Score
#1	Gemini 2.5 Pro Expe...	54.65 ± 1.46
#2	Gemini-2.5-pro-prev...	54.63 ± 0.55
#3	GPT-5.4-pro-28...	53.89 ± 0.82

View Full Ranking →

Source figure: VISTA benchmark shows Gemini 2.5 Pro at 54.65 and GPT-5.4-pro close behind at 53.89.

Pattern: Gemini variants hold every reported top position in AudioMultiChallenge text, AudioMultiChallenge audio, and VISTA.

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low

Near-ceiling honesty

Claude Opus 4.6



Claude Sonnet 4




MASK (Honesty)


Claude models occupy #1–#2 and score above 96%.

Low-score stress tests


Fortress #1 reported



Claude Sonnet 4




PropensityBench o3



Fortress (National Security) remains 8–13; PropensityBench is reported as lower=better.

Implication: alignment performance is task-dependent — Claude dominates honesty and professional reasoning, while national-security stress tests remain difficult.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains



#1-ranked lab/model by benchmark from source brief. Fortress #1 model is not named in the brief and is marked Other.

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

01

GPT-5.4 leads agentic and reasoning — but not multimodal.

OpenAI holds #1 on SWE Atlas, SWE-Bench Pro Public, HLE, HLE Text, EnigmaEval, MultiNRC, VisualToolBench, and PropensityBench.

02

Gemini dominates audio and vision-language.

Gemini variants lead AudioMultiChallenge text, AudioMultiChallenge audio, VISTA, and SciPredict.

03

Claude leads safety, honesty, and professional reasoning.

Claude Opus 4.6 scores 96.28 on MASK and ranks #1 on Legal, Finance, RLI, and SWE-Bench Pro Private.

04

Hard benchmarks are still unsaturated.

HLE tops out at 44.32, EnigmaEval at 23.82, and SciPredict at 25.27.

05

Muse Spark is a serious new entrant.

Top-three on 6+ benchmarks and #1 on MultiChallenge at 75.52 and MCP Atlas at 78.30.