

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of top-performing AI models across Scale Labs' benchmark suite spanning agentic coding, reasoning, safety, and multimodal capabilities.

Scale Labs | 100+ Models Evaluated

20+ Benchmarks Spanning Four Capability Domains

Evaluate 100+ Models

Agentic

Coding & Tool Use

- SWE Atlas
- MCP Atlas
- SWE-Bench Pro

Frontier

Reasoning & Knowledge

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- SciPredict

Safety

Alignment & Honesty

- MASK
- PropensityBench
- Fortress
- Professional Reasoning


Multimodal

Audio & Vision-Language

- AudioMultiChallenge
- VISTA
- VisualToolBench


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

GPT-5.4 variants lead most SWE tasks, while Claude Opus 4.6 strongly competes on private repos. Muse Spark shows strong utility in tool use.



GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%


Reasoning limits are tested here. Even top scores on HLE and EnigmaEval remain well below saturation, highlighting significant room for improvement.




Gemini Models Sweep Audio and Vision-Language Benchmarks

Google's Gemini models dominate multimodal understanding, securing top positions across both Audio and Vision-Language benchmarks.

Audio Output Benchmark





VISTA Vision-Language Benchmark



Claude Models Score ~96% on MASK Honesty; Fortress Scores Stay Low

Anthropic's Claude models dominate safety and professional reasoning benchmarks, yet specific national security evaluations reveal significant vulnerabilities across the board.



PropensityBench (Lower=Better): o3-2025 leads with 10.5 | Professional Reasoning: Claude Opus 4.6 leads Legal (52.3) and Finance (53.3)

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Domain	Benchmark	Leading Lab	#1 Model	Score
Agentic	SWE Atlas (Test Writing)	OpenAI	GPT-5.4-xHigh	44.4
Agentic	SWE-Bench Pro (Public)	OpenAI	GPT-5.4-pro	59.1
Agentic	SWE-Bench Pro (Private)	Anthropic	Claude Opus 4.6	47.1
Agentic	MCP Atlas (Tool Use)	Muse	Muse Spark	78.3
Frontier	Humanity's Last Exam	OpenAI	GPT-5.4-pro	44.3
Frontier	MultiChallenge	Muse	Muse Spark	75.5
Safety	MASK (Honesty)	Anthropic	Claude Opus 4.6	96.3
Multimodal	Audio Output	Google	Gemini 3.1 Flash Live	36.1
Multimodal	VISTA (Vision-Lang)	Google	Gemini 2.5 Pro	54.7

Muse Spark enters as a prominent competitor, capturing #1 in Tool Use (78.3%) and MultiChallenge (75.5%).

Key Takeaways: No Single Model Wins Everything

GPT-5.4 Leads Agentic & Reasoning, but Not Multimodal

OpenAI dominates SWE Atlas, SWE-Bench, and Humanity's Last Exam, but falls behind Gemini in audio/vision.

Gemini Dominates Audio & Vision

Google's models sweep all top-3 positions on AudioMultiChallenge and VISTA benchmarks.

Claude Leads Safety & Professional Reasoning

Anthropic's models reach near-ceiling (96%) on MASK honesty and take #1 in Legal and Finance reasoning.

Benchmarks Remain Far from Saturation

Crucial evaluations like HLE (44%), EnigmaEval (24%), and SciPredict (25%) show substantial room for improvement.

Muse Spark is a Formidable New Entrant

Ranking top 3 on 6+ benchmarks, Muse Spark captures #1 on MultiChallenge (75.5%) and MCP Atlas (78.3%).