

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across agentic coding, frontier reasoning, safety alignment, and multimodal understanding

Scale Labs | scale.com/leaderboard

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

Coding & Tool Use

SWE Atlas, SWE-Bench Pro, MCP Atlas

Frontier

Reasoning & Knowledge

Humanity's Last Exam, EnigmaEval, MultiChallenge

Safety

Alignment & Honesty


MASK, Fortress, PropensityBench

Multimodal

Audio & Vision-Language

AudioMultiChallenge, VISTA, VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro




Key Insights

GPT-5.4 variants lead on 3 of 4 coding benchmarks, including SWE-Bench Pro Public at 59.1%.

Claude Opus 4.6 tops private-repo evaluation (SWE-Bench Pro Private: 47.1%) and MCP Atlas.

Muse Spark — a new entrant — ranks #1 on MCP Atlas (78.3%) and #2 on SWE-Bench Pro Public.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%



Benchmark Saturation

44%

Humanity's Last Exam
— top score well below ceiling


24%

EnigmaEval — complex multi-step reasoning remains unsolved

25%

SciPredict — scientific forecasting still challenging

Gemini Models Sweep Audio and Vision-Language Benchmarks



Vision-Language Leaderboard

VISTA

Gemini 2.5 Pro Experimental

54.6

GPT-5.4-pro

53.89

VisualToolBench

GPT-5.4

29.17

Gemini-3.1-pro

28.97

Gemini holds all top-3 positions on AudioMultiChallenge (Text Output) and leads VISTA vision-language understanding.

Source: scale.com/leaderboard

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low

MASK (Honesty)

Near-ceiling performance on honesty detection

96.28

%

Claude Opus 4.6

96.13

%

Claude Sonnet 4



~96% of ceiling

Fortress (National Security)

Critically low scores across all models

8.24

#1 (unidentified model)

12.8

0

Claude Sonnet 4



~8–13 range — far from solved

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	#1 Lab / Model	Domain
SWE Atlas – Test Writing	OpenAI GPT-5.4-xHigh	Agentic
SWE Atlas – Codebase QnA	OpenAI GPT-5.4-xHigh	Agentic
SWE-Bench Pro (Public)	OpenAI GPT-5.4-pro	Agentic
SWE-Bench Pro (Private)	Anthropic Claude Opus 4.6	Agentic
MCP Atlas	Muse Spark	Agentic
Humanity's Last Exam	OpenAI GPT-5.4-pro	Frontier
EnigmaEval	OpenAI GPT-5.4-pro	Frontier
MultiChallenge	Muse Spark	Frontier
SciPredict	Google Gemini-3-pro	Frontier
AudioMultiChallenge	Google Gemini-3.1-flash-live	Multimodal
VISTA	Google Gemini 2.5 Pro	Multimodal
MASK (Honesty)	Anthropic Claude Opus 4.6	Safety
Professional Reasoning – Legal	Anthropic Claude Opus 4.6	Safety
Professional Reasoning – Finance	Anthropic Claude Opus 4.6	Safety

OpenAI
Google
Anthropic
Muse Spark

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

- 1 GPT-5.4 leads agentic coding and frontier reasoning**
Dominant on SWE Atlas, SWE-Bench Pro, Humanity's Last Exam, and EnigmaEval — but not multimodal.
- 2 Gemini dominates audio and vision-language**
Sweeps AudioMultiChallenge and leads VISTA, with all top-3 text-output slots held by Gemini variants.
- 3 Claude leads safety, honesty, and professional reasoning**
Near-ceiling MASK scores (~96%) and #1 on Legal and Finance reasoning — yet Fortress remains critically low.
- 4 Benchmarks remain far from saturation**
HLE at 44%, EnigmaEval at 24%, SciPredict at 25% — frontier models still have substantial room to improve.
- 5 Muse Spark emerges as a strong new competitor**
Top 3 on 6+ benchmarks including #1 MultiChallenge (75.5%), plus strong showings on MCP Atlas and SWE-Bench Pro.