

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated | 4 capability domains | Private + public datasets

Agentic Coding · Frontier Reasoning · Safety & Alignment · Multimodal

Source: scale.com/leaderboard | April 2026

 OpenAI

 Google

 Anthropic

 Muse Spark

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Private datasets prevent overfitting; open-source datasets ensure comparability. All scores include confidence intervals.

AGENTIC

SWE Atlas – Test Writing
SWE Atlas – Codebase QnA
MCP Atlas (Tool Use)
SWE-Bench Pro (Public)
SWE-Bench Pro (Private)

FRONTIER

Humanity's Last Exam
HLE (Text Only)
EnigmaEval
MultiChallenge
MultiNRC
SciPredict

SAFETY

MASK (Honesty)
PropensityBench
Fortress (Natl Security)
Prof. Reasoning – Legal
Prof. Reasoning – Finance
Remote Labor Index

MULTIMODAL

AudioMultiChallenge (Text)
AudioMultiChallenge (Audio)
VISTA (Vision-Language)
VisualToolBench

Suite maintained by Scale Labs — combining private & open-source evaluation datasets

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas & SWE-Bench Pro

Muse Spark emerges top-3 on 3/5 agentic benchmarks; Claude Opus 4.6 leads private repo evaluation

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

Note: Claude Opus 4.6 is the only model to lead a private-repo SWE-Bench evaluation, suggesting strong generalization.


GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Top scores on HLE (44%), EnigmaEval (24%), and SciPredict (25%) show frontier models still have vast room to improve

Humanity's Last Exam



EnigmaEval




MultiChallenge



Gemini Models Sweep Audio and Vision-Language Benchmarks

Google holds all top positions on AudioMultiChallenge (Text Output) and leads VISTA vision-language understanding



VisualToolBench: GPT-5.4 edges Gemini-3.1-pro (29.17 vs 28.97) — the tightest margin on any multimodal benchmark

Claude Models Score 96% on MASK Honesty; Fortress Scores Stay Below 13%

Safety alignment varies dramatically by task type — near-ceiling honesty vs. critically low national security scores

MASK (Honesty) — Near Ceiling

Claude Opus 4.6



Claude Sonnet 4



PropensityBench (lower = better)

o3-2025-04-16: 10.50 ±0.60 | Claude Sonnet 4: close follow

Remote Labor Index (RLI)

Claude Opus 4.6 (CoWo): 4.17 | Claude Opus 4.5: 3.75
Key insight: 96% honesty vs. 8–13% national security scores — safety is not a monolith

Fortress (National Security) — Critically Low

Top model



Claude Sonnet 4



Professional Reasoning (#1: Claude Opus 4.6)

Legal: 52.27 | Finance: 53.28 (Muse Spark #2 at 52.44)

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Heatmap: #1-ranked lab per benchmark — no single lab dominates all domains. Muse Spark emerges as notable new entrant.

AGENTIC	FRONTIER	SAFETY	MULTIMODAL
SWE Atlas – Test Writing GPT-5.4-xHigh	Humanity's Last Exam GPT-5.4-pro	MASK (Honesty) Claude Opus 4.6	AudioMulti (Text) Gemini-3-pro
SWE Atlas – Codebase QnA GPT 5.4 xHigh	HLE (Text Only) GPT-5.4-pro	PropensityBench o3-2025-04-16	AudioMulti (Audio) Gemini-3.1-flash
MCP Atlas (Tool Use) Muse Spark	EnigmaEval GPT-5.4-pro	Prof. Reasoning – Legal Claude Opus 4.6	VISTA Gemini 2.5 Pro
SWE-Bench Pro (Public) GPT-5.4-pro	MultiChallenge Muse Spark	Prof. Reasoning – Finance Claude Opus 4.6	VisualToolBench GPT-5.4
SWE-Bench Pro (Private) Claude Opus 4.6	MultiNRC GPT-5-pro	Remote Labor Index Claude Opus 4.6	
	SciPredict Gemini-3-pro		

Lab Color Key: OpenAI Google Anthropic Muse Spark

OpenAI #1 on 10 benchmarks | Google #1 on 4 | Anthropic #1 on 5 | Muse Spark #1 on 2

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

1

GPT-5.4 leads agentic coding and frontier reasoning

Dominant on SWE Atlas, SWE-Bench Pro (Public), HLE (44.3%), and EnigmaEval — but does not lead multimodal or safety.

2

Gemini sweeps audio and vision-language benchmarks

All top-3 on AudioMultiChallenge (Text Output) are Gemini variants. Leads VISTA at 54.65.

3

Claude leads safety, honesty, and professional reasoning

96%+ MASK honesty scores. #1 on Legal (52.27) and Finance (53.28) professional reasoning.

4

Benchmarks remain far from saturation

HLE: 44% | EnigmaEval: 24% | SciPredict: 25% — significant room for improvement across all labs.

5

Muse Spark is a strong new competitor

Top 3 on 6+ benchmarks spanning coding, tool use, and reasoning. #1 on MultiChallenge at 75.5%.