

Scale SEAL Leaderboards:

Frontier Model Performance Across 20+ AI Benchmarks

A comparative analysis of 100+ models from OpenAI, Google, Anthropic, Meta, and new entrants across Agentic, Frontier, Safety, and Multimodal capability domains.

Agentic

Frontier

Safety

Multimodal

April 2026 · Benchmark Analysis


20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Scale SEAL maintains private datasets to prevent overfitting, combined with open-source datasets for comparability.

Agentic		Frontier	Safety	100+ Models Evaluated	
SWE Atlas – Test Writing		Humanity's Last Exam	MASK (Honesty)	AudioMultiChallenge (Text)	
SWE Atlas – Codebase QnA		HLE (Text Only)	PropensityBench	AudioMultiChallenge (Audio)	
SWE-Bench Pro (Public)		EnigmaEval	Fortress (Nat. Security)	VISTA (Vision-Language)	
SWE-Bench Pro (Private)		MultiChallenge	Prof. Reasoning – Legal	VisualToolBench	
MCP Atlas (Tool Use)		MultiNRC (Multilingual)	Prof. Reasoning – Finance		
		SciPredict	Remote Labor Index		


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Scores shown are % task completion with 95% confidence intervals. MCP Atlas tests tool-use capability.



GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Even top-performing models score well below saturation, indicating these benchmarks remain challenging for frontier models.



No benchmark exceeds 76% at the top — EnigmaEval (24%) and SciPredict (25%) remain far from solved.


Gemini Models Sweep Audio and Vision-Language Benchmarks

Gemini variants hold top-3 on AudioMultiChallenge (text output) and lead VISTA vision-language understanding.


AudioMultiChallenge
(Text Output)



AudioMultiChallenge
(Audio Output)



VISTA
(Vision-Language)



Key insight: VisualToolBench is the only multimodal benchmark where GPT-5.4 edges Gemini-3.1-pro (29.17 vs 28.97).

Claude Models Score 95–96% on MASK Honesty; Fortress Scores Stay Low (8–13)

Safety alignment varies dramatically by task type — near-ceiling honesty does not imply national security compliance.

MASK (Honesty Benchmark) — Higher is Better



Claude Opus 4.6



Claude Sonnet 4

Near-ceiling performance: 96%+ honesty alignment

PropensityBench (lower = better): o3-2025-04-16 leads at 10.50 ±0.60

Fortress (National Security) — Higher is Better



(Best unnamed model)



Claude Sonnet 4

Critically low: 8–13% on national security tasks

Contrast: MASK honesty (96%) vs Fortress national security (8–13%) — same models, 80+ point gap in safety posture.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Heatmap: which lab's model ranks #1 on each benchmark. Color = leading lab/model. Muse Spark highlighted as notable new entrant.

Benchmark	Domain	Leading Lab	Top Model	Score
SWE Atlas – Test Writing	Agentic	OpenAI	GPT-5.4-xHigh (Codex CLI)	44.36
SWE Atlas – Codebase QnA	Agentic	OpenAI	GPT-5.4 xHigh (Codex)	40.80
MCP Atlas (Tool Use)	Agentic	Muse Spark	Muse Spark	78.30
SWE-Bench Pro (Public)	Agentic	OpenAI	GPT-5.4-pro (xHigh)	59.10
SWE-Bench Pro (Private)	Agentic	Anthropic	Claude Opus 4.6	47.10
Humanity's Last Exam	Frontier	OpenAI	GPT-5.4-pro	44.32
HLE (Text Only)	Frontier	OpenAI	GPT-5.4-pro	45.32
EnigmaEval	Frontier	OpenAI	GPT-5.4-pro	23.82
MultiChallenge	Frontier	Muse Spark	Muse Spark	75.52
MultiNRC (Multilingual)	Frontier	OpenAI	GPT-5-pro	65.20
SciPredict	Frontier	Google	Gemini-3-pro-preview	25.27
AudioMultiChallenge (Text Out)	Multimodal	Google	Gemini-3-pro-preview	54.65
AudioMultiChallenge (Audio Out)	Multimodal	Google	Gemini-3.1-flash-live	36.06
VISTA (Vision-Language)	Multimodal	Google	Gemini 2.5 Pro Experimental	54.65
VisualToolBench	Multimodal	OpenAI	GPT-5.4	29.17
MASK (Honesty)	Safety	Anthropic	Claude Opus 4.6	96.28
PropensityBench	Safety	OpenAI	o3-2025-04-16	10.50*
Fortress (Nat. Security)	Safety	Anthropic	Claude Sonnet 4	12.80
Prof. Reasoning – Legal	Safety	Anthropic	Claude Opus 4.6	52.27
Prof. Reasoning – Finance	Safety	Anthropic	Claude Opus 4.6	53.28

* PropensityBench: lower is better



OpenAI



Google



Anthropic



Muse Spark

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

Summary of competitive dynamics across Scale SEAL's 20+ benchmark suite.

GPT-5.4 Leads Agentic Coding & Frontier Reasoning

GPT-5.4-pro and variants claim #1 on SWE Atlas Test Writing (44.36%), SWE Atlas CB QnA (40.80%), SWE-Bench Pro Public (59.10%), Humanity's Last Exam (44.32%), HLE Text Only (45.32%), and EnigmaEval (23.82%). Leads are consistent but scores remain well below saturation.

Gemini Dominates Audio & Vision-Language

Gemini holds all top-3 spots on AudioMultiChallenge Text Output (54.65) and leads Audio Output (36.06) and VISTA vision-language (54.65). Only VisualToolBench escapes Gemini's sweep.

Claude Leads Safety, Honesty & Professional Reasoning

Claude Opus 4.6 scores 96.28% on MASK honesty and tops both Professional Reasoning benchmarks (Legal: 52.27%, Finance: 53.28%). Fortress national security remains low (8–13%) across all labs.

Muse Spark Emerges as a Strong New Entrant

Ranks top-3 on 6+ benchmarks: MCP Atlas #1 (78.30%), MultiChallenge #1 (75.52%), SWE-Bench Pro Public #2 (55.00%), Private #2 (44.70%), HLE #2 (40.56%), Humanity's Last Exam #2 (40.56%).

Benchmarks Remain Far from Saturated

EnigmaEval top score: 23.82%. SciPredict top score: 25.27%. HLE top score: 45.32%. These gaps confirm substantial room for improvement across all frontier models.