
Scale SEAL Leaderboards

Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across agentic, frontier, safety, and multimodal benchmarks

Scale Labs | Leaderboard Analysis | 2025

Agentic

Frontier

Safety

Multimodal

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

SWE Atlas – Test Writing

SWE Atlas – Codebase QnA

MCP Atlas (Tool Use)

SWE-Bench Pro (Public)

SWE-Bench Pro (Private)

Frontier

Humanity's Last Exam

HLE (Text Only)

EnigmaEval

MultiChallenge

MultiNRC

SciPredict

Safety

MASK (Honesty)

PropensityBench

Fortress (Nat'l Security)

Prof. Reasoning – Legal

Prof. Reasoning – Finance

Remote Labor Index

Multimodal


AudioMultiChallenge (Text)

AudioMultiChallenge (Audio)

VISTA (Vision-Language)

VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro



Benchmark	#1 Model	Score	#2 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70

KEY INSIGHT


GPT-5.4 variants hold #1 on 4/5 agentic coding benchmarks. Muse Spark is top-3 on 4/5. Claude Opus 4.6 leads SWE-Bench Pro (Private).

GPT-5.4-Pro Leads Humanity’s Last Exam at 44.3% — Scores Remain Below 50%

Humanity’s Last Exam



EnigmaEval



MultiChallenge



Benchmark	#1 Model	Score	#2
MultiNRC	GPT-5-pro	65.20	Gemini-3.1-pro
SciPredict	Gemini-3-pro-preview	25.27	Claude Opus 4.5
HLE (Text Only)	GPT-5.4-pro	45.32	Muse Spark

BENCHMARKS REMAIN UNSATURATED

Even the best models score well below 50% on key reasoning benchmarks:

- Humanity’s Last Exam: 44.3% (top)
- EnigmaEval: 23.8% (top)
- SciPredict: 25.3% (top)

These represent the frontier of unsolved AI capability.

Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge (Text Output)

Gemini-3-pro-preview



Gemini-2.5-pro



Gemini-2.5-flash



AudioMultiChallenge (Audio Output)

Gemini-3.1-flash-live-preview



GPT-realtime-1.5



VISTA (Vision-Language)

Gemini 2.5 Pro Exp.



GPT-5.4-pro



VisualToolBench: GPT-5.4 edges Gemini-3.1-pro (29.17 vs 28.97) — the only multimodal benchmark not led by Gemini.

Gemini holds all top-3 positions on AudioMultiChallenge (Text Output) and leads both VISTA and Audio Output leaderboards.

Claude Models Score 95–96% on MASK Honesty; Fortress Scores Stay Low

MASK (Honesty) — Near Ceiling

Claude Opus 4.6



Claude Sonnet 4



PropensityBench (lower = better)

o3-2025-04-16



Fortress (National Security) — Critically Low

#1 Model (unnamed)



Claude Sonnet 4



Professional Reasoning

Legal – Claude Opus 4.6



Finance – Claude Opus 4.6



Finance – Muse Spark

Safety alignment varies dramatically by task type: near-ceiling honesty (96%) vs. critically low national security scores (8–13).

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	OpenAI	Google	Anthropic	Muse Spark
SWE Atlas – Test Writing	★ #1			
SWE Atlas – Codebase QnA	★ #1			
MCP Atlas (Tool Use)				★ #1
SWE-Bench Pro (Public)	★ #1			
SWE-Bench Pro (Private)			★ #1	
Humanity’s Last Exam	★ #1			
EnigmaEval	★ #1			
MultiChallenge				★ #1
MultiNRC	★ #1			
SciPredict		★ #1		
AudioMultiChallenge (Text)		★ #1		
AudioMultiChallenge (Audio)		★ #1		
VISTA		★ #1		
VisualToolBench	★ #1			
MASK (Honesty)			★ #1	
PropensityBench	★ #1			
Fortress				
Prof. Reasoning – Legal			★ #1	

OpenAI: 7 #1s · Google: 4 #1s · Anthropic: 4 #1s · Muse Spark: 2 #1s · Fortress: no model above threshold

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier



GPT-5.4 leads agentic & reasoning

Dominates SWE Atlas, SWE-Bench Pro (Public), HLE, and EnigmaEval — but does not lead multimodal or safety benchmarks.



Gemini dominates audio & vision

Sweeps all top-3 spots on AudioMultiChallenge (Text), leads VISTA and Audio Output. Only VisualToolBench is contested.



Claude leads safety & honesty

96%+ on MASK honesty, #1 on both Professional Reasoning benchmarks (Legal & Finance), and leads SWE-Bench Pro (Private).



Key benchmarks remain unsaturated

HLE (44%), EnigmaEval (24%), SciPredict (25%) — frontier models still have substantial room for improvement on hard reasoning tasks.



Muse Spark: strong new entrant

Ranks top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5% and MCP Atlas at 78.3%.

Model specialization is the defining trend — choose models by task domain, not brand.