

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across Agentic Coding · Frontier Reasoning · Safety & Alignment ·
Multimodal Understanding

20+ Benchmarks Spanning Four Capability Domains — Evaluating 100+ Models

AGENTIC	FRONTIER	SAFETY	MULTIMODAL
Coding & Tool Use - SWE Atlas (Test Writing) - SWE Atlas (Codebase QnA) - MCP Atlas (Tool Use) - SWE-Bench Pro (Public) - SWE-Bench Pro (Private)	Reasoning & Knowledge - Humanity's Last Exam (HLE) - EnigmaEval - MultiChallenge - MultiNRC (Multilingual) - SciPredict	Alignment & Honesty - MASK (Honesty) - PropensityBench - Fortress (Nat'l Security) - Professional Reasoning (Legal) - Remote Labor Index (RLI)	Audio & Vision-Language - AudioMultiChallenge (Text) - AudioMultiChallenge (Audio) - VISTA Vision-Language - VisualToolBench

20+

Benchmarks

100+

Models

4

Capability Domains

5

Leading Labs

GPT-5.4 Dominates Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Score (% ± CI) | HIGHER = BETTER | Private evaluation prevents overfitting

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex -Xhigh	38.98
SWE Atlas – Codebase QnA	GPT-5.4-xHigh (Codex)	40.80	GPT-5.4-xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini-3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Benchmark Scores (% ± CI) | ALL SCORES BELOW 50% — significant headroom remains




⚠️ All frontier reasoning benchmarks remain far from saturation. Top HLE score of 44.3%, EnigmaEval at 23.8%, and SciPredict at 25.3% — substantial room for improvement.

Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge


Text Output Leaderboard

Higher = Better




Vision-Language: VISTA & VisualToolBench

VISTA Vision-Language




VisualToolBench



Gemini variants hold ALL top-3 positions on AudioMultiChallenge (Text Output) and lead VISTA vision-language understanding.

Claude Scores 95–96% on MASK Honesty; Fortress Scores Stay Critically Low



OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Heatmap: #1-ranked lab per benchmark | ● = domain leader

Agentic Coding		Frontier Reasoning		Safety & Alignment		Multimodal Audio/Vision	
SWE Atlas – Test Writing	●	GPT-5.4-xHigh (OpenAI)	●				
SWE Atlas – Codebase QnA	●	GPT-5.4-xHigh (OpenAI)	●				
MCP Atlas (Tool Use)	●	Muse Spark (New Entrant)	●				
SWE-Bench Pro (Public)	●	GPT-5.4-pro (OpenAI)	●				
SWE-Bench Pro (Private)	●	Claude Opus 4.6 (Anthropic)	●				
Humanity's Last Exam		GPT-5.4-pro (OpenAI)	●				
HLE (Text Only)		GPT-5.4-pro (OpenAI)	●				
EnigmaEval		GPT-5.4-pro (OpenAI)	●				
MultiChallenge		Muse Spark (New Entrant)	●				
MultiNRC		GPT-5-pro (OpenAI)	●				
SciPredict		Gemini-3-pro (Google)	●				
MASK Honesty				Claude Opus 4.6 (Anthropic)	●		
PropensityBench				o3-2025-04-16 (OpenAI)	●		
Fortress				Fortress Anon (New AI)	●		
Prof. Reasoning – Legal				Claude Opus 4.6 (Anthropic)	●		
Prof. Reasoning – Finance				Claude Opus 4.6 (Anthropic)	●		
AudioMultiChallenge (Text)						Gemini-3-pro (Google)	●
AudioMultiChallenge (Audio)						Gemini-3.1-flash (Google)	●
VISTA Vision-Language						Gemini-2.5-Pro (Google)	●
Domain Leaders: VisualToolBench	● OpenAI	● Google	● Anthropic	● Muse Spark	● Low AI		GPT-5.4 (OpenAI)

Muse Spark is a notable new entrant dominating tool use (MCP Atlas) and MultiChallenge reasoning.

No Single Model Wins Everywhere — Specialization Is the New Frontier

01 **GPT-5.4 Leads Agentic & Frontier Reasoning**

GPT-5.4 variants dominate SWE Atlas (44.4%), SWE-Bench Pro Public (59.1%), Humanity's Last Exam (44.3%), and EnigmaEval (23.8%). But Claude and Gemini consistently rank #2 — the gap is narrowing.

02 **Gemini Dominates Audio & Vision**

Gemini variants sweep all top-3 positions on AudioMultiChallenge (Text Output) and lead VISTA vision-language (54.65). GPT-5.4 is a close #2 on VISTA (53.89), making vision a two-horse race.

03 **Claude Leads Safety & Professional Reasoning**

Claude Opus 4.6 scores 96.3% on MASK honesty — near-ceiling performance. It also leads both Professional Reasoning benchmarks (Legal: 52.3%, Finance: 53.3%). But Fortress scores remain critically low (8–13%) across ALL labs.

04 **Muse Spark: A Strong New Entrant**

Muse Spark ranks #1 on MultiChallenge (75.5%) and MCP Atlas (78.3%), places top-3 on 6+ benchmarks including HLE (#2 at 40.6%), SWE-Bench Pro Public (#2 at 55%), and Finance Reasoning (#2 at 52.4%). Watches this new entrant closely.

05 **Benchmarks Remain Far From Saturation — Room for All**

HLE: 44.3%, EnigmaEval: 23.8%, SciPredict: 25.3%. Even the best frontier models score below 50% on the hardest reasoning benchmarks. The ceiling has not been reached — competitive differentiation will grow.