

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Be

100+ models evaluated across Agentic, Frontier, Safety, and Multimodal domains

Audience: ML researchers, AI lab leadership, engineering managers

Goal: Compare frontier model strengths, gaps, and domain specialization trends

Dataset scope: 20+ benchmarks, private + open sets, confidence intervals reported

Labs featured: OpenAI, Google, Anthropic, Muse Spark, others

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

SWE Atlas (Test Writing, Codebase QnA)
MCP Atlas
SWE-Bench Pro (Public/Private)

Frontier

Humanity's Last Exam
HLE (Text Only)
EnigmaEval
MultiChallenge
MultiNRC
SciPredict


Safety

MASK (Honesty)
PropensityBench
Fortress
Professional Reasoning
Remote Labor Index

Multimodal

AudioMultiChallenge (Text/Audio)
VISTA
VisualToolBench


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE



Key readouts

- OpenAI #1 on 3/4 coding subtasks
- Claude #1 on SWE-Bench Pro Private (47.1)
- Muse Spark consistently top-2 on SWE-Bench Pro
- GPT-5.4 variants lead both SWE Atlas tasks

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain E




Interpretation


- HLE leader is still <50% (44.32)
- EnigmaEval ceiling remains low (23.82)
- MultiChallenge higher, led by Muse Spark (75.52)
- Frontier reasoning remains unsaturated

Gemini Models Sweep Audio and Vision-Language Benchmarks


AudioMultiChallenge (Text Output) Top 3




AudioMultiChallenge (Audio Output) Top 2



VISTA Top 2



Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress



Variance across safety tasks

- MASK: Claude Opus 4.6 at 96.28
- Fortress top score only 8.24
- Claude Sonnet 4 Fortress: 12.80
- PropensityBench: o3 at 10.5 (lower is better)
- Alignment is benchmark-dependent

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

SWE Atlas Test Writing	OpenAI
SWE Atlas Codebase QnA	OpenAI
MCP Atlas	Muse Spark
SWE-Bench Pro Public	OpenAI
SWE-Bench Pro Private	Anthropic
Humanity's Last Exam	OpenAI
EnigmaEval	OpenAI
MultiChallenge	Muse Spark
SciPredict	Google
AudioMC Text	Google
AudioMC Audio	Google
VISTA	Google
MASK	Anthropic
Professional Reasoning Legal	Anthropic

#1 count (shown benchmarks)

- OpenAI: 6
- Google: 4
- Anthropic: 3
- Muse Spark: 2
- No single-lab monopoly

Key Takeaways: No Single Model Wins Everything — Specialization Is

- GPT-5.4 leads agentic coding and key frontier reasoning tasks, but does not dominate multimodal leaderboards.
- Gemini variants hold top positions across audio and vision-language benchmarks (AudioMultiChallenge, VISTA).
- Claude models lead safety and professional reasoning, with 96%+ MASK honesty performance.
- Hard benchmarks remain far from saturation: Humanity's Last Exam 44.32%, EnigmaEval 23.82%, SciPredict 25.27%.
- Muse Spark is a credible new entrant: top-3 on 6+ benchmarks and #1 on MultiChallenge (75.52%).