

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated across Agentic, Frontier, Safety, and Multimodal domains | Scale Labs

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Scale Labs combines private datasets to prevent overfitting with open-source datasets for comparability.

Agentic

Coding & Tool Use

SWE Atlas, SWE-Bench Pro, MCP Atlas

Frontier

Reasoning & Knowledge

Humanity's Last Exam, EnigmaEval, MultiChallenge

Safety

Alignment & Honesty

MASK, Fortress, PropensityBench


Multimodal

Audio & Vision-Language

AudioMultiChallenge, VISTA, VisualToolBench


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Grouped comparison of top-3 models across SWE Atlas subtasks, MCP Atlas, and SWE-Bench Pro (Public & Private).





GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Even top-performing models are far from saturation on frontier reasoning and scientific forecasting benchmarks.




Gemini Models Sweep Audio and Vision-Language Benchmarks

Google holds all top-3 positions on AudioMultiChallenge (Text Output) and leads VISTA vision-language understanding.



Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low

Safety alignment varies dramatically by task type: near-ceiling honesty scores contrast with critically low national-security robustness.




Key Safety Insights

- MASK (Honesty): Claude Opus 4.6 at 96.28% — near ceiling.
- Fortress (National Security): Best score only 8.24% — all models critically low.
- PropensityBench: o3 leads at 10.50 (lower = better).
- Professional Reasoning: Claude #1 on Legal (52.27%) and Finance (53.28%).
- RLI: Claude Opus 4.6 leads at 4.17.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Heatmap matrix mapping every benchmark to the lab holding the #1-ranked model. No single lab wins everywhere.

Benchmark #1 Rankings by Lab



Key Takeaways: No Single Model Wins Everything — ~~Specialization~~ Is the New Frontier

1. GPT-5.4 leads agentic coding and frontier reasoning

Dominant on SWE Atlas, SWE-Bench Pro, Humanity's Last Exam, and EnigmaEval — but not multimodal.

2. Gemini dominates audio and vision-language

Sweeps AudioMultiChallenge (all top-3 text-output spots) and leads VISTA at 54.65%.

3. Claude leads safety, honesty, and professional reasoning

Near-ceiling MASK scores (~96%), #1 on Legal and Finance reasoning, but Fortress remains critically low.

4. Benchmarks remain far from saturation

Top scores: HLE 44%, EnigmaEval 24%, SciPredict 25% — frontier models still have substantial room to improve.

5. Muse Spark emerges as a strong new competitor

Top-3 on 6+ benchmarks including #1 on MultiChallenge (75.5%), MCP Atlas, and SWE-Bench Pro.