

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across agentic, frontier, safety, and multimodal capabilities



OpenAI



Google



Anthropic



Muse Spark

Source: scale.com/leaderboard

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

AGENTIC

Coding & Tool Use

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)

FRONTIER

Reasoning & Knowledge

- Humanity’s Last Exam
- HLE (Text Only)
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

SAFETY

Alignment & Honesty

- MASK (Honesty)
- PropensityBench
- Fortress (Nat’l Security)
- Prof. Reasoning – Legal
- Prof. Reasoning – Finance
- Remote Labor Index

MULTIMODAL

Audio & Vision-Language

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA (Vision-Language)
- VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36 ±6.04	GPT-5.4-xHigh (Mini-SWE)	40.00 ±6.00	GPT-5.3-Codex -Xhigh	38.98 ±6.12
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80 ±5.10	GPT 5.4 xHigh (Mini-SWE)	36.30 ±4.90	Opus 4.6 (Claude Code)	33.30 ±5.00
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10 ±3.56	Muse Spark	55.00 ±3.60	Claude Opus 4.6	51.90 ±3.61
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10 ±6.07	Muse Spark	44.70 ±6.05	GPT-5.4-pro (xHigh)	43.40 ±6.03

OpenAI

Anthropic


Muse Spark

All scores include confidence intervals. Cell shading indicates lab affiliation.

- ▶ GPT-5.4 holds #1 on 3 of 4 agentic coding benchmarks
- ▶ Claude Opus 4.6 leads SWE-Bench Pro (Private) at 47.1%
- ▶ Muse Spark appears in top 3 on 3 benchmarks, including #2 on SWE-Bench Pro (Public)
- ▶ OpenAI variants sweep the top 2 positions on both SWE Atlas subtasks

GPT-5.4-Pro Leads Humanity’s Last Exam at 44.3% — Scores Remain Below 50%

Score



Unsaturated Benchmarks

Even the best models score well below 50% on HLE and EnigmaEval, revealing substantial headroom.

GPT-5.4-pro leads HLE

44.3% on Humanity’s Last Exam — the hardest general reasoning benchmark.

Muse Spark tops MultiChallenge


75.5% on MultiChallenge, beating Gemini and GPT-5.4.

EnigmaEval < 24%

Complex multi-step reasoning remains extremely challenging for all models.

SciPredict: Gemini-3-pro-preview leads at 25.3%, followed by Claude Opus 4.5 (23.1%).

Gemini Models Sweep Audio and Vision-Language Benchmarks




- ▶ Gemini sweeps all top-3 positions on AudioMultiChallenge (Text Output) with scores of 54.7, 46.9, and 40.0
- ▶ Gemini edges GPT-realtime-1.5 on Audio Output (36.1 vs 34.7); GPT-5.4 narrowly leads VisualToolBench (29.2 vs 29.0)
- ▶ VISTA is the closest contest: Gemini 2.5 Pro (54.7) vs GPT-5.4-pro (53.9) — a margin of less than 1 point

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low


MASK (Honesty) — Near Ceiling

MASK Score



Fortress (Nat'l Security) — Near Floor

Fortress Score



Additional Safety Benchmarks

PropensityBench (lower = better)

o3-2025-04-16: 10.50 ±0.60
Claude Sonnet 4: —

Professional Reasoning – Legal

Claude Opus 4.6: 52.27 ±0.66

Professional Reasoning – Finance

Claude Opus 4.6: 53.28 ±0.18
Muse Spark: 52.44 ±0.06

Remote Labor Index

Claude Opus 4.6 (CoWo): 4.17
Claude Opus 4.5: 3.75

KEY INSIGHT

Safety alignment varies dramatically by task type: Claude models achieve near-ceiling 96% on MASK honesty, yet Fortress national security scores stay critically low (8–13) across ALL models, revealing a fundamental gap in current alignment capabilities.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	OpenAI	Google	Anthropic	Muse Spark
SWE Atlas – Test Writing	#1			
SWE Atlas – Codebase QnA	#1			
MCP Atlas (Tool Use)				#1
SWE-Bench Pro (Public)	#1			
SWE-Bench Pro (Private)			#1	
Humanity’s Last Exam	#1			
HLE (Text Only)	#1			
EnigmaEval	#1			
MultiChallenge				#1
MultiNRC	#1			
SciPredict		#1		
AudioMultiChallenge (Text)		#1		
AudioMultiChallenge (Audio)		#1		
VISTA		#1		
VisualToolBench	#1			
MASK (Honesty)			#1	
PropensityBench	#1			
Fortress				
Prof. Reasoning – Legal			#1	
Prof. Reasoning – Finance			#1	
Remote Labor Index			#1	

#1 Rankings by Lab

OpenAI

9 #1 positions

Agentic coding, frontier reasoning

Google

4 #1 positions

Audio, vision, scientific reasoning

Anthropic

5 #1 positions

Safety, professional reasoning

Muse Spark

2 #1 positions

Tool use, multi-step reasoning

Fortress has no #1 entrant — scores remain critically low across all models (8–13 range).

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

GPT-5.4 leads agentic coding and frontier reasoning

#1 on SWE Atlas, SWE-Bench Pro (Public), and Humanity's Last Exam — but does not lead in multimodal or safety domains.

Gemini dominates audio and vision benchmarks

Sweeps all top-3 positions on AudioMultiChallenge (Text) and leads VISTA vision-language. Google's multimodal investment pays off.

Claude leads safety alignment and professional reasoning

96%+ on MASK honesty and #1 on both Professional Reasoning benchmarks (Legal and Finance). Safety specialization is Anthropic's edge.

Key benchmarks remain far from saturation

HLE (44%), EnigmaEval (24%), and SciPredict (25%) show substantial room for improvement — even the best models are well below ceiling.

Muse Spark emerges as a strong new competitor

Ranks top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5%.

Model specialization — not general dominance — is the defining trend of the current frontier landscape.