

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ Benchmarks

A data-driven analysis of 100+ models across Agentic, Frontier, Safety, and Multimodal domains

For ML Researchers | AI Lab Leadership | Engineering Managers

20+ Benchmarks Spanning Four Capability Domains

Scale Labs evaluates 100+ models across Agentic, Frontier, Safety, and Multimodal benchmarks

Agentic

- SWE Atlas — Test Writing
- SWE Atlas — Codebase QnA
- SWE-Bench Pro (Public/Private)
- MCP Atlas (Tool Use)

Frontier

- Humanity's Last Exam (HLE)
- EnigmaEval
- MultiChallenge
- MultiNRC • SciPredict

Safety

- MASK (Honesty)
- PropensityBench
- Fortress
- Professional Reasoning (Legal/Finance)

Multimodal

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA (Vision-Language)
- VisualToolBench

100+

Models Evaluated

20+

Benchmarks

4

Capability Domains

3

Leading Labs





GPT-5.4 Dominates Agentic Coding

Top scores on SWE Atlas and SWE-Bench Pro — Claude Opus 4.6 leads on private repos

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3%





Frontier reasoning scores remain well below 50% — benchmarks stay unsaturated



⚠️ **Top scores: HLE 44.3% | EnigmaEval 23.8% | SciPredict 25.3% — benchmarks remain far from saturated**

Gemini Models Sweep Audio and Vision-Language Benchmarks

Google holds top positions across AudioMultiChallenge (text & audio) and VISTA




★ Gemini-3-pro-preview leads AudioMultiChallenge (Text) at 54.65 — sweeps all top-3 positions in audio benchmarks


Claude Scores 95–96% on MASK; Fortress Scores Stay Critically Low

Safety alignment varies dramatically by task type — Claude dominates honesty, Fortress remains challenging for all models

MASK Honesty — Claude Sonnet 4: 96.13%



Fortress (National Security) — All models: 8–13%



MASK: 96% — Near-ceiling honesty performance
Claude Sonnet 4 & Opus 4.6 lead with near-perfect honesty benchmark scores

Fortress: 8–13% — Significant national security gap
All models score critically low on Fortress — substantial room for safety improvement

PropensityBench (lower=better): o3-2025-04-16 leads at 10.5 | Professional Reasoning – Legal: Claude Opus 4.6 (52.27%) | Finance: Claude Opus 4.6 (53.28%)

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

No single lab dominates — model specialization defines the competitive landscape



■ OpenAI — #1 on 7 benchmarks ■ Google — #1 on 4 benchmarks ■ Anthropic — #1 on 5 benchmarks ■ Muse — #1 on 2 benchmarks

Source: Scale Labs SEAL Leaderboards | scale.com/leaderboard

Key Takeaways: Specialization Is the New Frontier

- 1 GPT-5.4 leads agentic coding & frontier reasoning**
Top scores on SWE Atlas, SWE-Bench Pro, HLE, EnigmaEval — but not multimodal or safety
- 2 Gemini dominates audio & vision-language benchmarks**
Sweeps AudioMultiChallenge (text & audio output) and VISTA; all top-3 audio positions
- 3 Claude leads safety & professional reasoning**
MASK honesty near-ceiling at 96%; #1 on both Legal (52.27%) and Finance (53.28%) reasoning
- 4 Key benchmarks remain far from saturation**
HLE 44%, EnigmaEval 24%, SciPredict 25% — substantial room for frontier improvement
- 5 Muse Spark: a formidable new entrant**
Top 3 on 6+ benchmarks: #1 on MultiChallenge (75.5%) and MCP Atlas (78.3%)