

Scale SEAL Leaderboards: Frontier Model Performance

Comparative snapshot across agentic coding, frontier reasoning, safety alignment, and multimodal (audio + vision-language).

100+ models evaluated across Scale Labs' benchmark suite.

OpenAI

Google

Anthropic

Muse Spark

20+ Benchmarks Spanning Four Capability Domains Eval

Agentic

SWE Atlas (Test Writing)
SWE Atlas (Codebase QnA)
MCP Atlas (Tool Use)
SWE-Bench Pro (Public/Private)

Frontier

Humanity's Last Exam (HLE)
EnigmaEval
MultiChallenge
SciPredict
MultiNRC (Multilingual)


Safety

MASK (Honesty)
Fortress (NatSec)
PropensityBench
Professional Reasoning (Legal/Finance)
Remote Labor Index (RLI)

Multimodal

AudioMultiChallenge (Text/Audio)
VISTA (Vision-Language)
VisualToolBench
(20+ total benchmarks)

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE



Scores are leaderboard means; zero bars indicate model not in top-3 for that subtask.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Score

Humanity's Last Exam (HLE)



EnigmaEval



MultiChallenge



 **EnigmaEval**

Evaluating model performance on complex, multi-step reasoning tasks

1

gpt-5.4-pro...

NEW

ⓘ

23.82 ±2.43

EnigmaEval leaderboard (Scale)

Frontier reasoning: HLE, EnigmaEval, MultiChallenge

gemini-3.1-

NEW

ⓘ

19.76 ±2.27


Gemini Models Sweep Audio and Vision-Language Bench

Google

OpenAI

AudioMultiChallenge - Audio Output

Evaluating spoken dialogue systems in multi-turn interaction



VISTA

Vision-Language Understanding benchmark for multimodal models



Key leaderboard results (Scale SEAL)

AudioMultiChallenge (Text Output): Gemini-3-pro-preview 54.65 ±4.57 (Gemini-2.5-pro 46.90; Gemini-2.5-flash 40.04)
AudioMultiChallenge (Audio Output): Gemini-3.1-flash-live-preview 36.06 ±4.41 vs GPT-4o 34.73
VISTA: Gemini 2.5 Pro Experimental 54.65 ±1.46 vs GPT-5.4-pro 53.89 ±2.02
VisualToolBench: GPT-5.4 29.17 vs Gemini-3.1-pro 28.97

Claude Models Score 95–96% on MASK Honesty Bench

MASK (Honesty) — near ceiling

Claude Opus 4.6 — 96.28%




Claude Sonnet 4 — 96.13%



Honesty-style benchmarks are close to saturation for top Claude variants.

Fortress (National Security) — scores stay low

Fortress (score) — 8.24



Claude Sonnet 4 — 12.80



Alignment difficulty varies by task: high MASK does not imply high Fortress performance

OpenAI, Google, and Anthropic Each Claim #1 Across D

Benchmark	OpenAI	Google	Anthropic	Muse Spark
SWE Atlas – Test Writing	#1			
SWE Atlas – Codebase QnA	#1			
MCP Atlas (Tool Use)				#1
SWE-Bench Pro (Public)	#1			
SWE-Bench Pro (Private)			#1	
Humanity's Last Exam	#1			
EnigmaEval	#1			
MultiChallenge				#1
MultiNRC (Multilingual)	#1			
SciPredict		#1		
AudioMultiChallenge (Text Output)		#1		
AudioMultiChallenge (Audio Output)		#1		
VISTA (Vision-Language)		#1		
VisualToolBench	#1			
MASK (Honesty)			#1	
PropensityBench (lower=better)	#1			
Professional Reasoning – Legal			#1	
Professional Reasoning – Finance			#1	
Remote Labor Index (RLI)			#1	
Fortress (NatSec)				

Matrix shows #1-ranked lab/model per benchmark (only benchmarks listed in the brief). Heatmap: #1 lab/model per benchmark

Key Takeaways: No Single Model Wins Everything — Sp

Implications

GPT-5.4 leads agentic coding + frontier reasoning, but not multimodal.

Gemini dominates audio/vision leaderboards (AudioMultiChallenge + VISTA).

Claude leads safety/honesty (MASK ~96%) and professional reasoning (Legal/Finance).

Key benchmarks remain unsaturated: HLE ~44%, EnigmaEval ~24%, SciPredict ~25%.

Muse Spark is a notable new entrant: top-3 on 6+ benchmarks, incl. #1 on MultiChallenge (75.52%).

OpenAI

Google

Anthropic

Muse Spark

Conclusion: specialization across domains