

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

Comparative analysis of 100+ models across Agentic, Frontier, Safety, and Multimodal domains

Audience: ML researchers, AI lab leadership, engineering managers

Source: Scale SEAL leaderboard brief (all figures and rankings)


Objective: map performance leadership and specialization trends

Core finding: no single lab leads all capability domains

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic	Frontier	Safety	Multimodal
• SWE Atlas (Test Writing) • SWE Atlas (Codebase QnA) • SWE-Bench Pro Public • SWE-Bench Pro Private • MCP Atlas • MCPToolBench	• Humanity's Last Exam • EnigmaEval • MultiChallenge • SciPredict	• MASK (Honesty) • PropensityBench • Fortress • Professional Reasoning (Legal/Finance) • Remote Labor Index	• AudioMultiChallenge (Text) • AudioMultiChallenge (Audio) • VISTA


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Ber



Top #2 highlights:

- SWE Atlas TW: Claude 67.6
 - SWE Atlas QnA: Claude 72.0
 - SB Pro Public: Gemini 59.4
 - SB Pro Private: Claude 36.2
- OpenAI model family leads all 4 benchmarks.


GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below



Top models

- HLE: GPT-5.4-pro (44.33)
 - Enigma: GPT-5.4-pro (24.40)
 - MultiChallenge: Muse Spark (75.5)
 - SciPredict: Gemini-2.5-pro (25.23)
- Frontier reasoning remains unsaturated.


Gemini Models Sweep Audio and Vision-Language Benchmarks



Gemini sweep details

- AudioMultiChallenge (Text): top 3 all Gemini variants
 - AudioMultiChallenge (Audio): top 3 all Gemini variants
 - VISTA: Gemini-2.5-pro leads at 72.5
- Google holds #1 across audio + vision-language leaderboards.

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Score



Contrast view

- MASK near ceiling: ~96%
 - Fortress remains low: 8–13
 - PropensityBench (lower=better): o3 = 10.50
- Safety alignment varies sharply by task type.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

SWE Atlas TW
SWE Atlas QnA
SWE-Bench Pro Public
SWE-Bench Pro Private
MCP Atlas
MCPToolBench
HLE
EnigmaEval
MultiChallenge
SciPredict

AudioMulti Text
AudioMulti Audio
VISTA
MASK
PropensityBench
Fortress
Pro Reasoning Legal
Pro Reasoning Finance
RLI

#1 leader by benchmark
OpenAI: coding + frontier + tool use
Google: multimodal/audio/vision
Anthropic: safety + professional reasoning
Muse Spark: #1 on MultiChallenge

Key Takeaways: No Single Model Wins Everything — Specialization Is the Key

- GPT-5.4 leads agentic coding and key frontier reasoning (HLE 44.33, EnigmaEval 24.40), but does not lead multimodal.
- Gemini variants dominate audio and vision-language benchmarks, sweeping top positions in AudioMultiChallenge and leading VISTA.
- Claude models lead safety and professional reasoning (MASK ~96%; #1 in Legal and Finance reasoning), yet Fortress scores remain low.
- Core frontier benchmarks remain far from saturation: HLE 44%, EnigmaEval 24%, SciPredict 25%.
- Muse Spark is a credible new entrant: top-3 on 6+ benchmarks and #1 on MultiChallenge (75.5).