

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated across Agentic · Frontier · Safety · Multimodal

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

Coding & Tool Use

- SWE Atlas
- SWE-Bench Pro
- MCP Atlas
- Mini-SWE

Top lab: OpenAI

Frontier

Reasoning & Knowledge

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- SciPredict

Top lab: OpenAI

Safety

Alignment & Honesty

- MASK
- Fortress
- PropensityBench
- Professional Reasoning

Top lab: Anthropic

Multimodal

Audio & Vision-Language

- AudioMultiChallenge
- VISTA
- VisualToolBench


Top lab: Google

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

Note: Muse Spark (new entrant) ranks top-3 on MCP Atlas and both SWE-Bench Pro splits. Scores include confidence intervals.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%



Gemini Models Sweep Audio and Vision-Language Benchmarks




AudioMultiChallenge (Text Out)	Score
Gemini-3-pro-preview	54.65
Gemini-2.5-pro	46.90
Gemini-2.5-flash	40.04

AudioMultiChallenge (Audio Out)	Score
Gemini-3.1-flash-live-preview	36.06
GPT-realtime-1.5	34.73

Gemini holds all top-3 positions on AudioMultiChallenge (Text Output) and leads VISTA vision-language understanding.

VISTA (Vision-Language)	Score
Gemini 2.5 Pro Experimental	54.65
GPT-5.4-pro	53.89

VisualToolBench: GPT-5.4 edges Gemini-3.1-pro (28.97)



Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low

MASK (Honesty)

Near-ceiling performance on honesty alignment

Claude Opus 4.6

96.28 ±0.41

Claude Sonnet 4

96.13 ±0.57

Fortress (National Security)

Critically low scores across all models tested

#1 (unnamed)

8.24 ±1.93

Claude Sonnet 4

12.80 ±2.36

PropensityBench (lower = better): o3-2025-04-16 → 10.50 ±0.60 | Claude Sonnet 4 → —

Insight: Safety alignment varies dramatically by task type — honesty is near-solved, but national-security robustness remains a critical gap.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	OpenAI	Google	Anthropic	Muse Spark
SWE Atlas – Test Writing	✓			
SWE Atlas – Codebase QnA	✓			
MCP Atlas (Tool Use)				✓
SWE-Bench Pro (Public)	✓			
SWE-Bench Pro (Private)			✓	
Humanity's Last Exam	✓			
HLE (Text Only)	✓			
EnigmaEval	✓			
MultiChallenge				✓
MultiNRC (Multilingual)	✓			
SciPredict		✓		
AudioMultiChallenge		✓		
VISTA		✓		
VisualToolBench	✓			
MASK (Honesty)			✓	
PropensityBench	✓			
Fortress				
Prof. Reasoning – Legal			✓	
Prof. Reasoning – Finance			✓	
Matrix shows which lab's model ranks #1 on each benchmark. Blank = no #1 data or tie. Fortress has no clear #1 (8.74 unnamed). Remote Labor Index			✓	

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

1 GPT-5.4 leads agentic & reasoning

Dominates SWE Atlas, SWE-Bench Pro (Public), Humanity's Last Exam, and EnigmaEval — but does not lead multimodal or safety benchmarks.

2 Gemini dominates audio & vision

Sweeps AudioMultiChallenge (all top-3 text-output slots) and leads VISTA vision-language understanding.

3 Claude leads safety & professional reasoning

Near-ceiling MASK honesty scores (96%+) and #1 on Legal, Finance, and Remote Labor Index — yet Fortress scores remain critically low.

4 Benchmarks remain far from saturated

HLE ~44%, EnigmaEval ~24%, SciPredict ~25%. Frontier models still have substantial room for improvement on complex tasks.

5 Muse Spark is a strong new entrant

Ranks top-3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5% and MCP Atlas at 78.3%.