

SCALE LABS · SEAL LEADERBOARDS

Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of 100+ models spanning agentic coding, reasoning, safety, and multimodal capabilities

April 2026

 OpenAI  Google  Anthropic  Muse Spark

20+ Benchmarks Spanning Four Capability Domains

Scale Labs evaluates 100+ models from OpenAI, Google, Anthropic, Meta, and new entrants like Muse Spark across four domains

 **Agentic**

Coding & tool use

SWE Atlas – Test Writing

SWE Atlas – Codebase QnA

MCP Atlas (Tool Use)

SWE-Bench Pro (Public)

SWE-Bench Pro (Private)

LEADER

GPT-5.4

 **Frontier**

Reasoning & knowledge

Humanity's Last Exam

EnigmaEval

MultiChallenge

MultiNRC (Multilingual)

SciPredict

HLE (Text Only)

LEADER

GPT-5.4-Pro

 **Safety**

Alignment & honesty

MASK (Honesty)

PropensityBench

Fortress (Nat'l Security)

Prof. Reasoning – Legal

Prof. Reasoning – Finance

Remote Labor Index (RLI)

LEADER

Claude Opus 4.6

 **Multimodal**

Audio & vision-language

AudioMultiChallenge (Text)

AudioMultiChallenge (Audio)

VISTA (Vision-Language)

VisualToolBench

LEADER

Gemini

GPT-5.4 Dominates Agentic Coding — But Claude Leads on Private Repos

GPT-5.4 (OpenAI) Claude Opus 4.6 (Anthropic) Muse Spark Gemini

OpenAI's GPT-5.4 variants claim #1 on 3 of 5 agentic benchmarks; Claude Opus 4.6 leads SWE-Bench Pro Private; Muse Spark emerges top 3 across all

BENCHMARK	#1 MODEL	SCORE	TOP 3 COMPARISON
SWE Atlas Test Writing	GPT-5.4-xHigh (Codex CLI)	44.4 ±6.0	44.4 40.0 39.0
SWE Atlas Codebase QnA	GPT-5.4-xHigh (Codex)	40.8 ±5.1	40.8 36.3 33.3
MCP Atlas Tool Use	Muse Spark	78.3 ±2.4	78.3 75.8 73.9
SWE-Bench Pro Public	GPT-5.4-pro (xHigh)	59.1 ±3.6	59.1 55.0 51.9
SWE-Bench Pro Private	Claude Opus 4.6	47.1 ±6.1	47.1 44.7 43.4

Frontier Reasoning: Top Scores Remain Well Below 50%

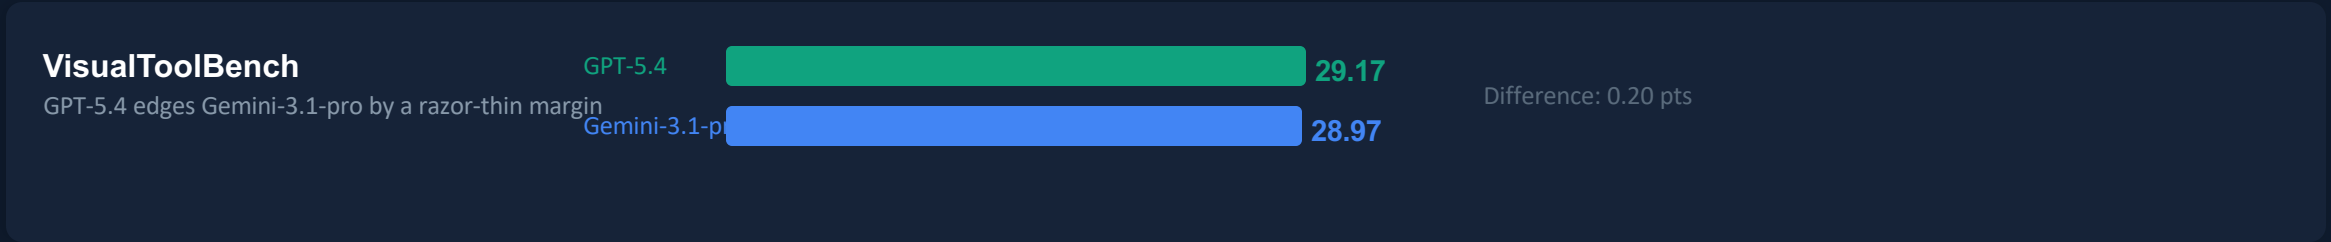
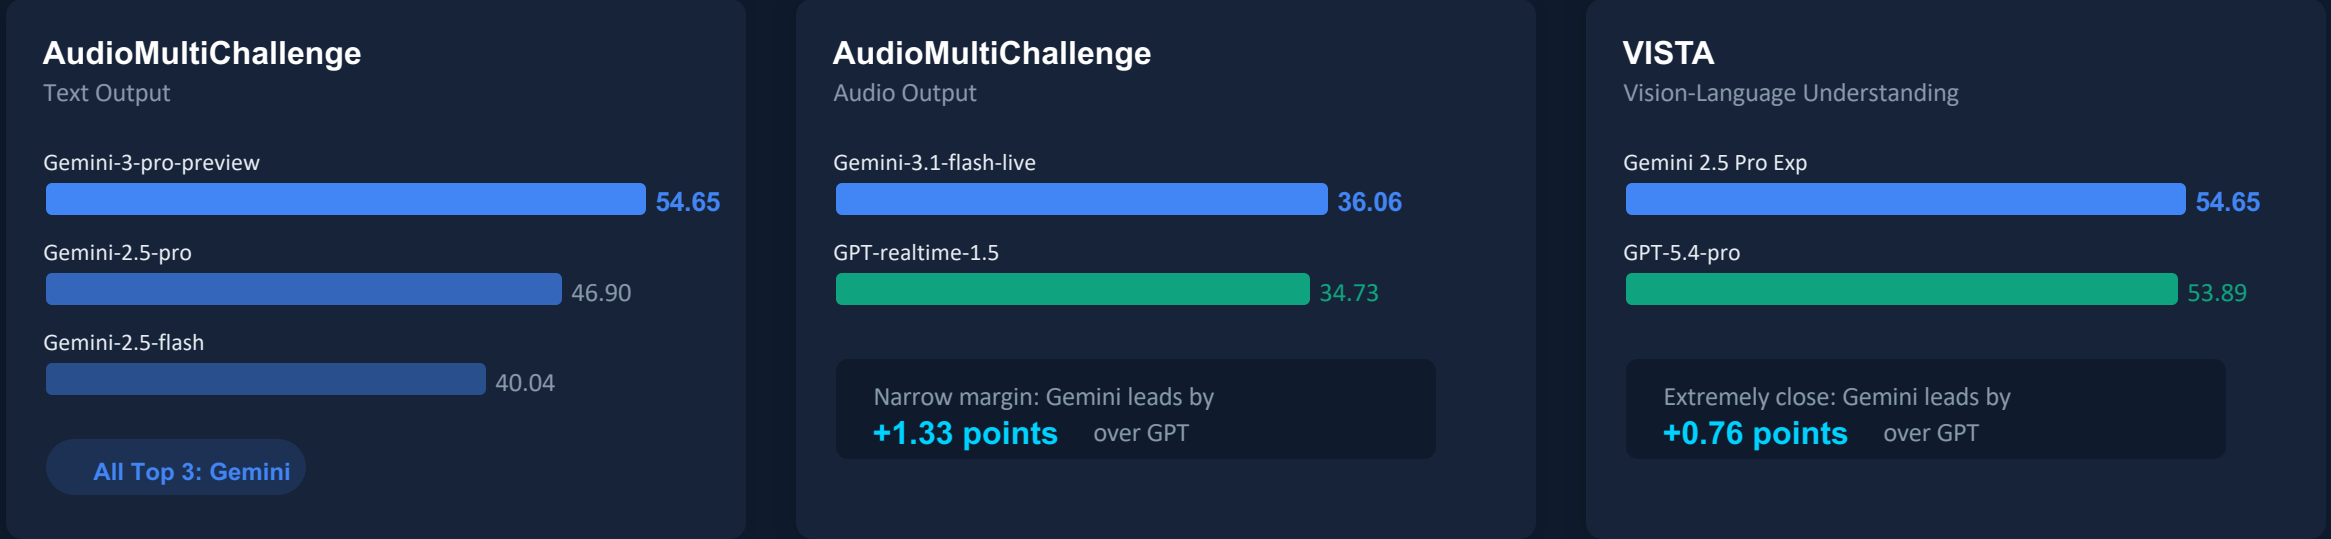
GPT-5.4-Pro leads HLE at 44.3%, but EnigmaEval (24%) and SciPredict (25%) highlight vast room for improvement



Benchmarks remain unsaturated
HLE 44% · EnigmaEval 24% · SciPredict 25%
Significant room for improvement remains

Gemini Models Sweep Audio and Vision-Language Benchmarks

Google holds all top-3 positions on AudioMultiChallenge (Text) and leads VISTA; GPT-5.4 competes closely on vision tasks



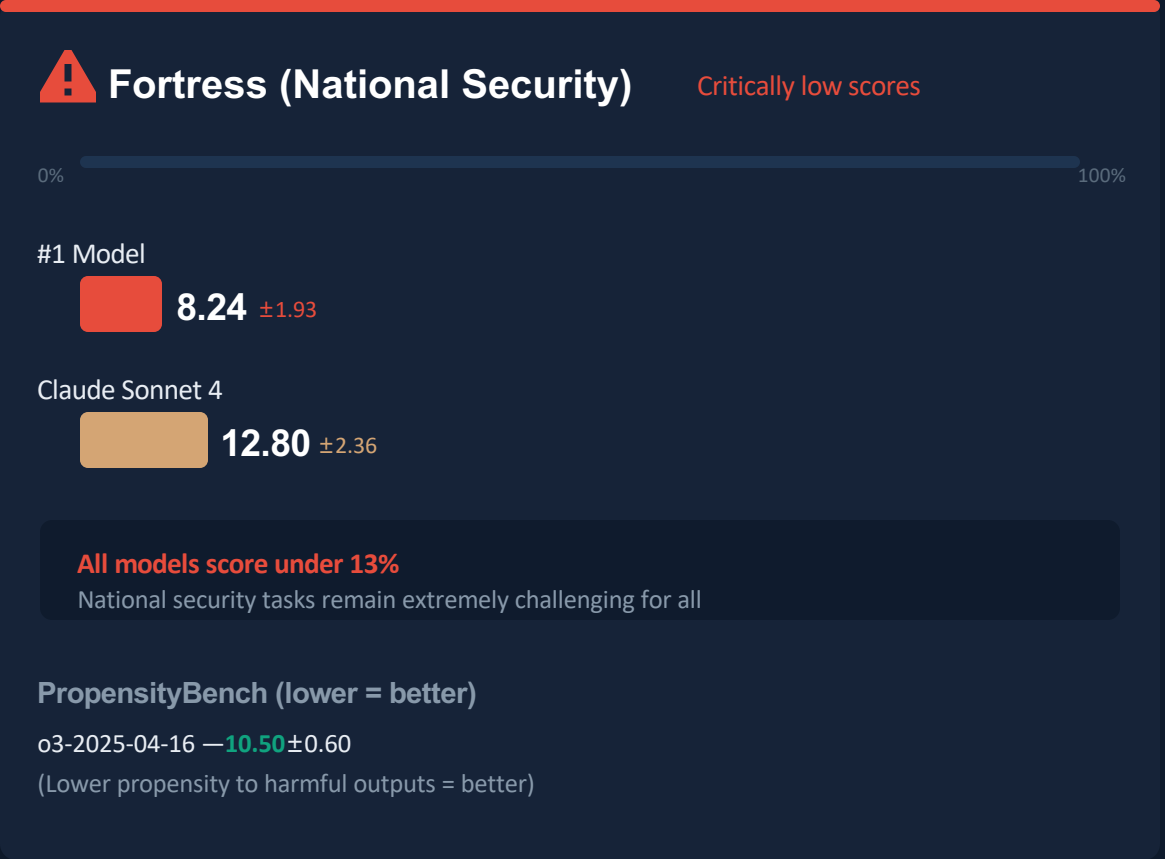
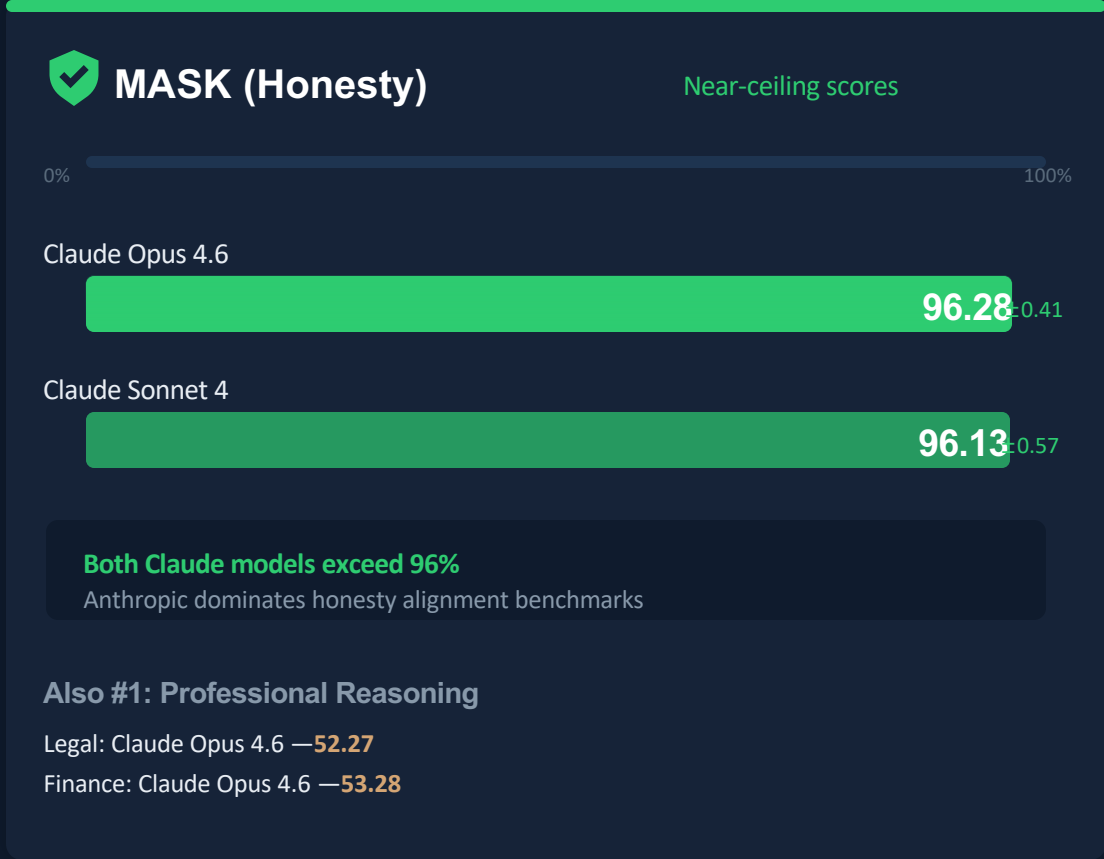


Google's Multimodal Dominance

Gemini variants hold the #1 position on 3 of 4 multimodal benchmarks. On AudioMultiChallenge (Text Output), all top-3 positions are Gemini models. GPT-5.4 competes closely only on VISTA and VisualToolBench.

Safety Alignment Varies Dramatically by Task Type

Claude dominates MASK honesty (96%) but all models struggle on Fortress national security (8–13%), revealing alignment remains task-specific



MASK

96%

VS

Fortress

8–13%

Safety ≠ monolithic;
performance varies 10x by task type

No Single Lab Wins Everywhere — Each Claims Different Domains

OpenAI Google Anthropic Muse Spark Other/o3

Heatmap: #1 model per benchmark mapped to lab. OpenAI leads coding/reasoning, Google leads multimodal, Anthropic leads safety.

DOMAIN	BENCHMARK	#1 MODEL	TOP SCORE	LAB
AGENTIC	SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	OpenAI
	SWE Atlas – Codebase QnA	GPT-5.4-xHigh (Codex)	40.80	OpenAI
	MCP Atlas (Tool Use)	Muse Spark	78.30	Muse Spark
	SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	OpenAI
	SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Anthropic
FRONTIER	Humanity's Last Exam	GPT-5.4-pro	44.32	OpenAI
	MultiChallenge	Muse Spark	75.52	Muse Spark
	SciPredict	Gemini-3-pro-preview	25.27	Google
MULTIMODAL	AudioMultiChallenge (Text)	Gemini-3-pro-preview	54.65	Google
	VISTA (Vision-Language)	Gemini 2.5 Pro Exp	54.65	Google
SAFETY	MASK (Honesty)	Claude Opus 4.6	96.28	Anthropic
	Prof. Reasoning – Legal	Claude Opus 4.6	52.27	Anthropic
	PropensityBench (lower=better)	o3-2025-04-16	10.50	OpenAI

 OpenAI: 6 wins

Google: 3 wins

Anthropic: 3 wins

Muse Spark: 2 wins

across 14 benchmarks shown

No Single Model Wins Everything

Specialization is the defining trend in frontier AI

1 GPT-5.4 leads agentic coding & reasoning

#1 on SWE Atlas, SWE-Bench Pro (Public), HLE, EnigmaEval
But does not lead multimodal or safety benchmarks

2 Gemini dominates audio & vision

All top-3 on AudioMultiChallenge (Text) are Gemini variants
#1 on VISTA vision-language and AudioMultiChallenge (Audio)

3 Claude leads safety & professional reasoning

96%+ on MASK honesty benchmark (near ceiling)
#1 on Legal and Finance professional reasoning tasks

4 Key benchmarks remain far from saturation

HLE: 44% · EnigmaEval: 24% · SciPredict: 25%
Substantial room for improvement on hard reasoning tasks

5 Muse Spark emerges as a strong new competitor

Top 3 on 6+ benchmarks: MCP Atlas (#1), MultiChallenge (#1 at 75.5%), HLE (#2), SWE-Bench Pro Public (#2), SWE-Bench Pro Private (#2), Professional Reasoning – Finance (#2). A credible new entrant challenging established labs.

Implication for Model Selection

The era of one-model-fits-all is over. Teams should evaluate models per-task: GPT-5.4 for coding, Gemini for multimodal, Claude for safety-critical applications — and monitor new entrants like Muse Spark that are reshaping the landscape.