

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of 100+ models across Agentic, Frontier, Safety, and Multimodal capabilities

SCALE LABS · scale.com/leaderboard



Scale SEAL Suite: Four Domains, 20+ Benchmarks, 100+ Models



Agentic
Coding & Tool Use

SWE Atlas
· Test Writing · Codebase Q&A

MCP Atlas (Tool Use)
SWE-Bench Pro
· Public Repository · Private Repository

Leader: OpenAI GPT-5.4 / Muse Spark



Frontier
Reasoning & Knowledge

Humanity's Last Exam (HLE)
EnigmaEval
MultiChallenge
MultiNRC (Multilingual)
SciPredict

Leader: OpenAI GPT-5.4-pro



Safety
Alignment & Honesty

MASK (Honesty)
PropensityBench
Fortress (National Security)
Professional Reasoning
· Legal · Finance
Remote Labor Index (RLI)

Leader: Anthropic Claude




Multimodal
Audio & Vision-Language

AudioMultiChallenge
· Text Output · Audio Output
VISTA (Vision-Language)
VisualToolBench


Leader: Google Gemini

Agentic Coding: Top-3 Models Across SWE Atlas and SWE-Bench Pro



GPT-5.4-Pro Leads Frontier Reasoning - Top Scores Below 50% Signal Unsaturated Benchmarks

Frontier Reasoning: Top-3 Models per Benchmark (Score %)



Multimodal Benchmarks: Google Gemini Dominates Audio and Vision

AudioMultiChallenge

Text Output



All Top-3 Text Output positions held by Gemini models

Audio Output



MCP Atlas (Tool Use) — Top Scores

Muse Spark: 78.30 | Claude Opus 4.6: 75.80 | Gemini 3.1 Pro: 73.90

Google Gemini Audio Dominance

Holds all Top 3 spots on Text Output (54.65, 46.90, 40.04)

Key Insight: Google Gemini holds a dominant multimodal position — audio sweep is comprehensive; GPT-5.4 remains competitive in vision tools.



Vision-Language Benchmarks

VISTA (Vision-Language Understanding)



Gemini leads by 0.76 pts — extremely close competition

VisualToolBench (Visual Tool Use)



GPT-5.4 edges Gemini on VisualToolBench (29.17 vs 28.97) — only vision benchmark Gemini doesn't win

Google Gemini Vision Leadership

- VISTA #1 at 54.65 (±1.46) — leads GPT-5.4-pro at 53.89
- Leads or matches GPT-5.4 on 3 of 4 multimodal benchmarks
- GPT-5.4 remains competitive — within 1pt on VISTA

Safety & Alignment: Claude Leads Honesty While Fortress Reveals Universal Weakness

 MASK Honesty Benchmark
Near-Ceiling Performance

96.28%

Claude Opus 4.6 — Rank #1




Anthropic models near ceiling — 96%+ honesty scores
Only ~4% gap from perfect honesty alignment

Anthropic #1 & #2

 Fortress (National Security)
Critically Low Scores — Universal Weakness

8.24

Best Model Score — No Lab Scores Above 13



PropensityBench (lower = better): o3-2025-04-16 leads at 10.50 ±0.60

Fortress scores remain critically low (8–13) across ALL models
National security alignment remains an industry-wide unsolved challenge

No Model Scores Above 13

Professional Reasoning — Claude Opus 4.6 #1 on Both Domains

Legal Reasoning
Claude Opus 4.6: 52.27 ±0.66

Finance Reasoning
Claude Opus 4.6: 53.28 ±0.18 | Muse Spark: 52.44

Remote Labor Index (RLI)
Claude Opus 4.6 (CoWo): 4.17 | Claude Opus 4.5: 3.75

No Single Lab Wins Everywhere — OpenAI Leads Coding/Reasoning, Google Owns Multimodal, Anthropic Dominates Safety

Competitive Landscape: #1 Rankings by Lab and Benchmark Domain

	AGENTIC					FRONTIER						SAFETY			MULTIMODAL			
	SWE-Atlas Test Write	SWE-Atlas QnA	MCP Atlas Tool Use	SWE-Bench Pro Public	SWE-Bench Pro Private	HLE Overall	HLE Text Only	Enigma Eval	Multi Challenge	Multi NRC	Sci Predict	MASK Honesty	Prof. Reason Legal	Prof. Reason Finance	Audio MC Text Out	Audio MC Audio Out	VISTA Vision-Lang	Visual ToolBench
OpenAI	1	1	—	1	3	1	1	1	3	1	—	—	—	—	—	2	2	1
Google	—	—	3	—	—	3	3	2	2	2	1	—	—	—	1	1	1	2
Anthropic	—	3	2	3	1	—	—	—	—	—	2	1	1	1	—	—	—	—
Muse Spark	—	—	1	2	2	2	2	—	1	—	—	—	—	2	—	—	—	—

#1 Rankings: OpenAI: 10 × #1 Google: 5 × #1 Anthropic: 5 × #1 Muse Spark: 2 × #1 (new!)

OpenAI Google Anthropic Muse Spark | Cell intensity: dark=1st, medium=2nd, light=3rd, gray=not in top 3

Strategic Implication: Model selection should be domain-specific — no single vendor provides best-in-class across all capability types. Muse Spark's emergence as a new entrant challenging incumbents across 6+ benchmarks signals a more competitive multivendor landscape.

Specialization Is the Defining Competitive Trend in Frontier AI

01 OpenAI leads Agentic Coding & Frontier Reasoning

GPT-5.4 ranks #1 on SWE Atlas (Test Writing: 44.4, Codebase QnA: 40.8), SWE-Bench Pro Public (59.1), and Humanity's Last Exam (44.3%) — but not multimodal.

10 benchmark #1 rankings across agentic + frontier domains

GPT-5.4 also leads MultiNRC multilingual at 65.2%

02 Google Gemini sweeps Audio & Vision benchmarks

All Top-3 AudioMultiChallenge Text Output spots held by Gemini (54.65, 46.90, 40.04). Leads VISTA vision-language at 54.65 and Audio Output at 36.06.

5 benchmark #1 rankings — comprehensive multimodal dominance

GPT-5.4 edges Gemini on VisualToolBench (29.17 vs 28.97)

03 Claude dominates Safety & Professional Reasoning

Claude Opus 4.6 scores 96.28% on MASK Honesty (near ceiling). #1 on Legal (52.27) and Finance (53.28) professional reasoning. Leads private repo coding.

5 benchmark #1 rankings — strongest safety + alignment profile

MASK #2: Claude Sonnet 4 at 96.13% — Anthropic holds top 2 spots

04 Key benchmarks remain far from saturation

Top scores: HLE 44.3% · EnigmaEval 23.8% · SciPredict 25.3%

These results indicate substantial headroom for improvement — frontier models are not close to solving complex multi-step reasoning and scientific forecasting.

Combined private datasets prevent overfitting; scores reflect genuine capability gaps

05 Muse Spark is a serious new entrant — top 3 on 6+ benchmarks

#1 MultiChallenge (75.52) · #1 MCP Atlas Tool Use (78.30)

#2 HLE (40.56) · #2 SWE-Bench Pro Public (55.0) & Private (44.7)

Competing with incumbents across agentic, reasoning, and professional reasoning benchmarks simultaneously.

Signals the competitive AI landscape is expanding beyond the OpenAI/Google/Anthropic triopoly

Source: Scale Labs SEAL Leaderboard · scale.com/leaderboard · April 2026

All scores include confidence intervals. Private datasets used to prevent benchmark overfitting.