

SCALE SEAL LEADERBOARDS

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

100+ Models | 4 Capability Domains | Comprehensive Private + Open Benchmarks

● OpenAI ● Google ● Anthropic ● Muse Spark

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Private datasets prevent overfitting. All scores include confidence intervals. Combined private + open-source for comparability.

AGENTIC

Coding & Tool Use

SWE Atlas

Test Writing | Codebase QnA

MCP Atlas

Tool Use

SWE-Bench Pro

Public | Private

5 benchmarks

FRONTIER

Reasoning & Knowledge

Humanity's Last Exam (HLE)

EnigmaEval

MultiChallenge

MultiNRC

SciPredict

6 benchmarks

SAFETY

Alignment & Honesty

MASK

Honesty benchmark

PropensityBench

Lower = better

Fortress

National Security | Prof. Reasoning (Legal, Finance) | RLI

6 benchmarks

MULTIMODAL

Audio & Vision-Language

AudioMultiChallenge

Text Output | Audio Output

VISTA

Vision-Language understanding


VisualToolBench

Visual tool use

3 benchmarks

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

GPT-5.4 variants hold #1 on 3 of 5 agentic benchmarks. Muse Spark enters top 3 on 3 benchmarks including #1 on MCP Atlas.




GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Key benchmarks remain unsaturated: HLE (44%), EnigmaEval (24%), SciPredict (25%) — substantial room for improvement.

Humanity's Last Exam



EnigmaEval



MultiChallenge



Additional Frontier Benchmarks

HLE (Text Only)
GPT-5.4-pro: 45.32%
Muse Spark: 40.92%

MultiNRC (Multilingual)
GPT-5-pro: 65.20%
Gemini-3.1-pro: 64.74%

SciPredict
Gemini-3-pro-preview: 25.27%

Benchmarks Far From Saturation

44% HLE top score — under half

24% EnigmaEval — models struggle

25% SciPredict — far from solved

Even the best models have substantial room for improvement


Gemini Models Sweep Audio and Vision-Language Benchmarks

All top-3 positions on AudioMultiChallenge (Text Output) are held by Gemini variants. VISTA is nearly tied.



Claude Models Score 95–96% on MASK Honesty; Fortress Scores Stay Critically Low

Safety alignment varies dramatically by task type — near-ceiling honesty vs. critically low national security performance.



Additional Safety & Professional Reasoning Results		
PropensityBench	Professional Reasoning	Remote Labor Index (RLI)
(Lower = better)		
10.50 o3-2025-04-16	Legal: 52.27 Claude Opus 4.6	4.17 Claude Opus 4.6 (CoWo)
Claude Sonnet 4 — noted	Finance: 53.28 Claude Opus 4.6 52.44 Muse Spark	3.75 Claude Opus 4.5
Claude leads all safety and professional reasoning benchmarks — but Fortress scores show no model is safe enough for national security tasks		

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

No single lab dominates all domains. Each leads in its specialty. Muse Spark emerges as a notable cross-domain competitor.

BENCHMARK	OPENAI	GOOGLE	ANTHROPIC	MUSE SPARK
AGENTIC CODING				
SWE Atlas – Test Writing	#1 44.4			
SWE Atlas – Codebase QnA	#1 40.8			
MCP Atlas (Tool Use)				#1 78.3
SWE-Bench Pro (Public)	#1 59.1			
SWE-Bench Pro (Private)			#1 47.1	
FRONTIER REASONING				
Humanity's Last Exam	#1 44.3			
EnigmaEval	#1 23.8			
MultiChallenge				#1 75.5
SciPredict		#1 25.3		
SAFETY & ALIGNMENT				
MASK (Honesty)	#1 96.3		Prof. Reasoning (Legal + Finance)	#1
MULTIMODAL				
		AUDIO + VISION		
AudioMultiChallenge (Text + Audio)	#1 54.7	VISTA + VisualToolBench	Google #1 VISTA	OpenAI #1 VTB

#1 Rankings by Lab

Across all benchmarks

5

OpenAI

HLE, EnigmaEval, SWE Atlas x2, VisualToolBench

4

Google

Audio x2, VISTA, SciPredict

4

Anthropic

MASK, Prof. Reasoning x2, SWE-Bench Pro Private

2

Muse Spark

MCP Atlas, MultiChallenge

New Entrant

Top-3 on 6+ benchmarks

No Single Model Wins Everything — Specialization Is the New Frontier



The defining trend: Model specialization

Each leading lab excels in its domain — no single model wins across all capabilities. Choose models for your specific use case.

● OpenAI: Agentic + Reasoning ● Google: Multimodal ● Anthropic: Safety ● Muse Spark: Cross-domain