

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of top-performing AI models across agentic coding,
reasoning, safety, and multimodal capabilities

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Scale Labs combines private and open-source datasets to prevent overfitting and ensure comparability.



Agentic

Focus Areas:

- Software Engineering
- Tool Use (MCP)
- Codebase QnA

Key Benchmarks:

SWE Atlas, SWE-Bench Pro,
MCP Atlas



Frontier

Focus Areas:

- Complex Reasoning
- Scientific Forecasting
- Multilingual Knowledge

Key Benchmarks:

Humanity's Last Exam,
EnigmaEval, MultiChallenge



Safety

Focus Areas:

- Alignment & Honesty
- National Security
- Professional Reasoning

Key Benchmarks:

MASK, PropensityBench,
Fortress, RLI



Multimodal

Focus Areas:


- Audio Output
- Vision-Language
- Visual Tool Use

Key Benchmarks:

AudioMultiChallenge, VISTA,
VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

OpenAI variants lead strongly, but Anthropic's Claude Opus 4.6 competes effectively on private repositories.



GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

These benchmarks test the absolute limits of model reasoning, indicating significant room for improvement.

Humanity's Last Exam



MultiChallenge




EnigmaEval



EnigmaEval

Evaluating model performance on complex, multi-step reasoning tasks



[View Full Ranking →](#)

Scale Labs - EnigmaEval Leaderboard


Gemini Models Sweep Audio and Vision-Language Benchmarks

Google's Gemini 2.5 and 3.1 series hold top positions across all three major multimodal leaderboards.

AudioMultiChallenge

Gemini sweeps top 3 positions for text output; leads audio output.

Evaluating spoken dialogue systems in multi-turn interaction



Text Output #1: Gemini-3-pro-preview (54.65)
Audio Output #1: Gemini-3.1-flash-live (36.06)

VISTA (Vision-Language)

Gemini 2.5 Pro edges out GPT-5.4-pro in visual reasoning.

for multimodal models



VISTA #1: Gemini 2.5 Pro (54.65)
VISTA #2: GPT-5.4-pro (53.89)

Claude Models Score 96% on Honesty; Fortress Scores Stay Dangerously Low

Anthropic's Claude leads safety benchmarks, but alignment performance varies dramatically by task type.

The Ceiling: Honesty

Models are near-perfect on basic honesty tests (MASK).

Claude Opus 4.6 (MASK)

96.3%

Professional Reasoning

Legal (Claude Opus 4.6)

52.3%

Finance (Claude Opus 4.6)

53.3%

The Floor: National Security

Complex national security risks remain a major blind spot.

13

Claude Sonnet 4 (Fortress)

PropensityBench (Lower = Better)

Measuring model propensity to engage in harmful behaviors.

10

o3-2025-04-16 (Best Score)

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

No single lab wins everywhere. Model specialization is the defining trend of 2026.

	OpenAI	Anthropic	Google	Muse Spark
Agentic Coding	GPT-5.4 (Atlas/SWE) #1 on 3 Benchmarks	Claude Opus 4.6 #1 on SWE-Bench Private		Muse Spark #1 on MCP Atlas (78.3%)
Frontier Reasoning	GPT-5.4-pro #1 on HLE, EnigmaEval, NRC			Muse Spark #1 on MultiChallenge
Multimodal			Gemini 3.1 & 2.5 #1 on Audio & Vision	
Safety & Honesty		Claude Models #1 on MASK (Honesty)		
Prof. Reasoning		Claude Opus 4.6 #1 on Legal, Finance, RLI		

No Single Model Wins Everything — Specialization Is the New Frontier

Summary of the 2026 Scale SEAL Leaderboard findings across 100+ evaluated models.



- **OpenAI Dominates Core Tasks:** GPT-5.4 leads agentic coding & frontier reasoning.
- **Google Owns Multimodal:** Gemini models sweep audio and vision-language benchmarks.
- **Anthropic Leads Safety:** Claude 4.6 leads alignment and professional reasoning.
- **Muse Spark Is a Top Contender:** A strong new competitor ranking top-3 across 6+ benchmarks.
- **Benchmarks Remain Unsaturated:** HLE (44%), EnigmaEval (24%) show models are far from ceiling.