
Scale SEAL Leaderboards Frontier Model Performance Across 20+ Benchmarks

Comparative analysis of 100+ models across agentic coding, frontier reasoning, safety alignment, and multimodal capabilities

SOURCE

Scale Labs Leaderboard

DATE

April 2026

KEY INSIGHT

Scale Labs evaluates 100+ models with private datasets to prevent overfitting — all scores include confidence intervals

 Agentic

Coding and tool-use benchmarks

SWE Atlas

Test Writing & Codebase QnA

MCP Atlas

Tool use & agent coordination

SWE-Bench Pro

Public & Private repository tasks

 Frontier

Reasoning and knowledge limits

Humanity's Last Exam

Expert-level reasoning

MultiChallenge / MultiNRC

General & multilingual reasoning

EnigmaEval

Complex multi-step puzzles

SciPredict

Scientific forecasting

 Safety

Alignment and honesty evaluation

MASK

Honesty & truthfulness

Fortress

National security scenarios

PropensityBench

Sycophancy detection (lower=better)

Professional Reasoning

Legal & Finance domain expertise

 Multimodal

Audio and vision-language understanding

AudioMultiChallenge

Text & Audio output understanding

VisualToolBench


Visual tool-use tasks

VISTA

Vision-language comprehension

KEY INSIGHT

GPT-5.4 variants lead 3 of 4 benchmarks; Claude Opus 4.6 excels on private repos; Muse Spark emerges as a new contender



Performance Insights

GPT-5.4 Dominance

Leads 3 of 4 agentic benchmarks, including both SWE Atlas subtasks and SWE-Bench Pro Public.

Claude on Private Code

Opus 4.6 takes #1 on SWE-Bench Pro Private (47.1) — stronger on unseen repositories.


Muse Spark Emerges

New entrant places top 3 on both SWE-Bench splits and #1 on MCP Atlas tool use at 78.3.

All scores include $\pm$ confidence intervals

KEY INSIGHT

Even top scores are well below saturation — EnigmaEval under 24%, SciPredict under 26%, indicating vast room for improvement



Benchmark Saturation

Far From Solved

Top scores remain well below ceiling. HLE leader at only 44%, EnigmaEval at 24%, SciPredict at 25% — frontier models still have substantial room to improve.

Muse Spark #1 on MultiChallenge


The new entrant leads at 75.5%, ahead of Gemini-3.1-pro (71.4) and GPT-5.4-pro (69.2).

Google Leads SciPredict

Gemini-3-pro-preview at 25.27, with Claude Opus 4.5 (23.05) and Opus 4.1 (22.22) trailing.

KEY INSIGHT

Gemini holds all top-3 positions on AudioMultiChallenge Text Output; VISTA is competitive with GPT-5.4-pro within 1 point

AudioMultiChallenge**Text Output****Audio Output****Gemini Sweep**

All top-3 positions on AudioMultiChallenge **Text Output** held by Gemini

Vision-Language**VISTA****VisualToolBench****Tight Race on VISTA**

Gemini leads by only 0.76 points (54.65 vs 53.89);
VisualToolBench is effectively tied (29.17 vs 28.97)

KEY INSIGHT Safety alignment varies dramatically by task type — MASK honesty is near-solved, while national security (Fortress) remains critically unsolved at 8–13

Near-Ceiling Performance

MASK (Honesty)



PropensityBench (lower = better)



MASK is Near-Solved

Both Claude variants score above 96% — honesty alignment is approaching ceiling

Critically Low Scores

Fortress (National Security)



Professional Reasoning



Remote Labor Index



Fortress: Unsolved

National security scores of 8–13 show alignment remains critically weak in high-stakes domains

06

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

LANDSCAPE

KEY INSIGHT No single lab wins everywhere — Muse Spark emerges as a notable new entrant with top-3 placements on 6+ benchmarks

Benchmark	OpenAI	Google	Anthropic	Muse Spark	Other
SWE Atlas Test Writing	GPT-5.4 44.4				
SWE Atlas Codebase QnA	GPT-5.4 40.8				
MCP Atlas (Tool Use)				78.3	
SWE-Bench Pro Public	GPT-5.4 59.1			#2: 55.0	
SWE-Bench Pro Private			Claude 47.1	#2: 44.7	
Humanity's Last Exam	GPT-5.4 44.3			#2: 40.6	
MultiChallenge				75.5	
SciPredict		Gemini 25.3			
AudioMultiChallenge		Gemini 54.7			
VISTA		Gemini 54.7			
MASK (Honesty)			Claude 96.3		

Colored call = #1 ranked lab/model on that benchmark

Light tint = top-3 placement (not #1)

 GPT-5.4 Leads Agentic & Reasoning

Dominates SWE Atlas, SWE-Bench Pro Public, Humanity's Last Exam, and EnigmaEval — but does not lead multimodal or safety benchmarks.

 Gemini Sweeps Audio & Vision

Holds all top-3 positions on AudioMultiChallenge Text Output. Leads VISTA vision-language at 54.65. Only VisualToolBench is competitive with OpenAI.

 Claude Leads Safety; Fortress Is Weak

96%+ on MASK honesty, #1 on Professional Reasoning (Legal & Finance). Yet Fortress national security scores remain critically low (8–13).

 Benchmarks Remain Far From Saturation

HLE leader at 44%, EnigmaEval at 24%, SciPredict at 25%. These benchmarks still challenge frontier models — substantial room to improve.

 Muse Spark: A Strong New Competitor

Ranks top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5%, #2 on HLE at 40.6%, and strong showings on both SWE-Bench Pro splits and MCP Atlas. A notable new entrant challenging established labs.

Conclusion: Model specialization is the defining trend — labs must choose their battlegrounds.