
Scale SEAL Leaderboards

Frontier Model Performance Across 20+ AI Benchmarks

100+ models evaluated | 4 capability domains | Agentic · Frontier · Safety · Multimodal

Source: scale.com/leaderboard

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Scale Labs combines private datasets (to prevent overfitting) with open-source datasets for comparability. All scores include confidence intervals.

AGENTIC

Coding & Tool Use

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)

FRONTIER

Reasoning & Knowledge

- Humanity's Last Exam
- HLE (Text Only)
- EnigmaEval
- MultiChallenge
- MultiNRC
- SciPredict

SAFETY

Alignment & Honesty

- MASK (Honesty)
- PropensityBench
- Fortress (Nat'l Security)
- Prof. Reasoning – Legal
- Prof. Reasoning – Finance
- Remote Labor Index

MULTIMODAL

Audio & Vision-Language

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA (Vision-Language)
- VisualToolBench


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas & SWE-Bench Pro

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40

Note: Muse Spark (new entrant) places top 3 on MCP Atlas, SWE-Bench Pro (Public & Private), and 6+ other benchmarks.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%

Top scores on frontier reasoning benchmarks remain well below saturation, indicating significant room for improvement.



Gemini Models Sweep Audio and Vision-Language Benchmarks

AUDIO BENCHMARKS

Benchmark	Model	Score	Rank
AudioMultiChallenge (Text Output)	Gemini-3-pro-preview	54.65	#1
AudioMultiChallenge (Audio Output)	Gemini-3.1-flash-live	36.06	#1
AudioMultiChallenge (Text #2, #3)	Gemini-2.5-pro / flash	46.90 / 40.04	#2, #3


VISION-LANGUAGE BENCHMARKS

Benchmark	Model	Score	Rank
VISTA	Gemini 2.5 Pro Exp.	54.65	#1
VISTA	GPT-5.4-pro	53.89	#2
VisualToolBench	GPT-5.4	29.17	#1 (close)

KEY INSIGHT

Google’s Gemini variants hold all top-3 positions on AudioMultiChallenge (Text Output) and lead VISTA vision-language understanding. GPT-5.4-pro is the closest competitor on VISTA (53.89 vs 54.65) and edges Gemini on VisualToolBench (29.17 vs 28.97).

Claude Scores 96% on MASK Honesty; Fortress Scores Stay Low Across All Models



PROFESSIONAL REASONING — Claude Leads Legal & Finance

Benchmark	#1 Model	#2 Model	Score Gap
Legal	Claude Opus 4.6 (52.27)	—	Clear lead
Finance	Claude Opus 4.6 (53.28)	Muse Spark (52.44)	+0.84

Safety alignment varies dramatically by task: near-ceiling honesty (96%) vs. critically low national security (8–13).

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	Lab	#1 Model
SWE Atlas – Test Writing	OpenAI	GPT-5.4-xHigh
SWE Atlas – Codebase QnA	OpenAI	GPT 5.4 xHigh
MCP Atlas (Tool Use)	Muse Spark	Muse Spark
SWE-Bench Pro (Public)	OpenAI	GPT-5.4-pro
SWE-Bench Pro (Private)	Anthropic	Claude Opus 4.6
Humanity's Last Exam	OpenAI	GPT-5.4-pro
HLE (Text Only)	OpenAI	GPT-5.4-pro
EnigmaEval	OpenAI	GPT-5.4-pro
MultiChallenge	Muse Spark	Muse Spark
MultiNRC	OpenAI	GPT-5-pro
SciPredict	Google	Gemini-3-pro
AudioMultiChall. (Text)	Google	Gemini-3-pro
AudioMultiChall. (Audio)	Google	Gemini-3.1-flash
VISTA	Google	Gemini 2.5 Pro
VisualToolBench	OpenAI	GPT-5.4
MASK (Honesty)	Anthropic	Claude Opus 4.6
Prof. Reasoning – Legal	Anthropic	Claude Opus 4.6
Prof. Reasoning – Finance	Anthropic	Claude Opus 4.6
Remote Labor Index	Anthropic	Claude Opus 4.6

#1 WINS BY LAB

● OpenAI
8 benchmark(s) – #1

● Google
4 benchmark(s) – #1

● Anthropic
5 benchmark(s) – #1

● Muse Spark
2 benchmark(s) – #1

No single lab dominates all domains — model specialization is the defining competitive dynamic.

Key Takeaways: No Single Model Wins Everything

Specialization Is the New Frontier



GPT-5.4 leads agentic coding & frontier reasoning

Dominates SWE Atlas, SWE-Bench Pro (Public), Humanity's Last Exam (44.3%), and EnigmaEval — but does not lead multimodal or safety.



Gemini sweeps audio and vision-language

All top-3 positions on AudioMultiChallenge (Text Output) are Gemini variants. Leads VISTA at 54.65. Strongest multimodal portfolio.



Claude leads safety and professional reasoning

96%+ on MASK honesty. #1 on Legal (52.27) and Finance (53.28) professional reasoning. Fortress scores remain critically low across all models (8–13).



Benchmarks remain far from saturation

HLE top score: 44%. EnigmaEval: 24%. SciPredict: 25%. Substantial room for improvement across frontier reasoning and scientific forecasting.



Muse Spark: strong new competitor

Top 3 on 6+ benchmarks spanning coding, tool use, and reasoning. #1 on MultiChallenge at 75.5% and MCP Atlas at 78.3%.