

Scale SEAL Leaderboards

Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated • 4 capability domains • Agentic • Frontier • Safety • Multimodal

Source: scale.com/leaderboard

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

AGENTIC

Coding & Tool Use

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- SWE-Bench Pro (Public)
- SWE-Bench Pro (Private)
- MCP Atlas (Tool Use)

FRONTIER

Reasoning & Knowledge

- Humanity's Last Exam
- EnigmaEval
- MultiChallenge
- MultiNRC (Multilingual)
- SciPredict

SAFETY

Alignment & Honesty

- MASK (Honesty)
- PropensityBench
- Fortress (Nat'l Security)
- Professional Reasoning – Legal
- Professional Reasoning – Finance
- Remote Labor Index

MULTIMODAL

Audio & Vision-Language


- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA
- VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Bench Pro

Benchmark	#1 Model	Score	#2 Model	Score	#3 Model	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36	GPT-5.4-xHigh (Mini-SWE)	40.00	GPT-5.3-Codex-Xhigh	38.98
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80	GPT 5.4 xHigh (Mini-SWE)	36.30	Opus 4.6 (Claude Code)	33.30
MCP Atlas (Tool Use)	Muse Spark	78.30	Claude Opus 4.6	75.80	Gemini 3.1 Pro Preview	73.90
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10	Muse Spark	55.00	Claude Opus 4.6	51.90
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10	Muse Spark	44.70	GPT-5.4-pro (xHigh)	43.40


Key Insight: GPT-5.4 variants hold #1 on 3/5 agentic benchmarks. Claude Opus 4.6 leads SWE-Bench Pro (Private) — the only benchmark where a non-OpenAI model tops coding. Muse Spark appears in top 3 on 3/5 benchmarks.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%



Benchmarks remain far from saturation: HLE tops out at 44.3%, EnigmaEval at 23.8%, SciPredict at 25.3%. These benchmarks still challenge the best frontier models — substantial headroom remains. Muse Spark is the only non-OpenAI/Google model to reach #1 on a frontier benchmark (MultiChallenge at 75.5%).


Gemini Models Sweep Audio and Vision-Language Benchmarks




Gemini holds all top-3 positions on AudioMultiChallenge (Text Output) and leads both Audio Output and VISTA. GPT-5.4 edges Gemini-3.1-pro on VisualToolBench (29.17 vs 28.97) — the only multimodal benchmark where a non-Google model takes #1. Audio benchmarks remain well below saturation.

Claude Models Score 95–96% on MASK Honesty; Fortress Scores Stay Low

MASK (Honesty) — Near Ceiling



Fortress & PropensityBench — Critically Low



Safety alignment varies dramatically by task type: MASK honesty scores approach ceiling (96%), while Fortress national security scores remain in single digits (8–13). PropensityBench (lower = better) shows o3 at 10.5. Claude also leads Professional Reasoning — Legal (52.27%) and Finance (53.28%).

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

		OpenAI	Google	Anthropic	Muse Spark
SWE Atlas – Test Writing	GPT-5.4-xHigh	● #1	—	—	—
SWE Atlas – Codebase QnA	GPT 5.4 xHigh	● #1	—	—	—
MCP Atlas (Tool Use)	Muse Spark	—	—	—	● #1
SWE-Bench Pro (Public)	GPT-5.4-pro	● #1	—	—	—
SWE-Bench Pro (Private)	Claude Opus 4.6	—	—	● #1	—
Humanity's Last Exam	GPT-5.4-pro	● #1	—	—	—
EnigmaEval	GPT-5.4-pro	● #1	—	—	—
MultiChallenge	Muse Spark	—	—	—	● #1
SciPredict	Gemini-3-pro	—	● #1	—	—
AudioMultiChallenge (Text)	Gemini-3-pro	—	● #1	—	—
AudioMultiChallenge (Audio)	Gemini-3.1-flash	—	● #1	—	—
VISTA	Gemini 2.5 Pro	—	● #1	—	—
MASK (Honesty)	Claude Opus 4.6	—	—	● #1	—
PropensityBench	o3	● #1	—	—	—
Fortress	Claude Sonnet 4	—	—	● #1	—
Prof. Reasoning – Legal	Claude Opus 4.6	—	—	● #1	—
Prof. Reasoning – Finance	Claude Opus 4.6	—	—	● #1	—

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

GPT-5.4 leads agentic coding and frontier reasoning

Holds #1 on SWE Atlas, SWE-Bench Pro (Public), Humanity's Last Exam, and EnigmaEval — but does not lead multimodal or safety benchmarks.

Gemini dominates audio and vision-language understanding

Sweeps all top-3 positions on AudioMultiChallenge (Text Output) and leads VISTA. Also leads SciPredict — the only frontier benchmark won by Google.

Claude leads safety alignment and professional reasoning

Near-ceiling MASK honesty scores (96%+), #1 on both Professional Reasoning benchmarks (Legal & Finance), and tops SWE-Bench Pro (Private) — the only coding benchmark not led by OpenAI.

Key benchmarks remain far from saturation

Humanity's Last Exam (44%), EnigmaEval (24%), SciPredict (25%) — frontier models still have substantial room for improvement on the hardest benchmarks.

New entrant Muse Spark is a strong competitor

Ranks top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5% and #1 on MCP Atlas at 78.3%. Signals increasing competitive pressure.