

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across Agentic, Frontier, Safety, and Multimodal domains (Scale Labs leaderboards)

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

- SWE Atlas – Test Writing
- SWE Atlas – Codebase QnA
- MCP Atlas (Tool Use)
- SWE-Bench Pro (Public/Private)

Frontier

- Humanity's Last Exam (HLE)
- EnigmaEval
- MultiChallenge
- SciPredict

Safety

- MASK (Honesty)
- Fortress (National Security)
- PropensityBench (lower=better)
- Professional Reasoning (Legal/Finance)

Multimodal

- AudioMultiChallenge (Text Output)
- AudioMultiChallenge (Audio Output)
- VISTA (Vision-Language)
- VisualToolBench


Benchmark	#1	Score	#2	Score	#3	Score
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)	44.36 ±6.04	GPT-5.4-xHigh (Mini-SWE)	40.00 ±6.00	GPT-5.3-Codex-Xhigh	38.98 ±6.12
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)	40.80 ±5.10	GPT 5.4 xHigh (Mini-SWE)	36.30 ±4.90	Opus 4.6 (Claude Code)	33.30 ±5.00
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)	59.10 ±3.56	Muse Spark	55.00 ±3.60	Claude Opus 4.6	51.90 ±3.61
SWE-Bench Pro (Private)	Claude Opus 4.6	47.10 ±6.07	Muse Spark	44.70 ±6.05	GPT-5.4-pro (xHigh)	43.40 ±6.03

Interpretation

- OpenAI leads SWE Atlas and SWE-Bench Pro (Public)
- Anthropic leads SWE-Bench Pro (Private) on private repos
- Muse Spark places top-2 on both SWE-Bench Pro leaderboards

Scores as reported in source brief and (+/-) indicates confidence interval






Benchmark	#1 Model	Score
Humanity's Last Exam	GPT-5.4-pro	44.32 ±1.95
EnigmaEval	GPT-5.4-pro	23.82 ±2.43
MultiChallenge	Muse Spark	75.52 ±4.05

Unsaturated benchmarks

- HLE tops out at 44–45%
- EnigmaEval remains under 24%
- SciPredict tops out at 25.27% (Gemini-3-pro-preview)

🔊 AudioMultiChallenge - Audio Output

Evaluating spoken dialogue systems in multi-turn interaction



• AudioMultiChallenge (Text Output): Gemini-3-pro-preview leads at 54.65 ±4.57

• AudioMultiChallenge (Audio Output): Gemini-3.1-flash-live-preview leads at 36.06 ±4.41

• VISTA: Gemini 2.5 Pro Experimental leads at 54.65 ±1.46 (GPT-5.4-pro close at 53.89 ±2.02)


👁️ VISTA

Vision-Language Understanding benchmark for multimodal models




[View Full Ranking →](#)

MASK (Honesty)



Near-ceiling honesty

Fortress (NatSec)



Low national-security safety

Additional safety signal

- PropensityBench (lower=better): o3-2025-04-16 at 10.50 ±0.60

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	#1 Lab / Model
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex CLI)
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex)
MCP Atlas (Tool Use)	Muse Spark
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh)
SWE-Bench Pro (Private)	Claude Opus 4.6
Humanity's Last Exam	GPT-5.4-pro
EnigmaEval	GPT-5.4-pro
MultiChallenge	Muse Spark
MultiNRC (Multilingual)	GPT-5-pro
SciPredict	Gemini-3-pro-preview
AudioMultiChallenge (Text)	Gemini-3-pro-preview
AudioMultiChallenge (Audio)	Gemini-3.1-flash-live-preview
VISTA	Gemini 2.5 Pro Experimental
VisualToolBench	GPT-5.4
MASK (Honesty)	Claude Opus 4.6
PropensityBench	o3-2025-04-16
Fortress (NatSec)	#1 not listed (score 8.24 ±1.93)
Professional Reasoning – Legal	Claude Opus 4.6
Professional Reasoning – Finance	Claude Opus 4.6

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

- GPT-5.4 leads agentic coding and key frontier reasoning tasks, but not multimodal.
- Gemini dominates audio and vision-language (top positions across AudioMultiChallenge + VISTA).
- Claude leads safety/honesty (MASK ~96%) and professional reasoning (Legal + Finance).
- Benchmarks remain unsaturated: HLE ~44–45%, EnigmaEval ~24%, SciPredict ~25%.
- Muse Spark is a strong new entrant: #1 on MultiChallenge at 75.52% and top-3 on SWE-Bench Pro + MCP Atlas.