

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

100+ models evaluated across Agentic, Frontier, Safety, and Multimodal benchmarks

Comparative analysis for ML researchers, AI lab leadership, and engineering managers

Scale Labs benchmark suite reveals domain-specific leaders rather than a single universal winner.

Data window: Scale SEAL results in source brief



20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic

- SWE Atlas Test Writing
- SWE Atlas Codebase QnA
- SWE-Bench Pro (Public/Private)
- MCP Atlas
- Terminal-Bench
- RE-Bench
- GadgetBench

Frontier

- Humanity's Last Exam
- EnigmaEval
- SciPredict
- MultiChallenge
- ARC-AGI-2
- MMLU-Pro

Safety

- MASK (Honesty)
- PropensityBench
- Fortress
- Professional Reasoning (Legal/Finance)
- Remote Labor Index

Multimodal

- AudioMultiChallenge (Text)
- AudioMultiChallenge (Audio)
- VISTA
- MMMU
- SimpleVQA


GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE-Ber

Top-3 model scores from source brief

Benchmark	#1	#2	#3
SWE Atlas Test Writing	GPT-5.4: 72.0	Claude Opus 4.6: 64.0	Gemini-3.1-pro: 62.0
SWE Atlas Codebase QnA	GPT-5.4: 48.2	Claude Opus 4.6: 45.5	Gemini-3.1-pro: 44.6
SWE-Bench Pro (Public)	GPT-5.4 (high): 75.8	Claude Opus 4.6 (thinking): 75.5	Muse Spark: 74.0
SWE-Bench Pro (Private)	GPT-5.4 (high): 68.6	Claude Opus 4.6 (thinking): 67.6	Muse Spark: 67.1


Insight: OpenAI leads all four agentic coding tasks, with Claude close behind and Muse Spark entering top-3 on both SWE-Bench Pro tracks.

GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below




Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge
(Text)




Audio Text Output
#1 Gemini-2.5-pro
#2 Gemini-3.1-pro
#3 Gemini-2.5-flash

AudioMultiChallenge
(Audio)




Audio Audio Output
#1 Gemini-3.1-pro
#2 Gemini-2.5-pro
#3 Gemini-2.5-flash

VISTA



VISTA
#1 Gemini-2.5-pro
#2 o3-2025-04-16
#3 Gemini-2.5-flash

Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Score



Contrast in alignment performance

- MASK (honesty): near ceiling at 96%
- Fortress (nat. security): very low (8–13)

PropensityBench (lower=better)

- #1 o3-2025-04-16: 10.50 ±0.60

Implication: alignment capability is highly benchmark-dependent.

OpenAI, Google, and Anthropic Each Claim #1 Across Different Domains

Benchmark	#1 Lab/Model Family
SWE Atlas TW	OpenAI
SWE Atlas QnA	OpenAI
SWE-Pro Public	OpenAI
SWE-Pro Private	OpenAI
MCP Atlas	OpenAI
Terminal-Bench	OpenAI
HLE	OpenAI
EnigmaEval	OpenAI
SciPredict	OpenAI
SimpleVQA	OpenAI
AudioMC Text	Google
AudioMC Audio	Google
VISTA	Google
MASK	Anthropic
PR-Legal	Anthropic
PR-Finance	Anthropic
RLI	Anthropic
MultiChallenge	Muse Spark
MultiChallenge	Anthropic
MultiChallenge	Muse Spark

#1 share across listed benchmarks:

- OpenAI: 10
- Google: 3
- Anthropic: 4
- Muse Spark: 1

No single lab dominates all domains.
Specialization now defines frontier competition.

Key Takeaways: No Single Model Wins Everything — Specialization Is the Key

- GPT-5.4 leads agentic coding and frontier reasoning (e.g., SWE Atlas, SWE-Bench Pro, HLE 44.3, EnigmaEval 24.1).
- Gemini variants dominate multimodal/audio benchmarks, sweeping AudioMultiChallenge leaderboards and leading VISTA.
- Claude models lead safety and professional reasoning, with 96%+ MASK scores and #1 on legal/finance reasoning.
- Critical benchmarks remain unsaturated: HLE 44.3%, EnigmaEval 24.1%, SciPredict 25.0%.
- Muse Spark emerges as a strong new entrant, with top-3 placement on 6+ benchmarks and #1 on MultiChallenge (75.5%).

Implication for deployment strategy: choose model portfolios by domain-specific performance, not aggregate branding.