

Scale SEAL Leaderboards: Frontier Model Performance Across 20+ Benchmarks

A comparative analysis of top-
performing AI models across Scale
Labs' benchmark suite

Evaluating 100+ models across
agentic, frontier, safety, and
multimodal capabilities.

20+ Benchmarks Spanning Four Capability Domains Evaluate 100+ Models

Agentic (Coding & Tool)

SWE Atlas
MCP Atlas
SWE-Bench Pro

Frontier (Reasoning)

Humanity's Last Exam
EnigmaEval
MultiChallenge
SciPredict

Safety (Alignment)

MASK
PropensityBench
Fortress
Prof. Reasoning
RLI


Multimodal (Audio/Vision)

AudioMultiChallenge
VISTA
VisualToolBench

GPT-5.4 Dominates Agentic Coding: Top Scores on SWE Atlas and SWE- Bench Pro

Benchmark	#1 Model	#2 Model	#3 Model
SWE Atlas – Test Writing	GPT-5.4-xHigh (Codex) (44.36)	GPT-5.4-xHigh (Mini-SWE) (40.00)	GPT-5.3-Codex-Xhigh (38.98)
SWE Atlas – Codebase QnA	GPT 5.4 xHigh (Codex) (40.80)	GPT 5.4 xHigh (Mini-SWE) (36.30)	Opus 4.6 (Claude Code) (33.30)
SWE-Bench Pro (Public)	GPT-5.4-pro (xHigh) (59.10)	Muse Spark (55.00)	Claude Opus 4.6 (51.90)
SWE-Bench Pro (Private)	Claude Opus 4.6 (47.10)	Muse Spark (44.70)	GPT-5.4-pro (xHigh) (43.40)


GPT-5.4-Pro Leads Humanity's Last Exam at 44.3% — Scores Remain Below 50%




Note: EnigmaEval (under 24%) and Humanity's Last Exam show frontier models are well below saturation.

Gemini Models Sweep Audio and Vision-Language Benchmarks

AudioMultiChallenge (Text Output)




VISTA (Vision-Language)



Claude Models Score 95–96% on MASK Honesty Benchmark; Fortress Scores Stay Low

Score



Alignment Variance:

Safety alignment varies dramatically by task type. Honesty is near ceiling, while National Security (Fortress) is critically low.

PropensityBench:

o3-2025-04-16 scores 10.50 (lower is better).

OpenAI, Google, and Anthropic

Each Claim #1 Across Different

Benchmark	#1 Model	Leading Lab
SWE Atlas Test Writing (Coding)	GPT-5.4-xHigh	OpenAI
Humanity's Last Exam (Reasoning)	GPT-5.4-pro	OpenAI
AudioMultiChallenge (Audio)	Gemini-3-pro-preview	Google
VISTA (Vision-Language)	Gemini 2.5 Pro Exp	Google
MASK (Honesty/Safety)	Claude Opus 4.6	Anthropic
Prof. Reasoning – Legal	Claude Opus 4.6	Anthropic
SWE-Bench Pro (Private)	Claude Opus 4.6	Anthropic
MultiChallenge	Muse Spark	Muse

Key Takeaways: No Single Model Wins Everything — Specialization Is the New Frontier

- GPT-5.4 leads agentic/reasoning but does not dominate multimodal tasks.
- Gemini dominates audio/vision, holding top positions across AudioMultiChallenge and VISTA.
- Claude leads safety/honesty, scoring 96% on MASK and #1 on Professional Reasoning.
- Key benchmarks like HLE (44%), EnigmaEval (24%), and SciPredict (25%) remain far from saturation, showing substantial room for improvement.
- New entrant Muse Spark is a strong competitor, ranking top 3 on 6+ benchmarks spanning coding, tool use, and reasoning — including #1 on MultiChallenge at 75.5%.